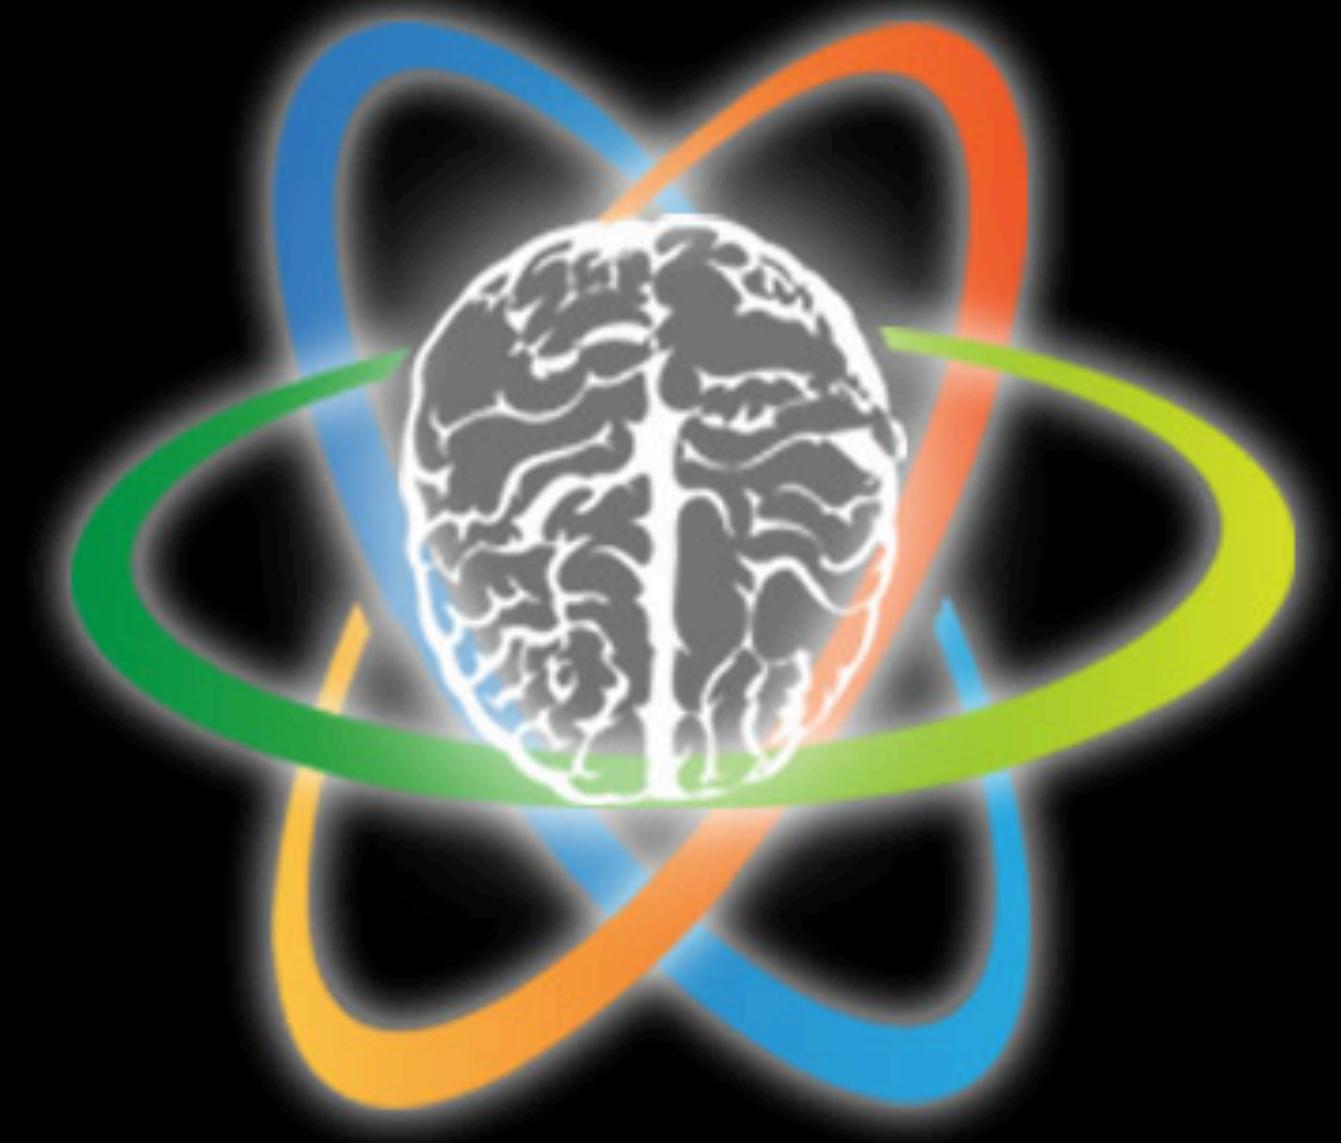




# Neural Networks Jumpstart

Daniel Hinojosa

NFUS



# ÜBERCONF

THE ULTIMATE CONFERENCE FOR SOFTWARE DEVELOPERS AND ARCHITECTS

July 15 - 18, 2025 — Denver

Save \$250 - Register by March 31st

**Gateway Software Symposium**

April 4 - 5, 2025

**Central Ohio Software Symposium**

April 11 - 12, 2025

**New England Software Symposium**

May 2 - 3, 2025

**CODE REMIX**

Summit Miami | May 12 - 14, 2025

**Great Lakes Software Symposium**

May 16 - 17, 2025

**ÜBERCONF**

July 15 - 18, 2025

**Central Iowa Software Symposium**

September 5 - 6, 2025

**Twin Cities Software Symposium**

September 12 - 13, 2025

**Northern Virginia Software Symposium**

September 19 - 20, 2025

**Pacific Northwest Software Symposium**

October 10 - 11, 2025

**TECHLEADER SUMMIT**

December 3 - 5, 2025

**ARCHCONF**

December 8 - 11, 2025

Each NFJS conference registration now comes bundled with an NFJS Virtual Workshop Subscription!

# Applications

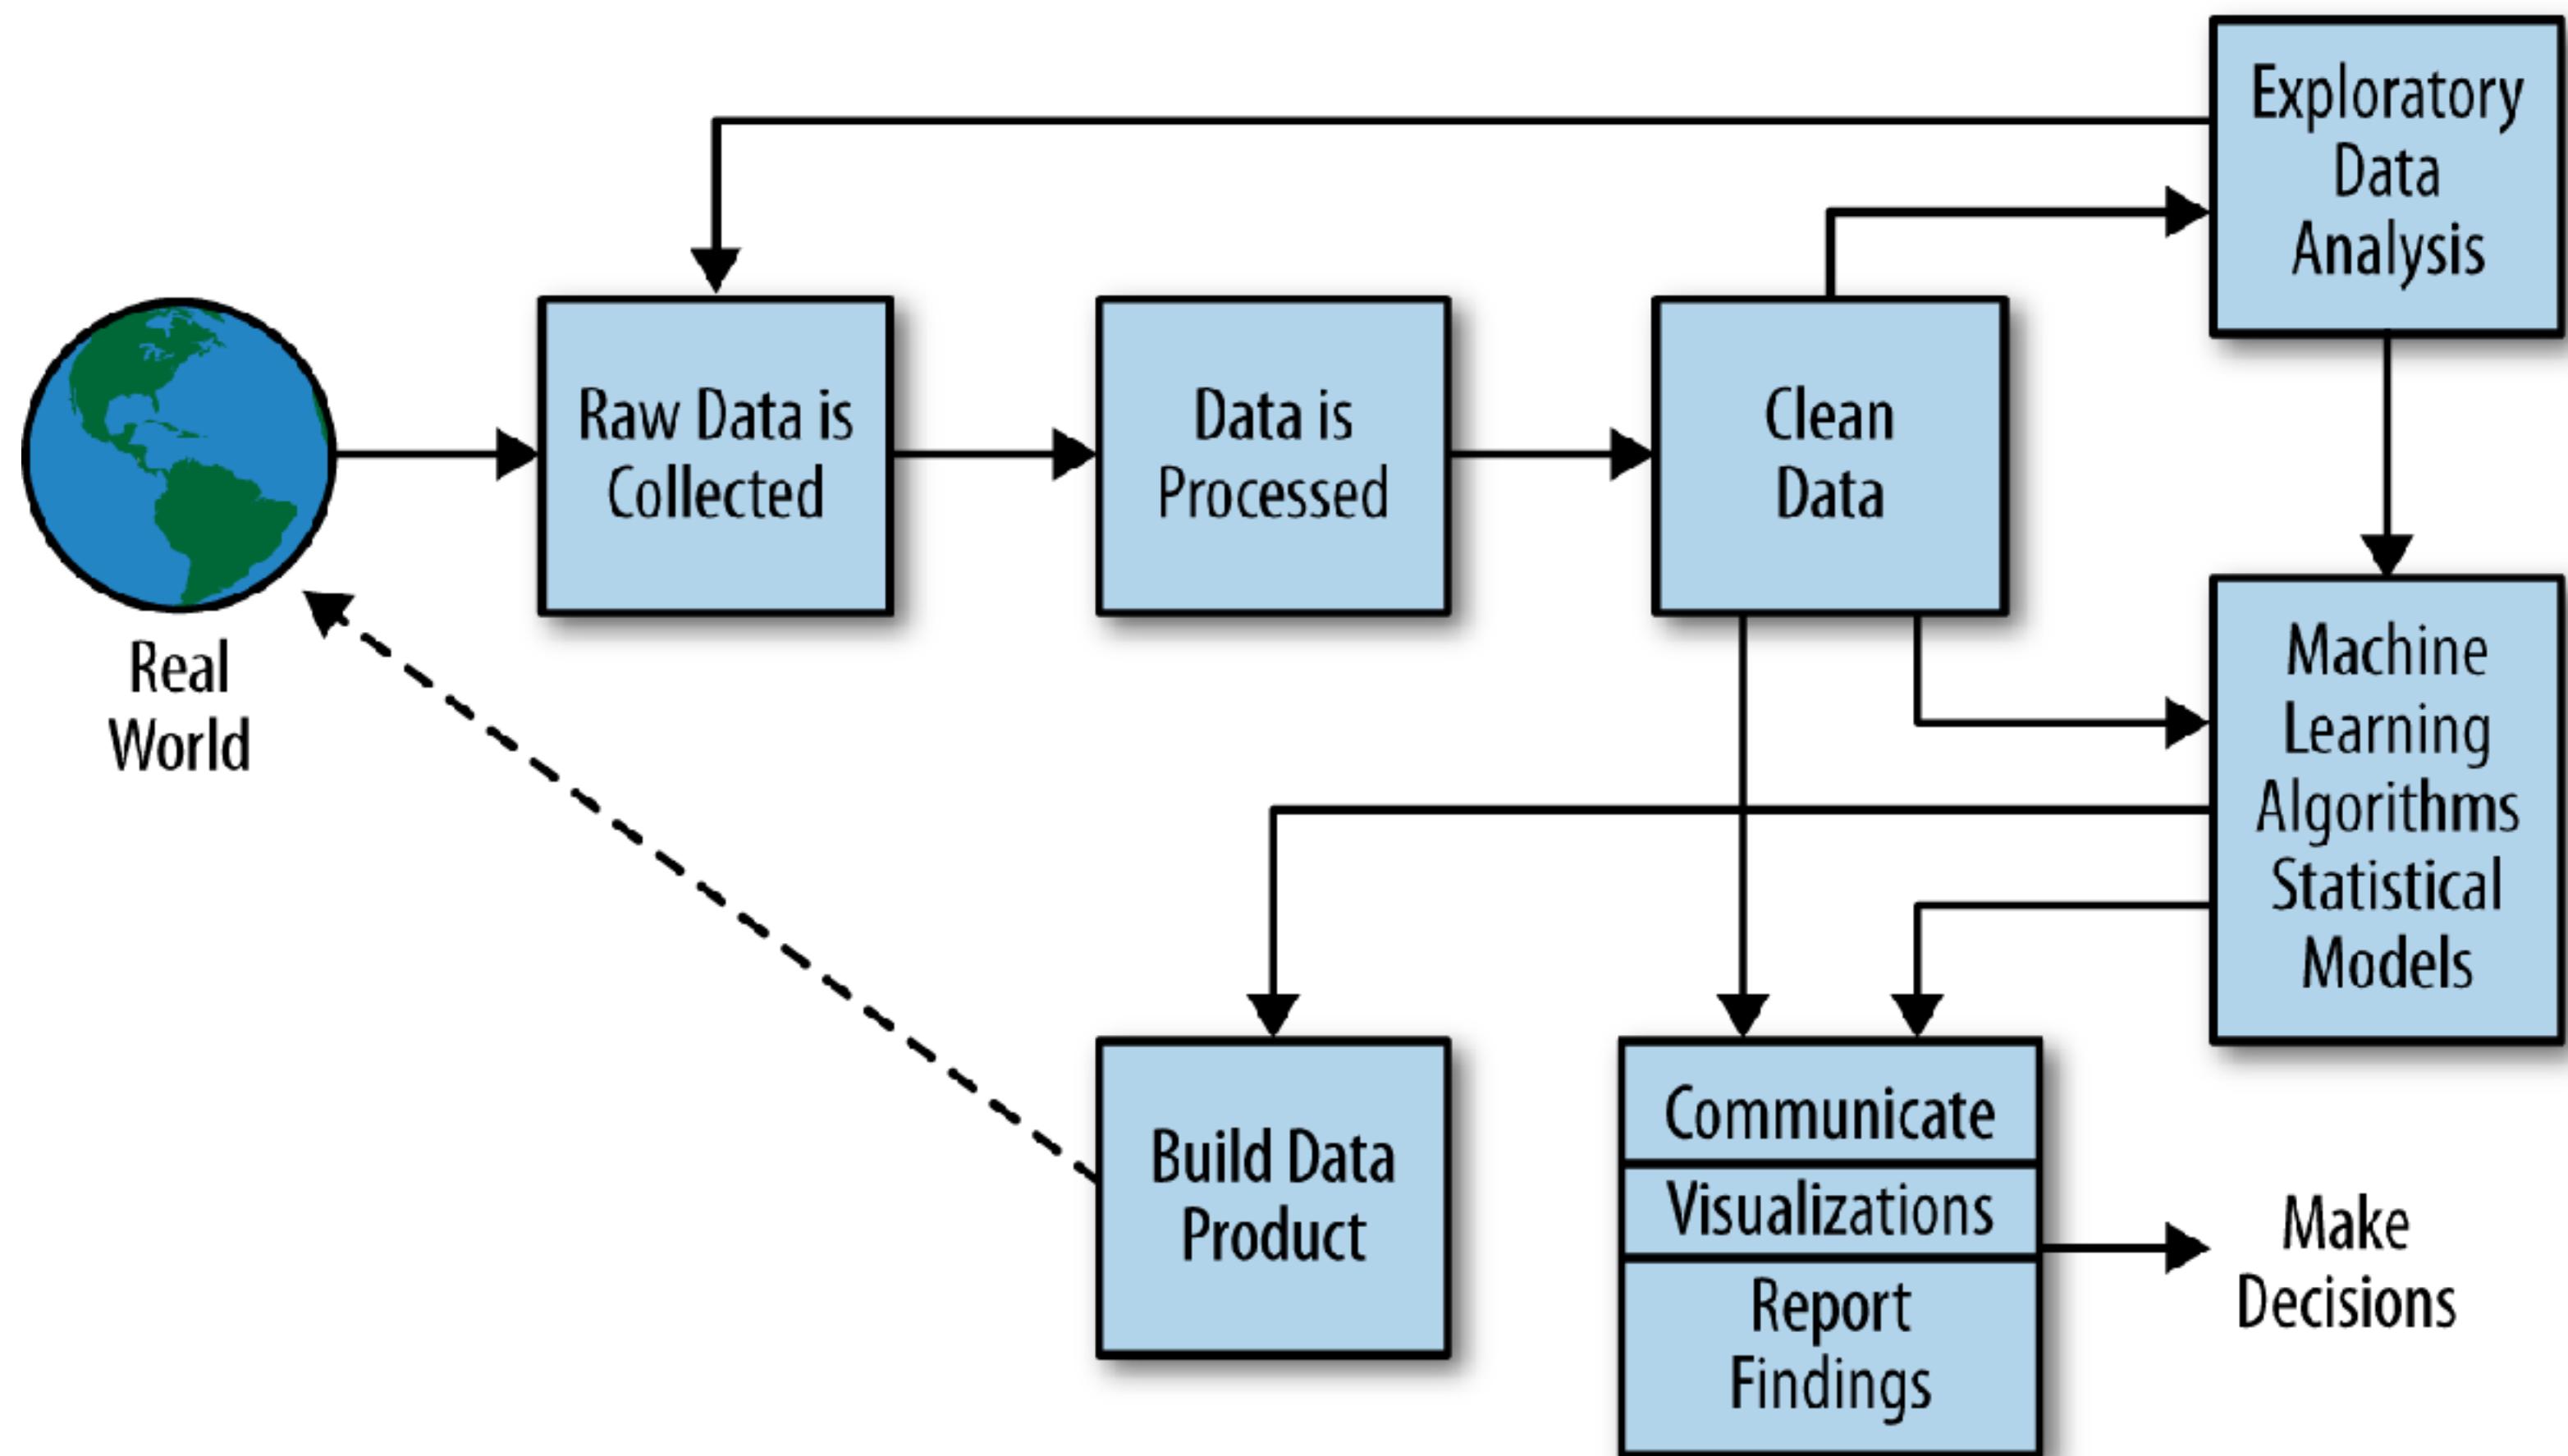
Fraud Detection

Financial Approvals

Picture Recognition

Your Idea Here

# Process of Machine Learning



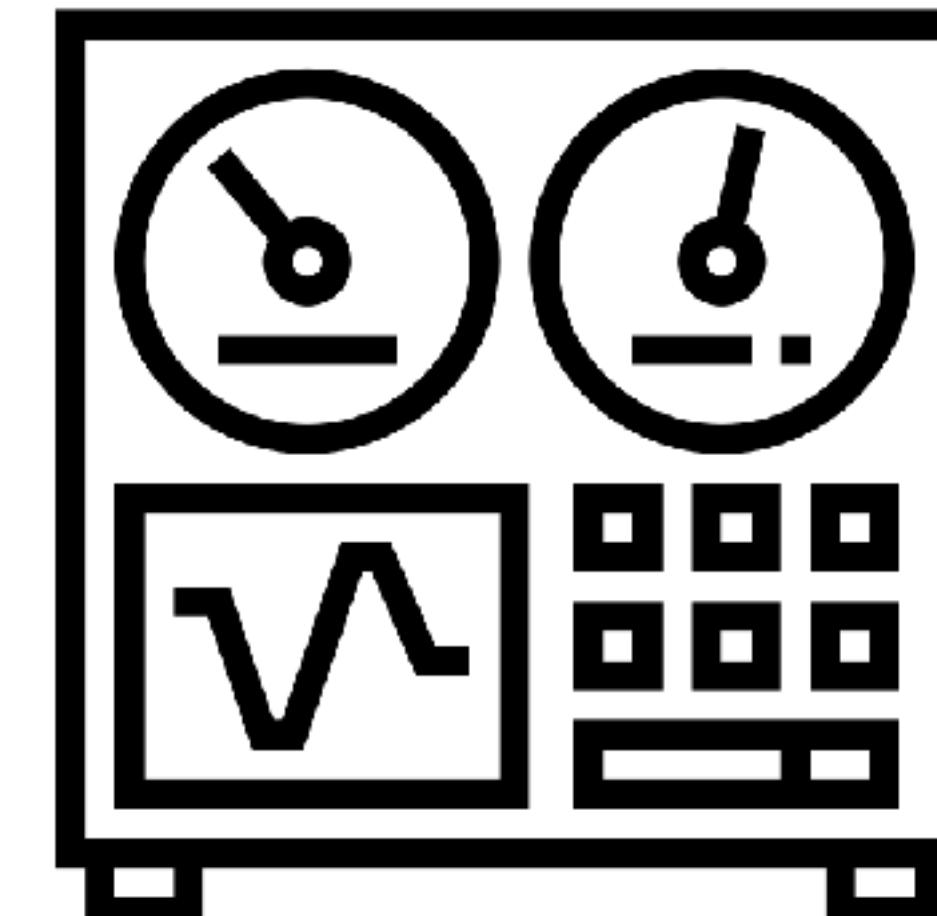
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1

training phase



<b>id</b>	<b>age</b>	<b>sex</b>	<b>trestbps</b>	<b>risk</b>
0	20	1	145	1
1	45	1	130	1
2	43	0	130	0
3	33	1	150	0

**data**

**target**

training →

testing ↗

<b>id</b>	<b>age</b>	<b>sex</b>	<b>trestbps</b>	<b>target</b>
0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1
10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1

<b>id</b>	<b>age</b>	<b>sex</b>	<b>trestbps</b>	<b>target</b>
0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1
10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1

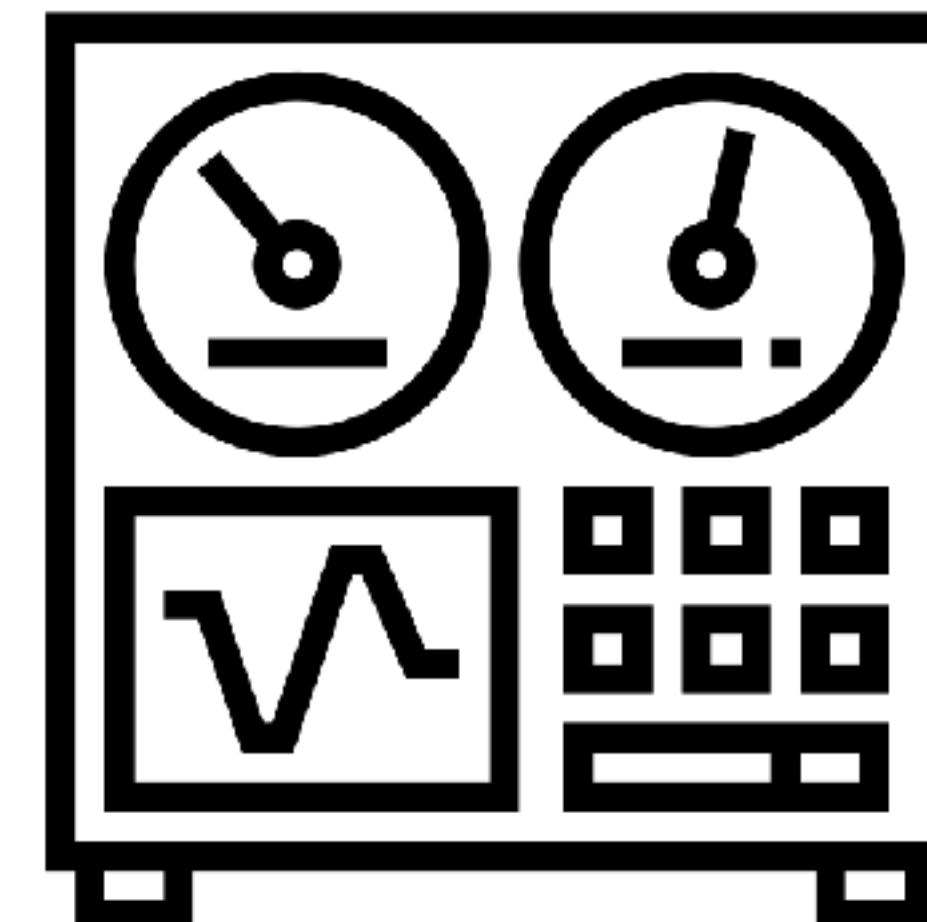
**random distribution between training and testing**

# training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

# testing

10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1



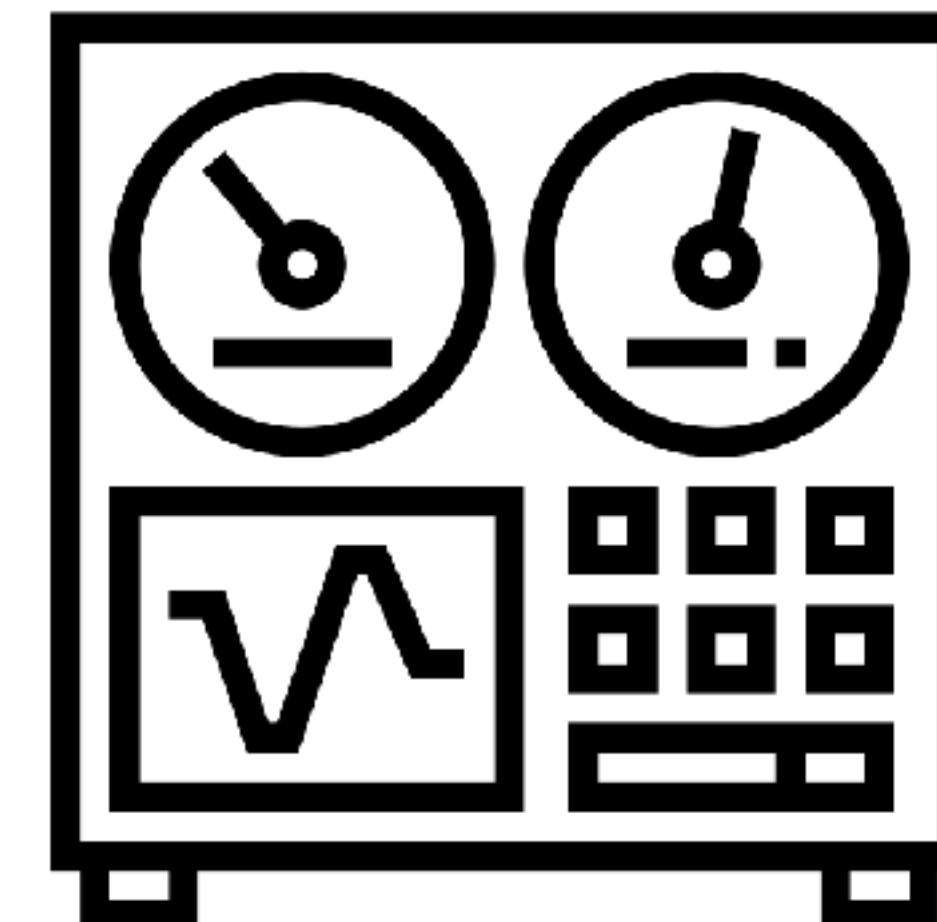
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1

model



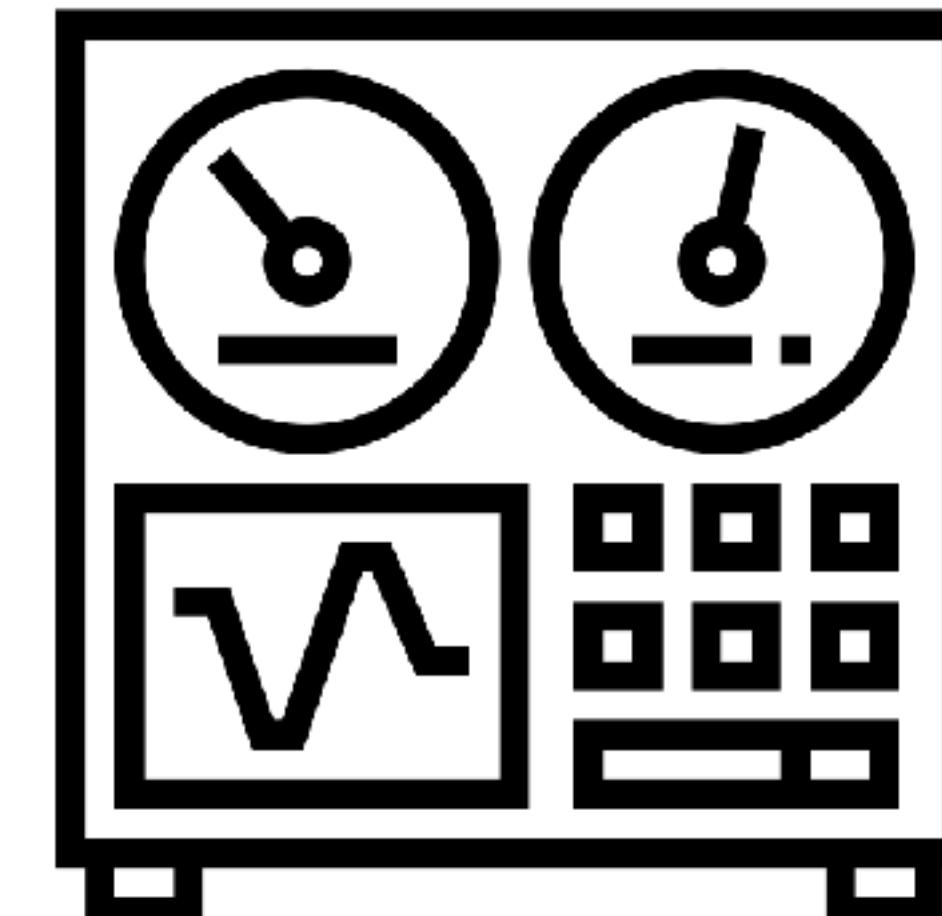
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1

training phase

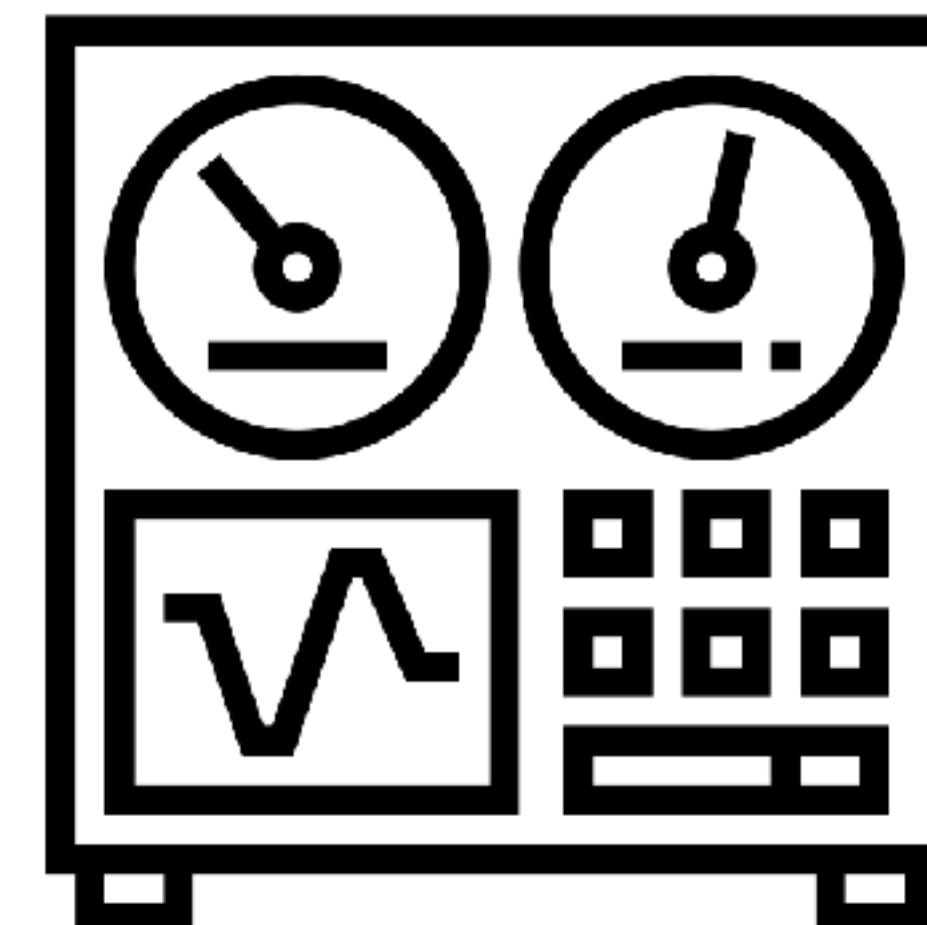


# training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

# testing

10	54	1	239	1
11	48	0	275	0
12	49	1	266	1
13	64	1	211	0
14	58	0	283	1



# training

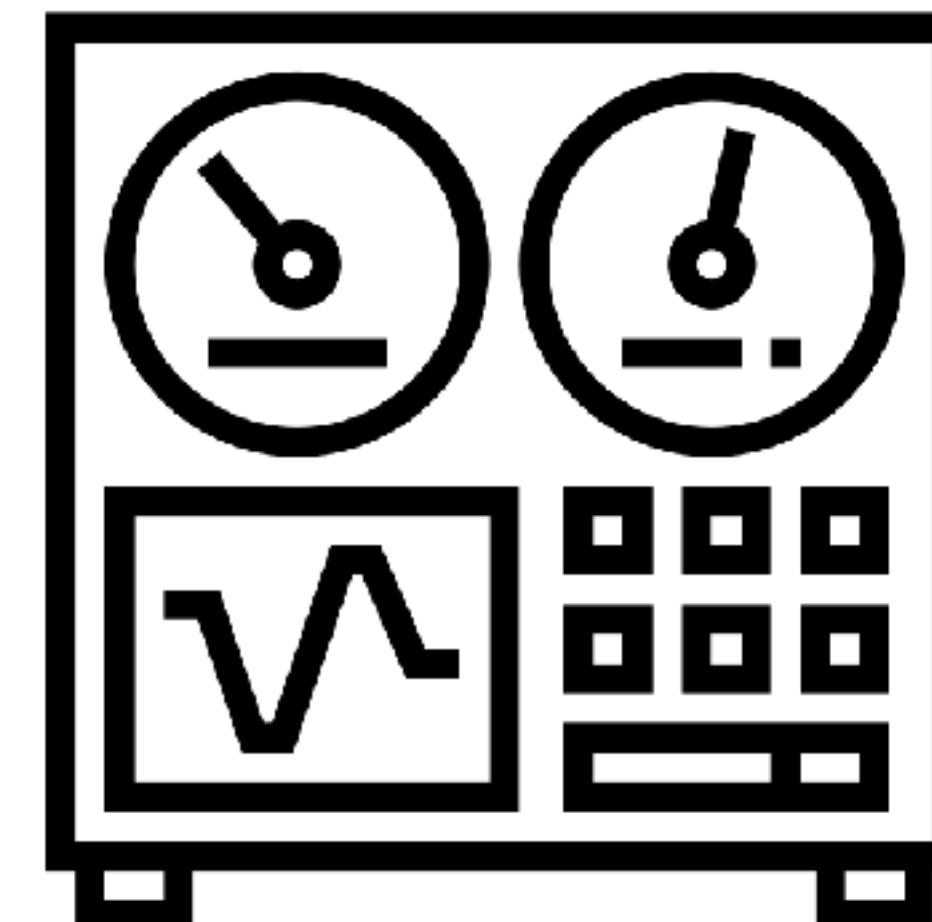
0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

# testing

10	54	1	239
11	48	0	275
12	49	1	266
13	64	1	211
14	58	0	283

# actual

1
0
1
0
1



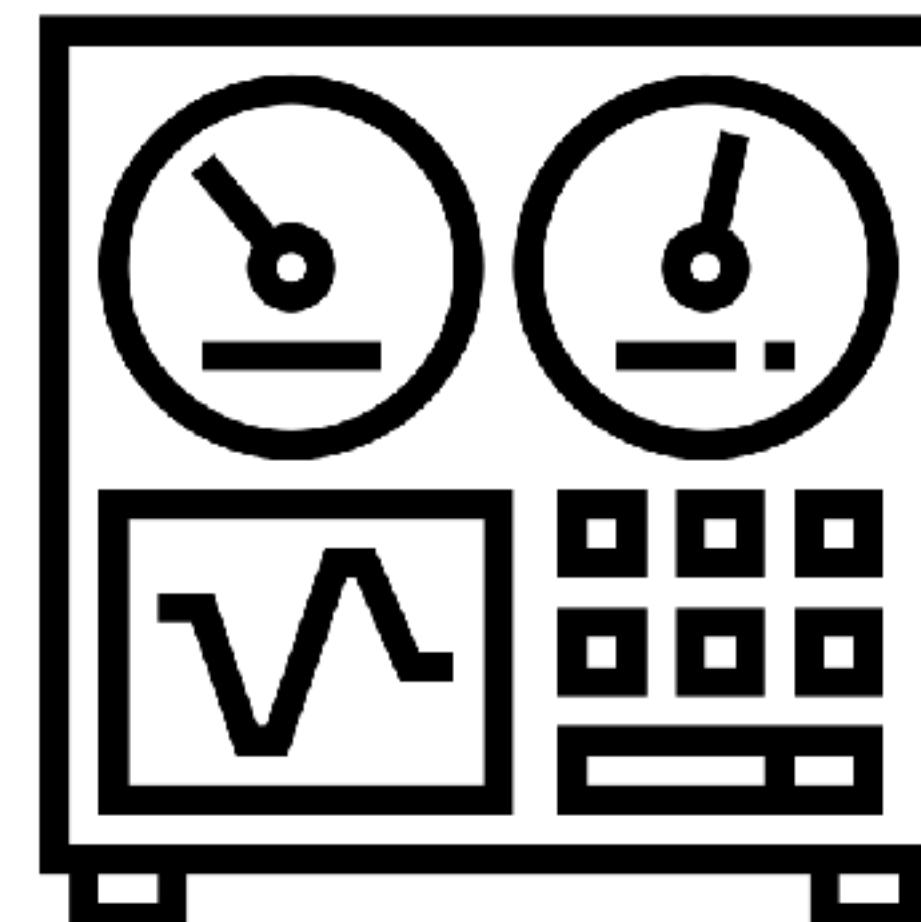
reserve the actual later for verification

## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239
11	48	0	275
12	49	1	266
13	64	1	211
14	58	0	283



testing phase

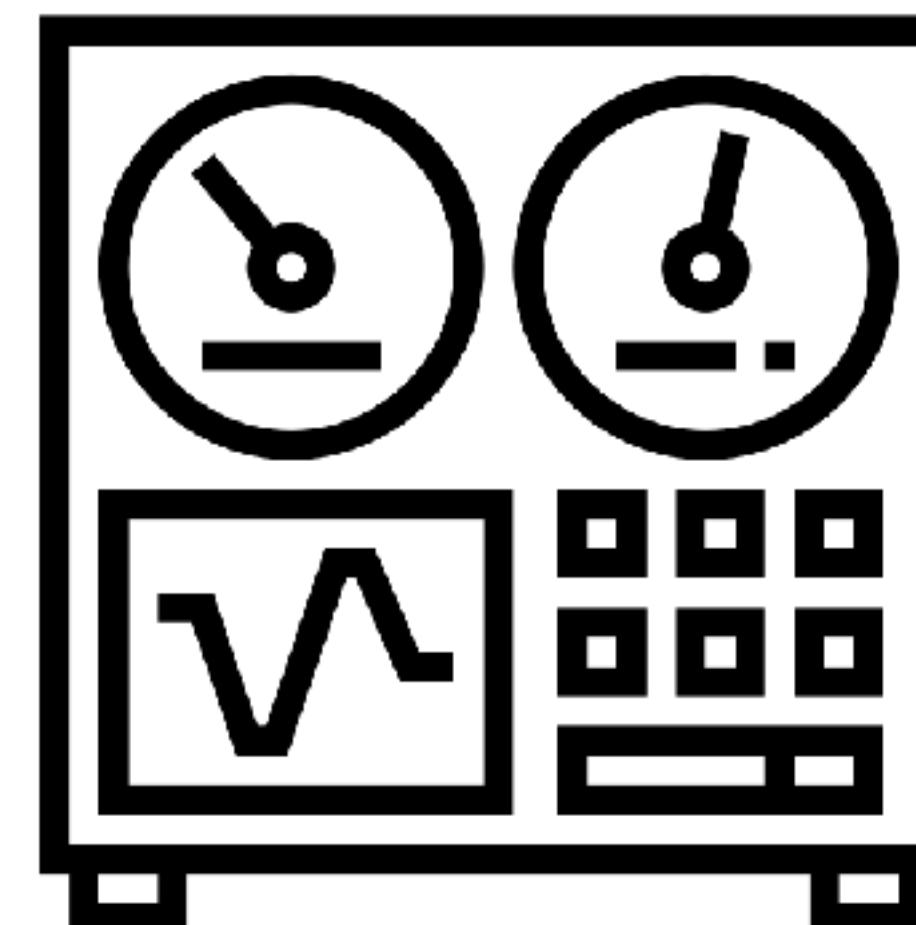
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239
11	48	0	275
12	49	1	266
13	64	1	211
14	58	0	283

!!!!!!



generated  
result

1
0
1
1
1

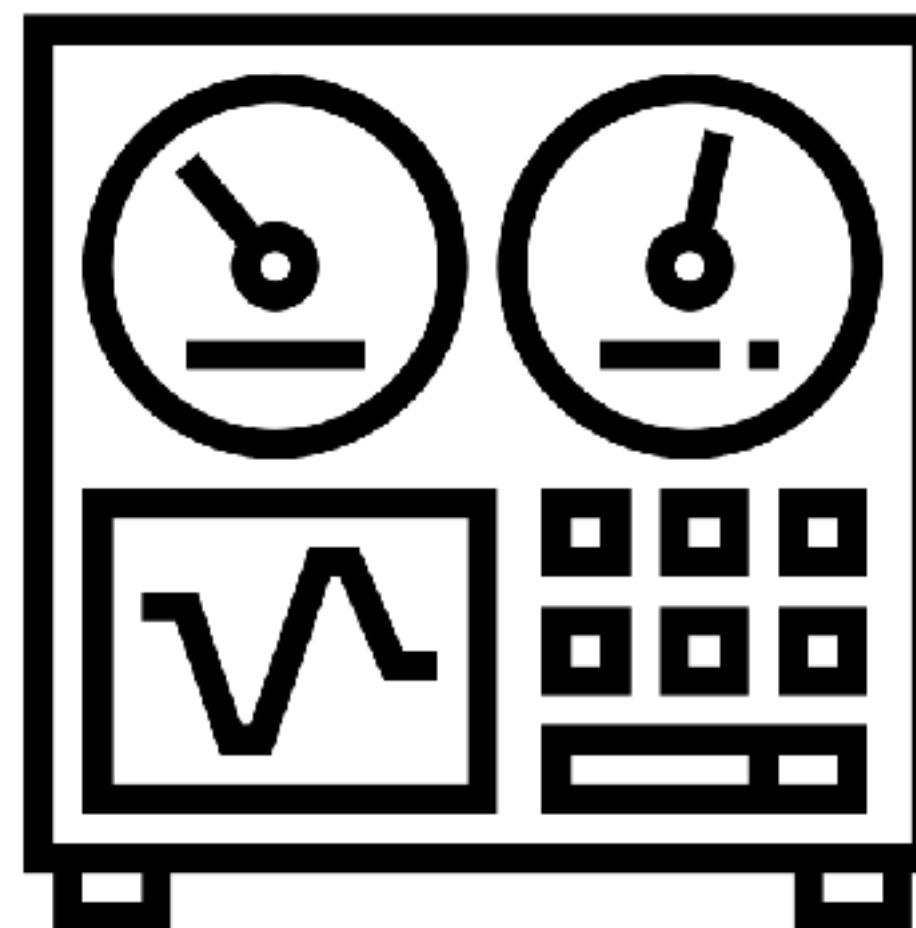
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239
11	48	0	275
12	49	1	266
13	64	1	211
14	58	0	283

!!!!!!



generated

result

actual

1
0
1
1
1

1
0
1
0
1

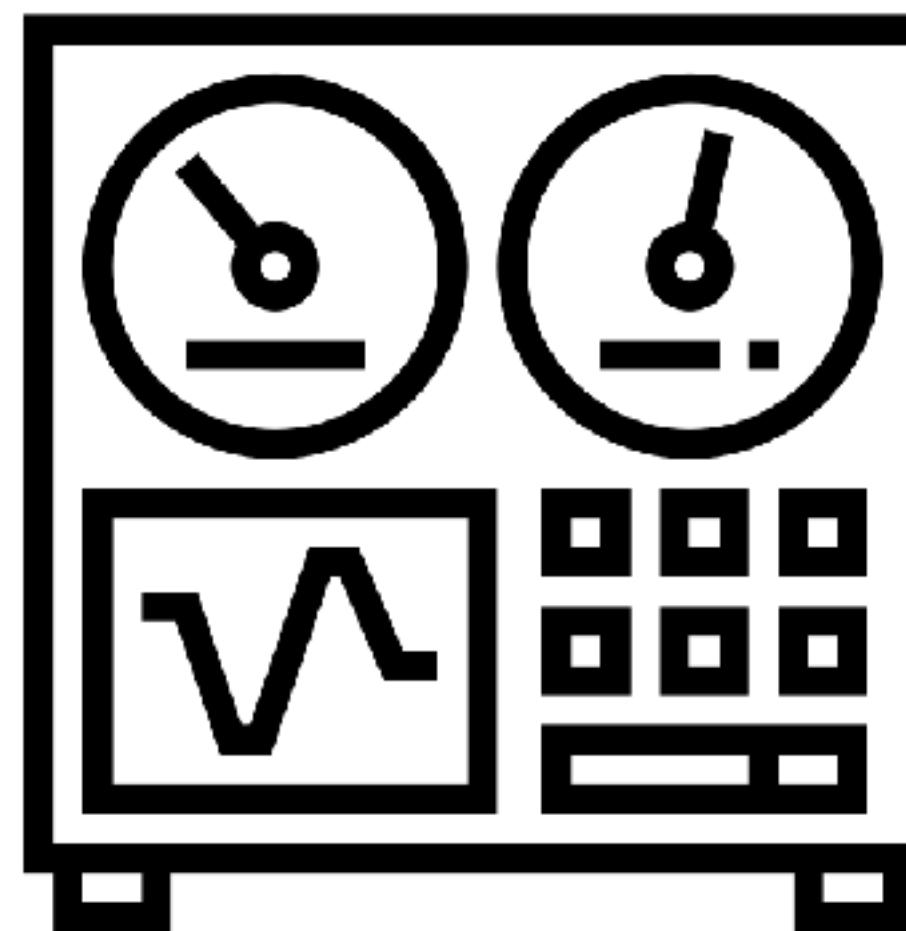
## training

0	63	1	233	1
1	37	1	250	0
2	41	0	204	1
3	56	1	236	1
4	57	0	354	0
5	57	1	192	0
6	56	0	294	0
7	44	1	263	1
8	52	1	199	1
9	57	1	168	1

## testing

10	54	1	239
11	48	0	275
12	49	1	266
13	64	1	211
14	58	0	283

!!!!!!



generated

result

actual

1
0
1
1
1

1
0
1
0
1

How did we do?

generated  
result      actual

1
0
1
1
1

1
0
1
0
1

How did we do?

# Tensor

Scalar

Rank-0

1

Vector

Rank-1

1
2

Matrix

Rank-2

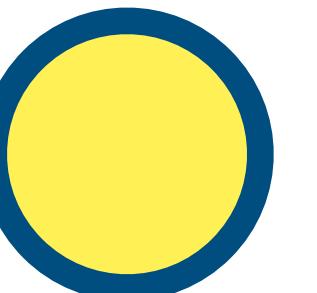
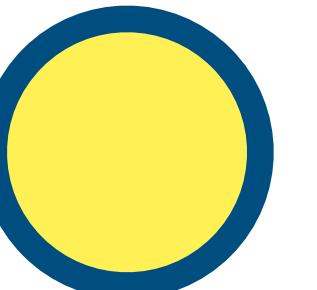
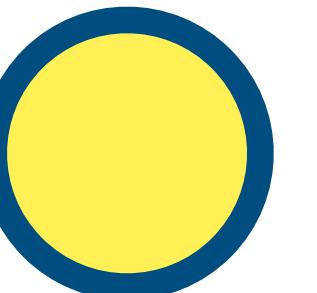
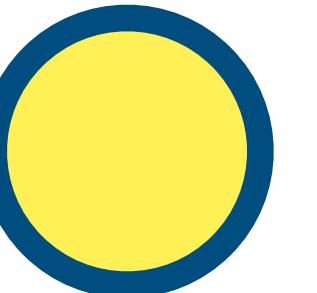
1	1
2	1

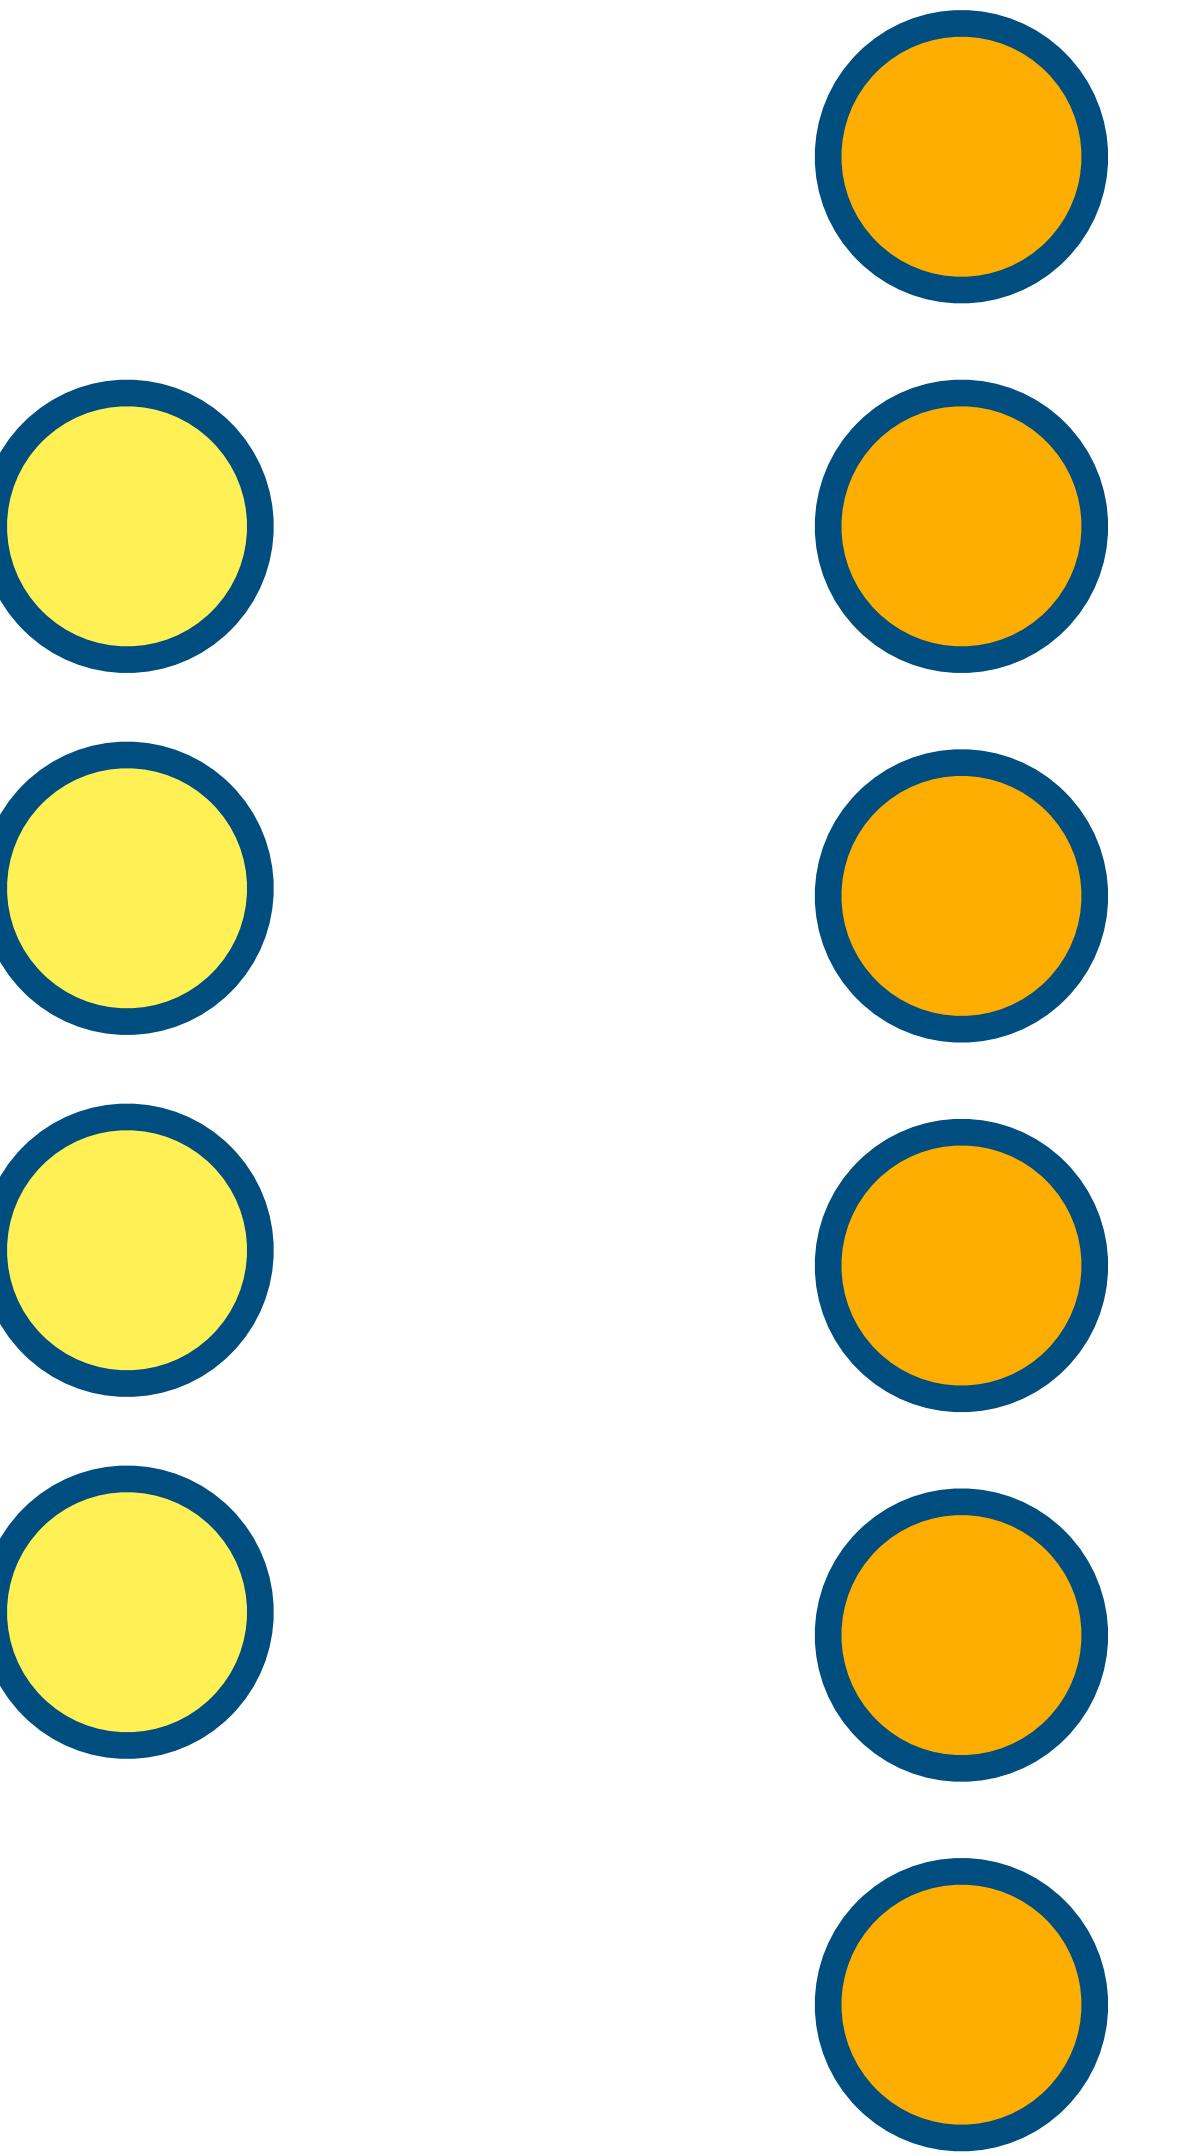
A tensor is a mathematical entity that lives in a structure and interacts with other mathematical entities. If one *transforms* the other entities in the structure in a regular way, then the tensor *must obey a related transformation rule*.

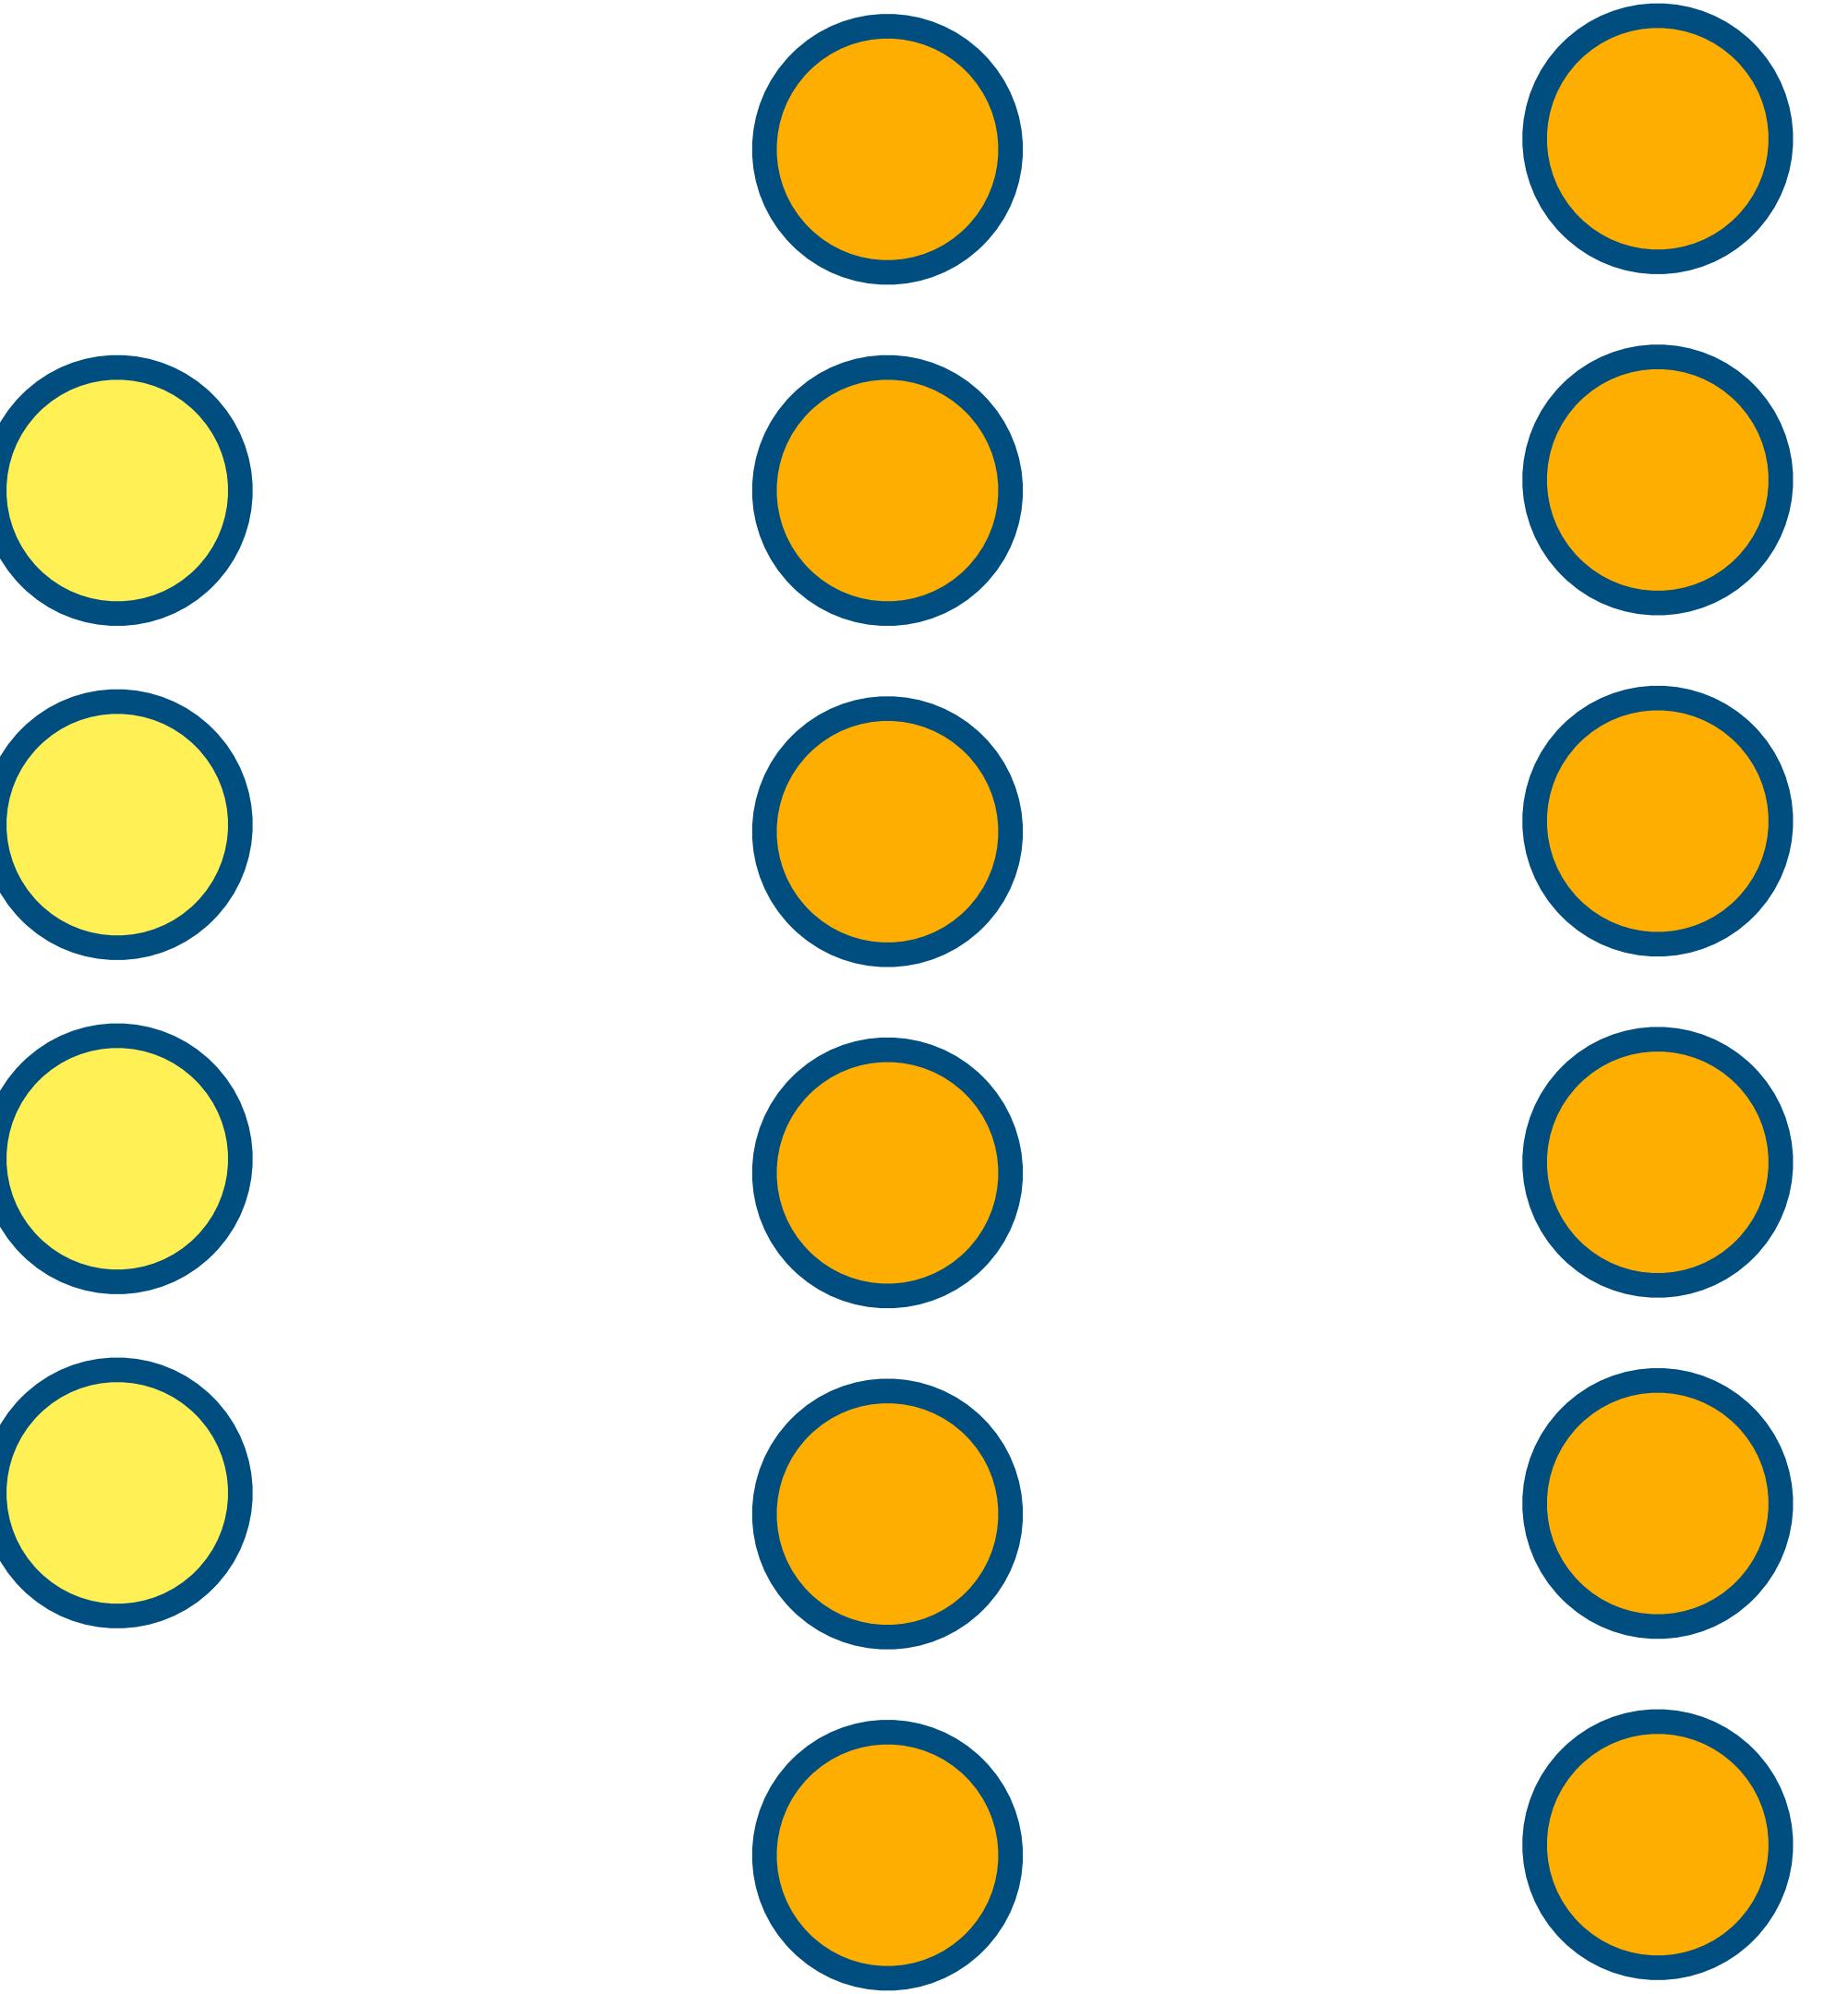
8 9	7 8 9	8 9	1	5	2	3	4 6	6
3	2	1 2				1 2		2
5	6	4 6				4 5	4 5 6	5 6
7 8 9	8 9	7 9	7 8 9	8 9	7	7	7	7
1 2	2		3	3	3	1 2		
5	5	4	6	6	6	5	8 9	
7		7	7	7	7			
2	2 3		2 3	2 3		2		
4	4	5	6	6	7	4	9 1	
8	8					8		
6	2	7	8	1 2	1	2		
4	9			9	4 5	4 5	4 5	3
2		2	2 3	2 3		3		
4		4 5 6	4 5 6	6	4 5 6	4 5	4 5 6	5 6
8 9	8 9	9	9	9	9	7 8	7	7 8
2		2	2 3		3			
5		5	5	4	5	9	9 1	5
8		8	7	8	8	7 8	7	7 8
1				1	1			
5	5	7	5		5			
8 9	8 9	7 9	7 8 9	8 9	8 9	6	6	4
7	2	1 2	2 3	1 2 3	1 3		3	
4 5			5 6	6	5 6	5	5	5
8 9	8 9	9	8 9	8 9	8 9	8	8	8

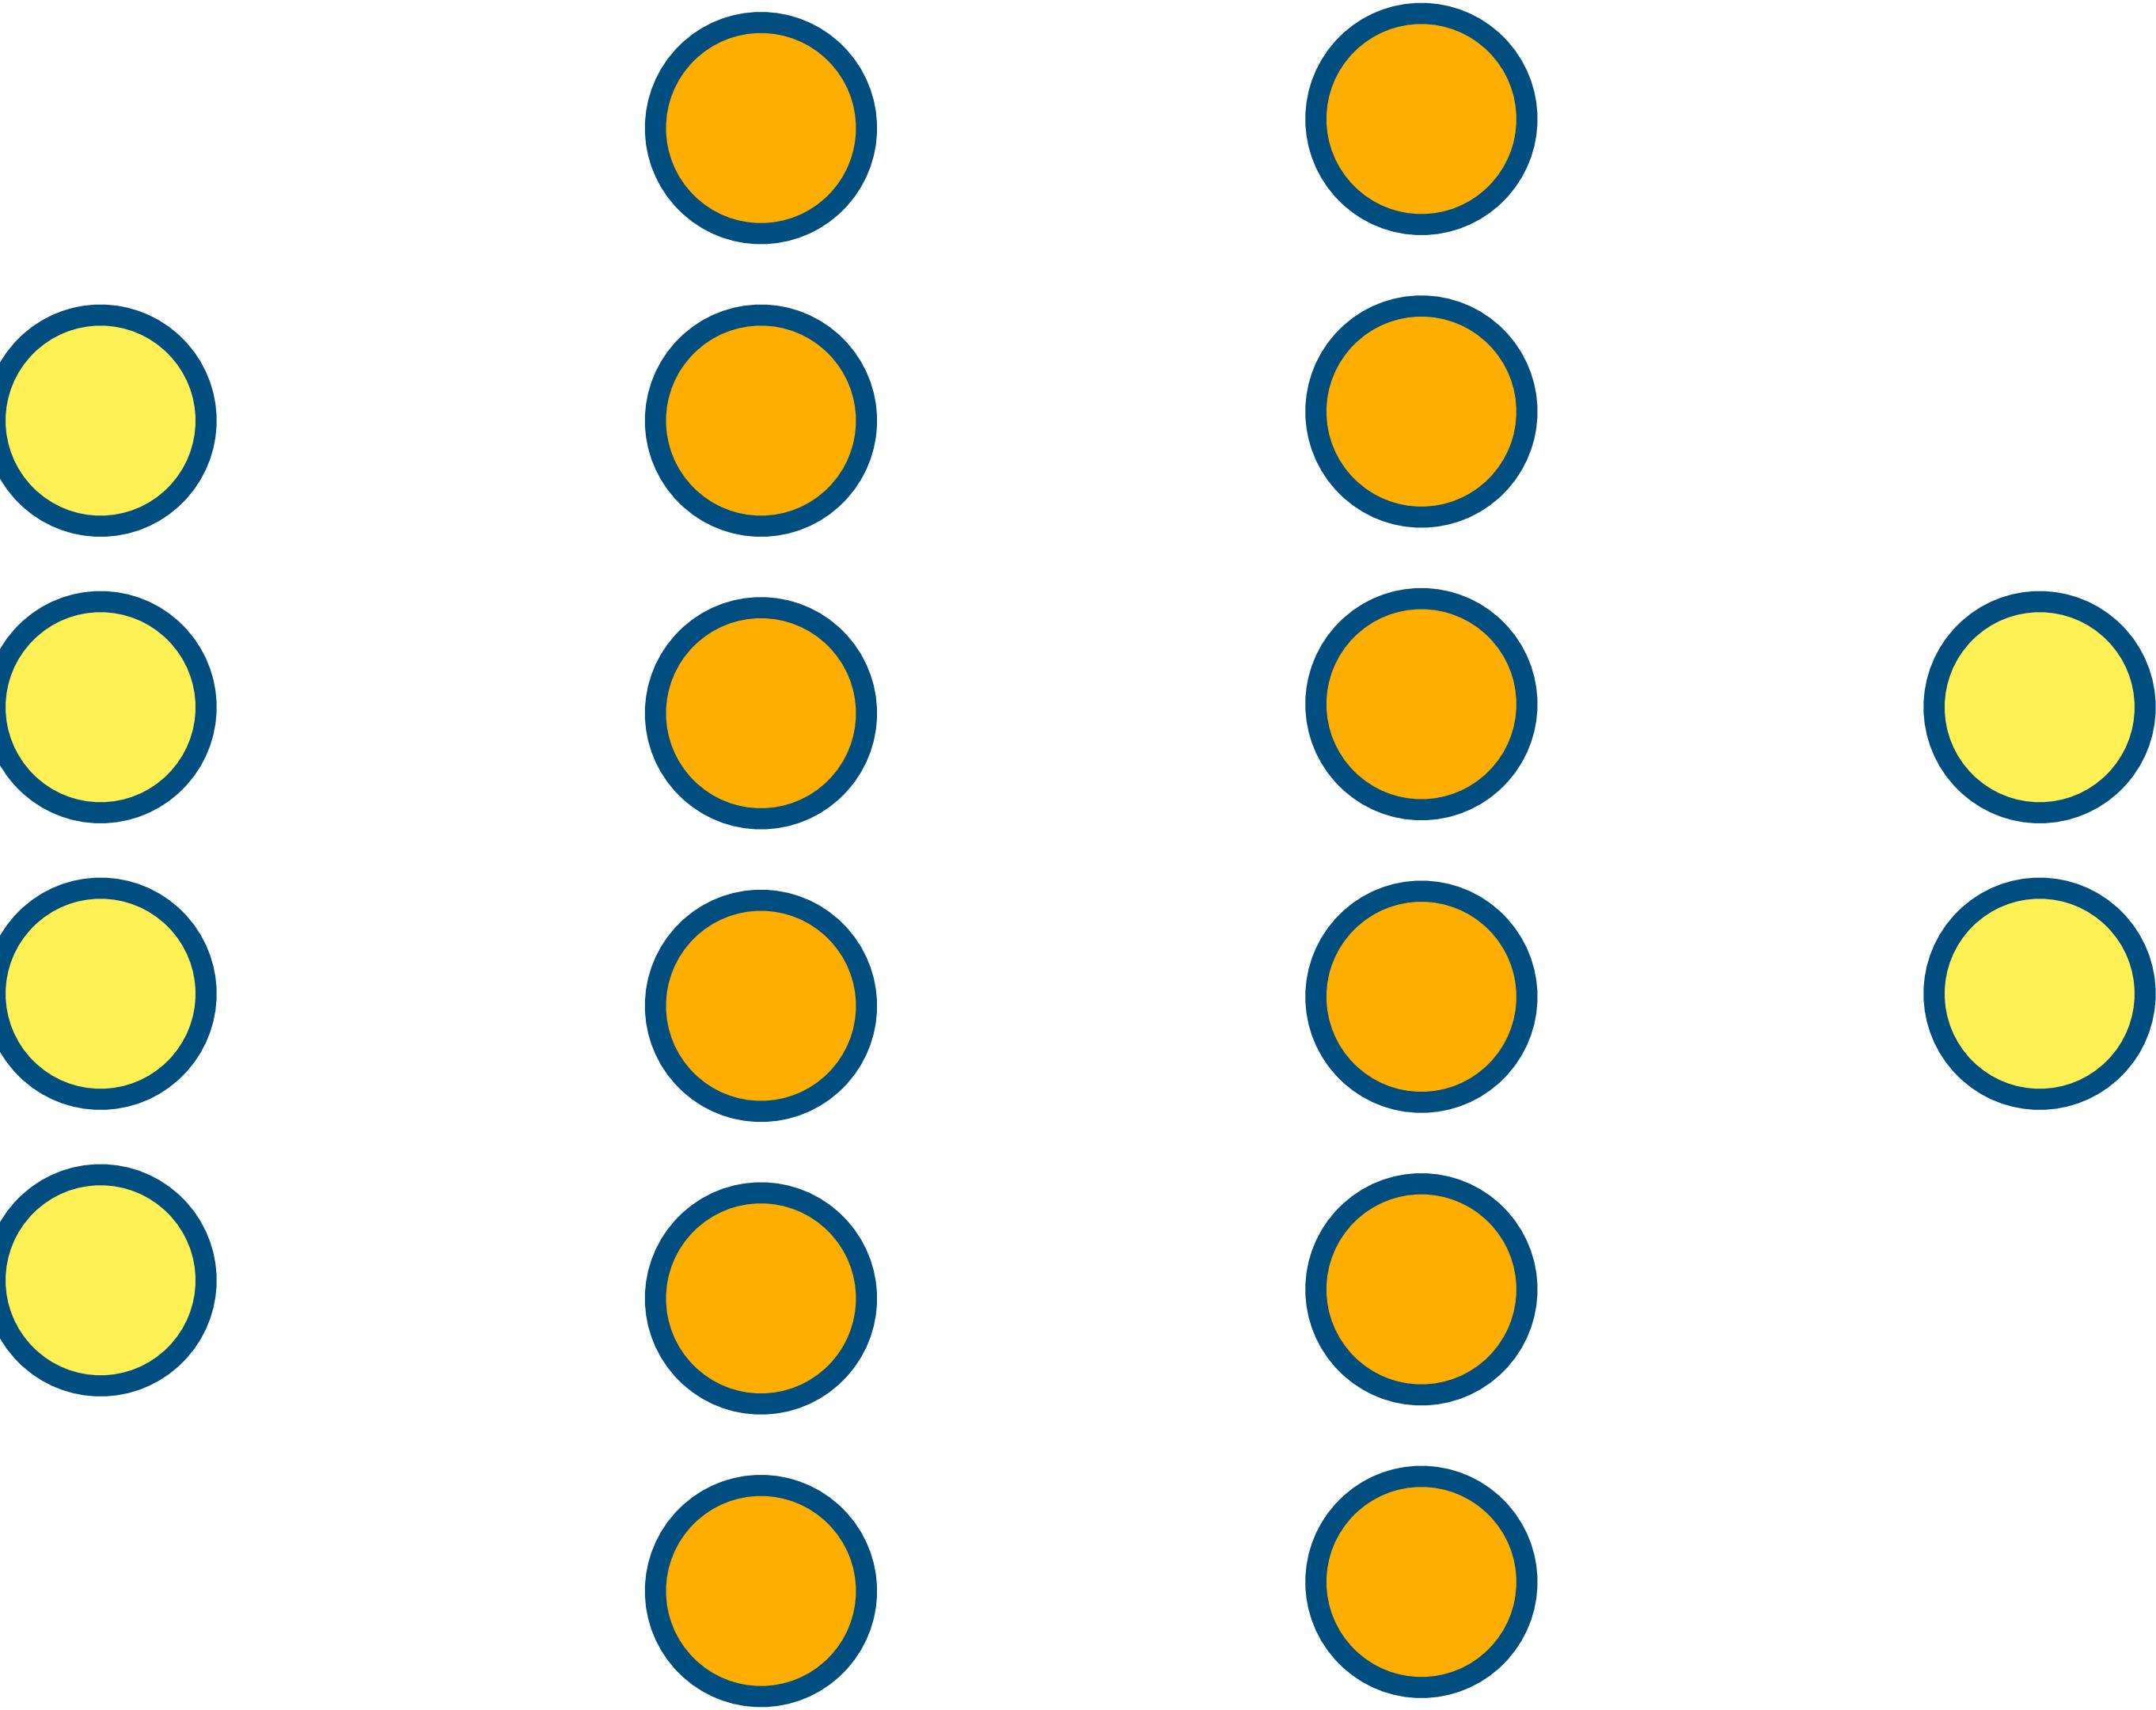
A 2-D Tensor?

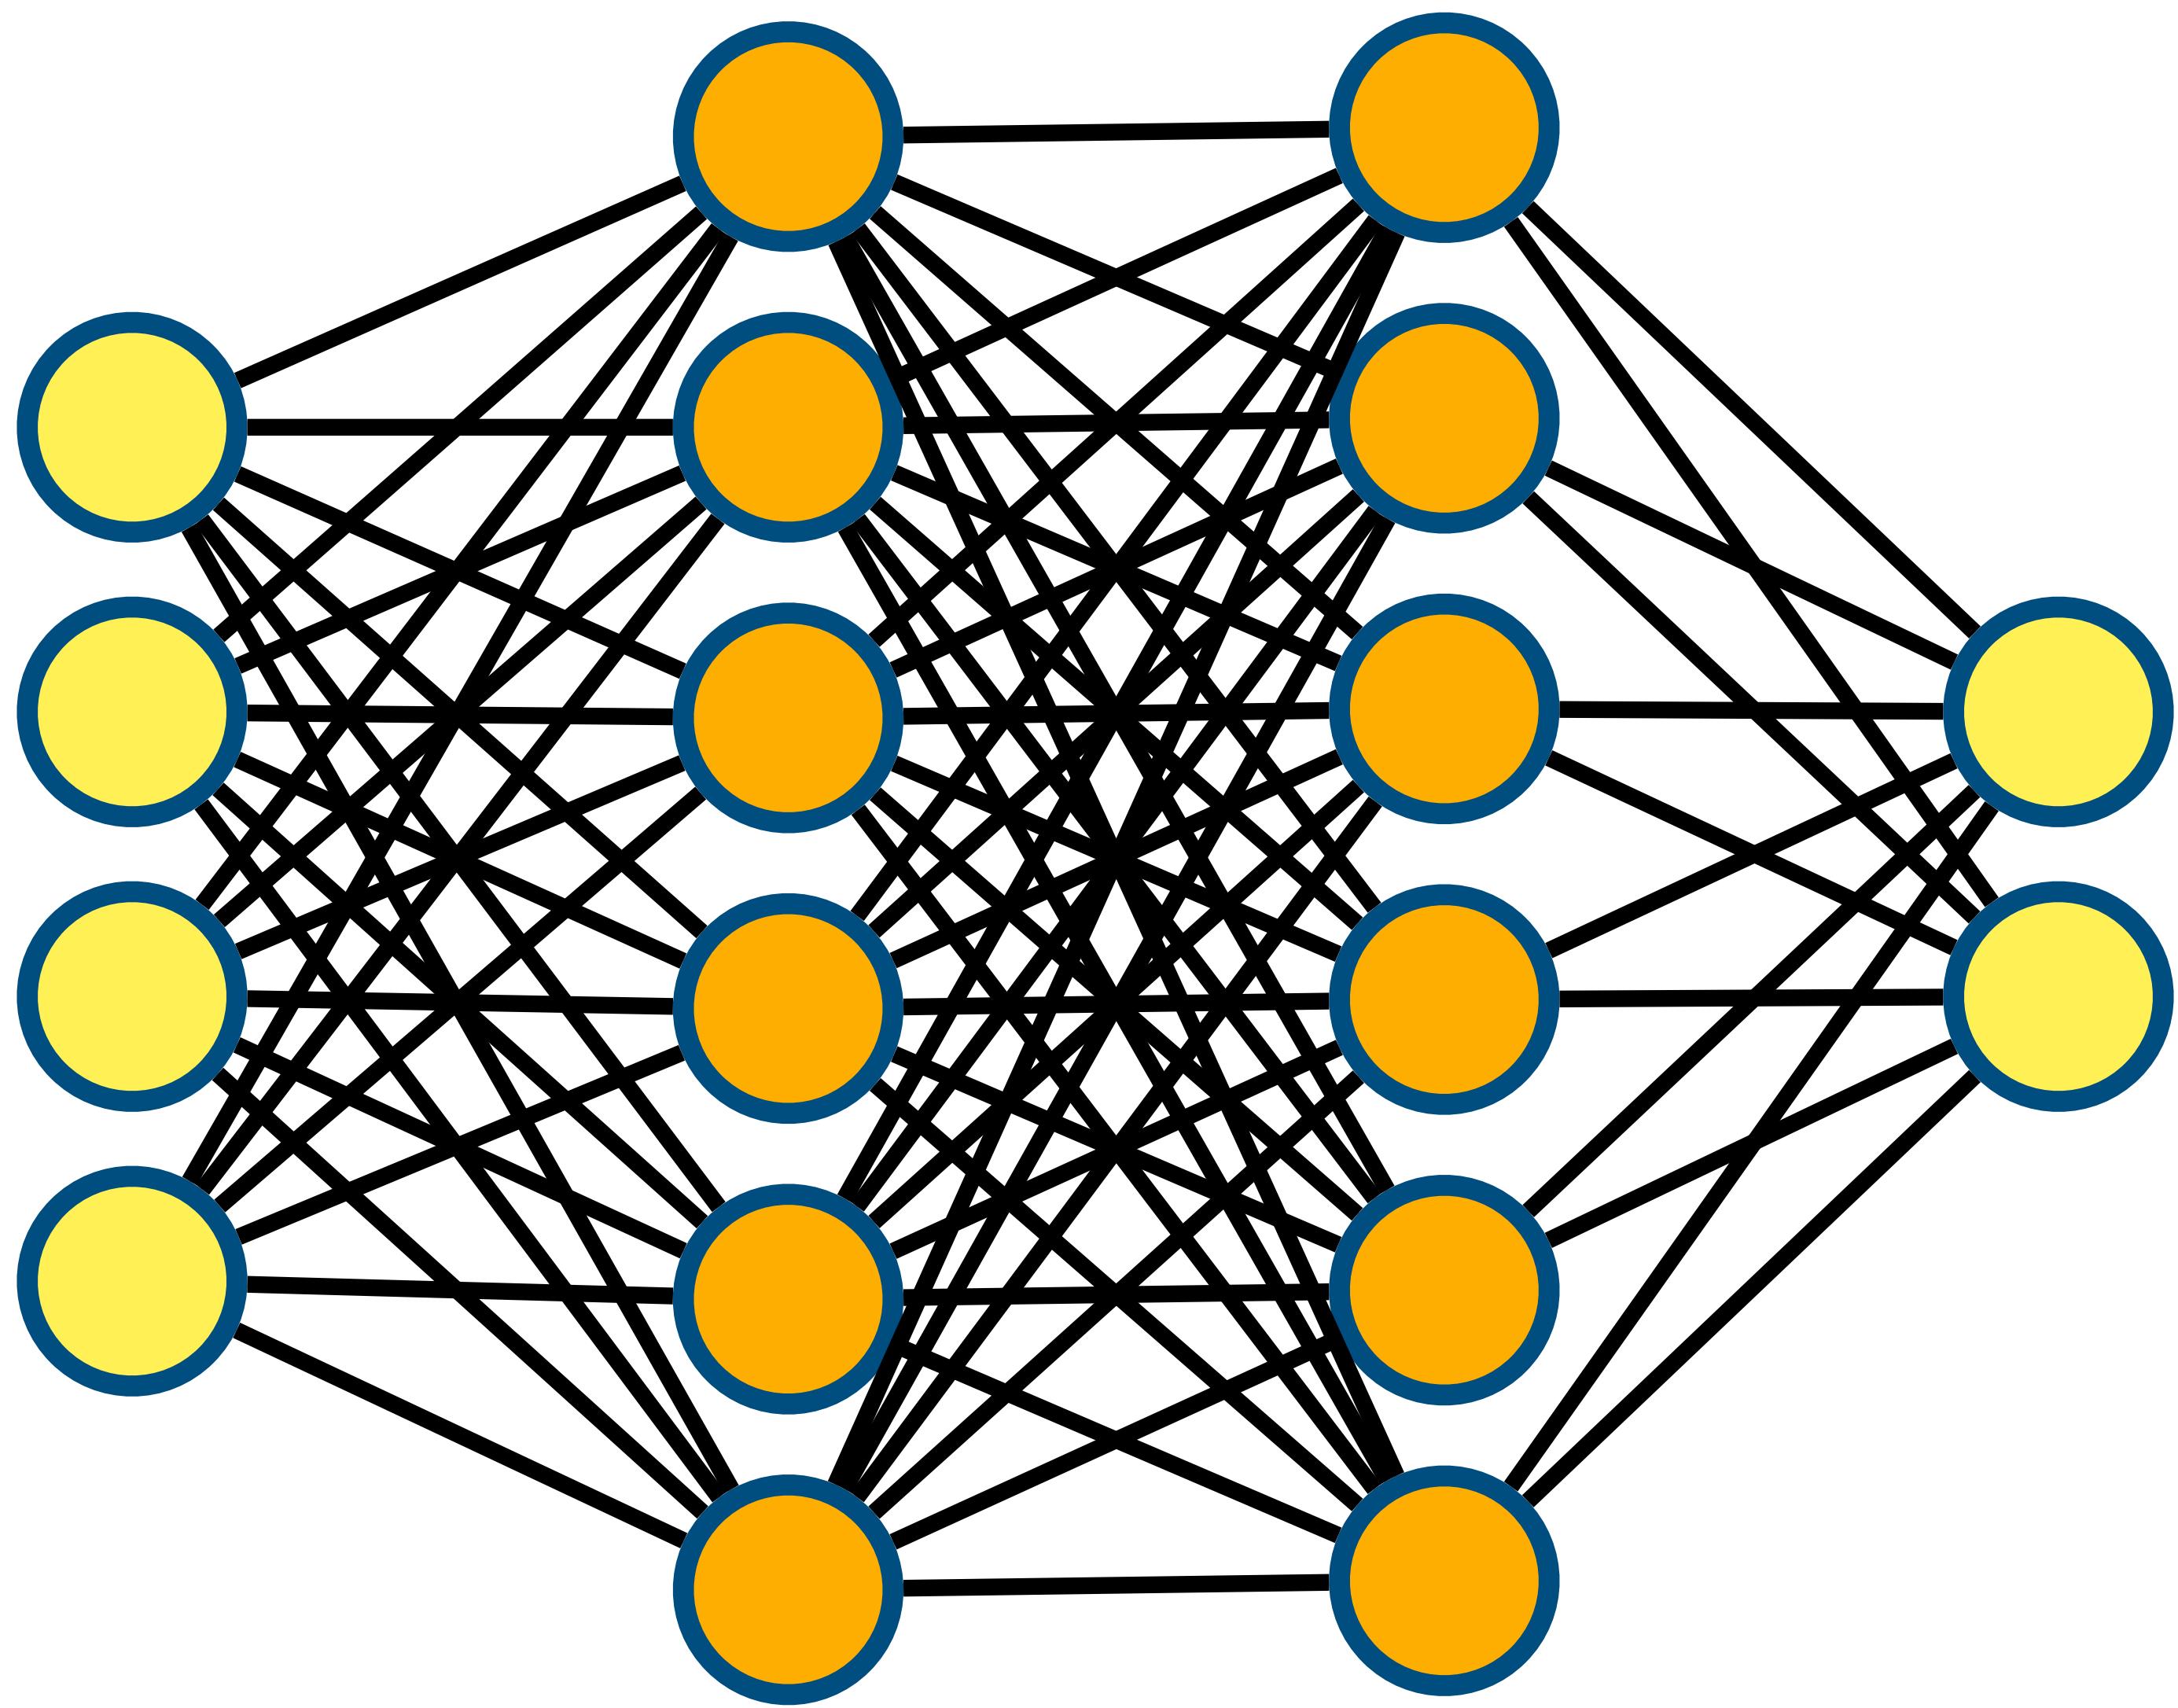
# Neural Networks





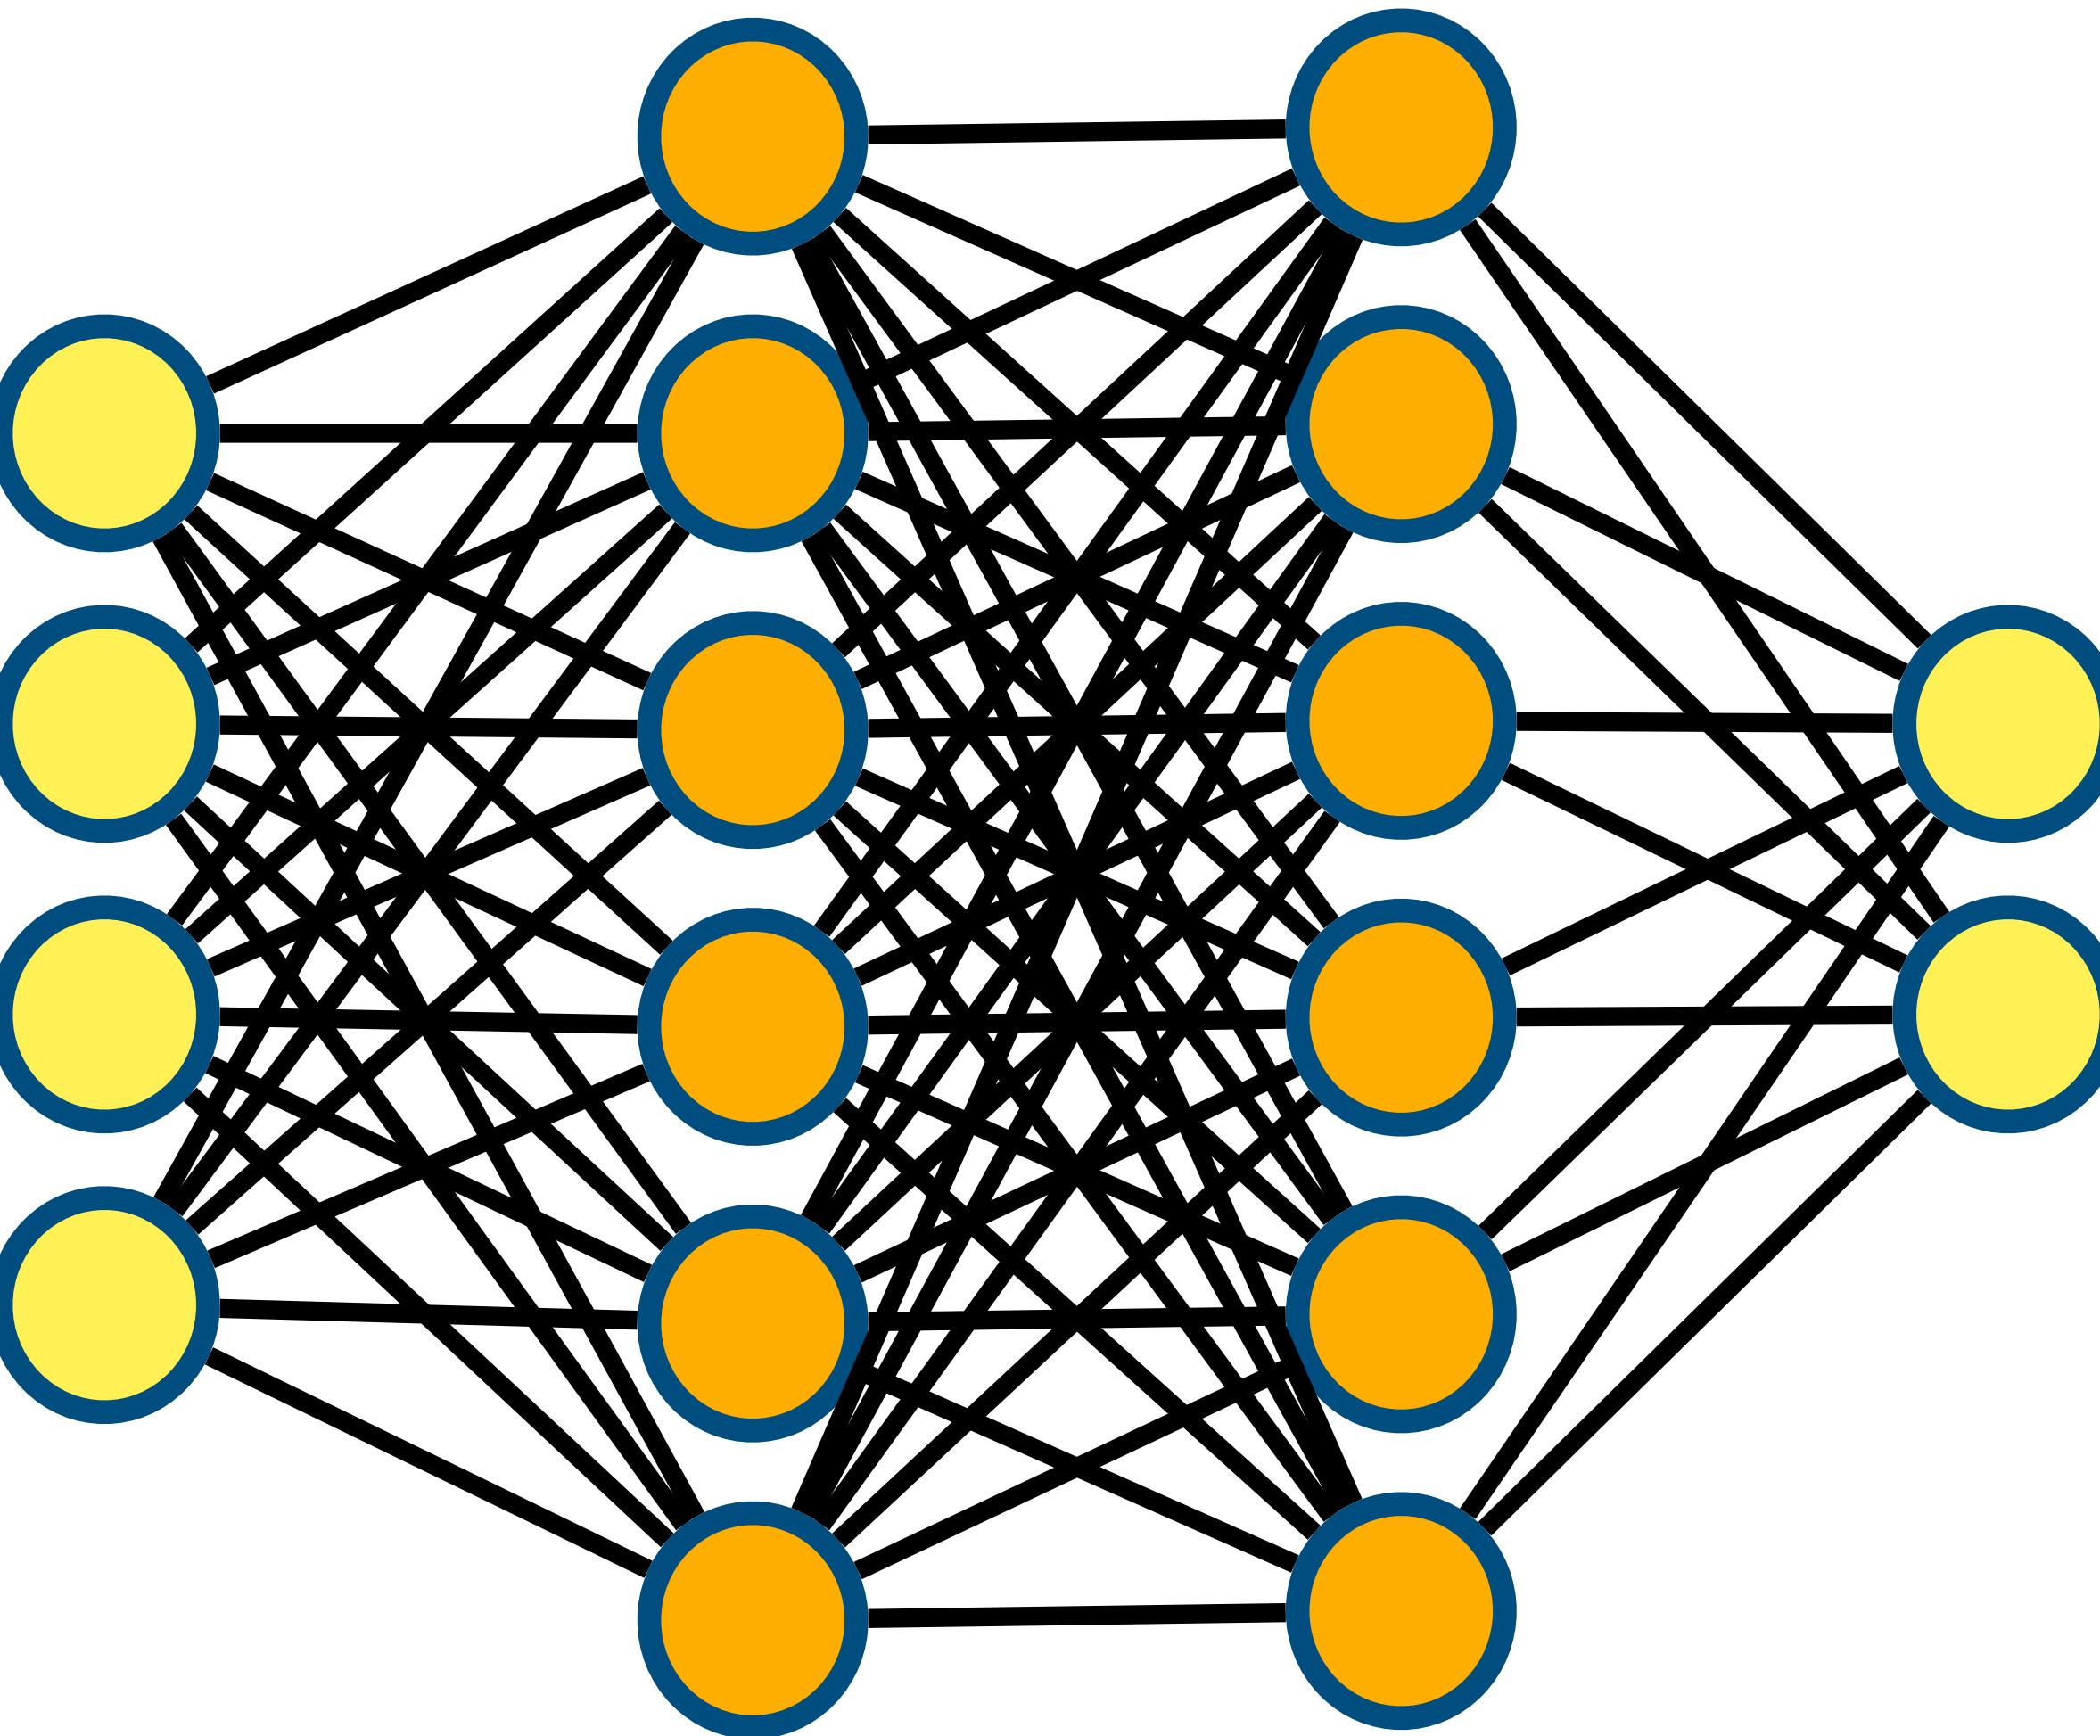






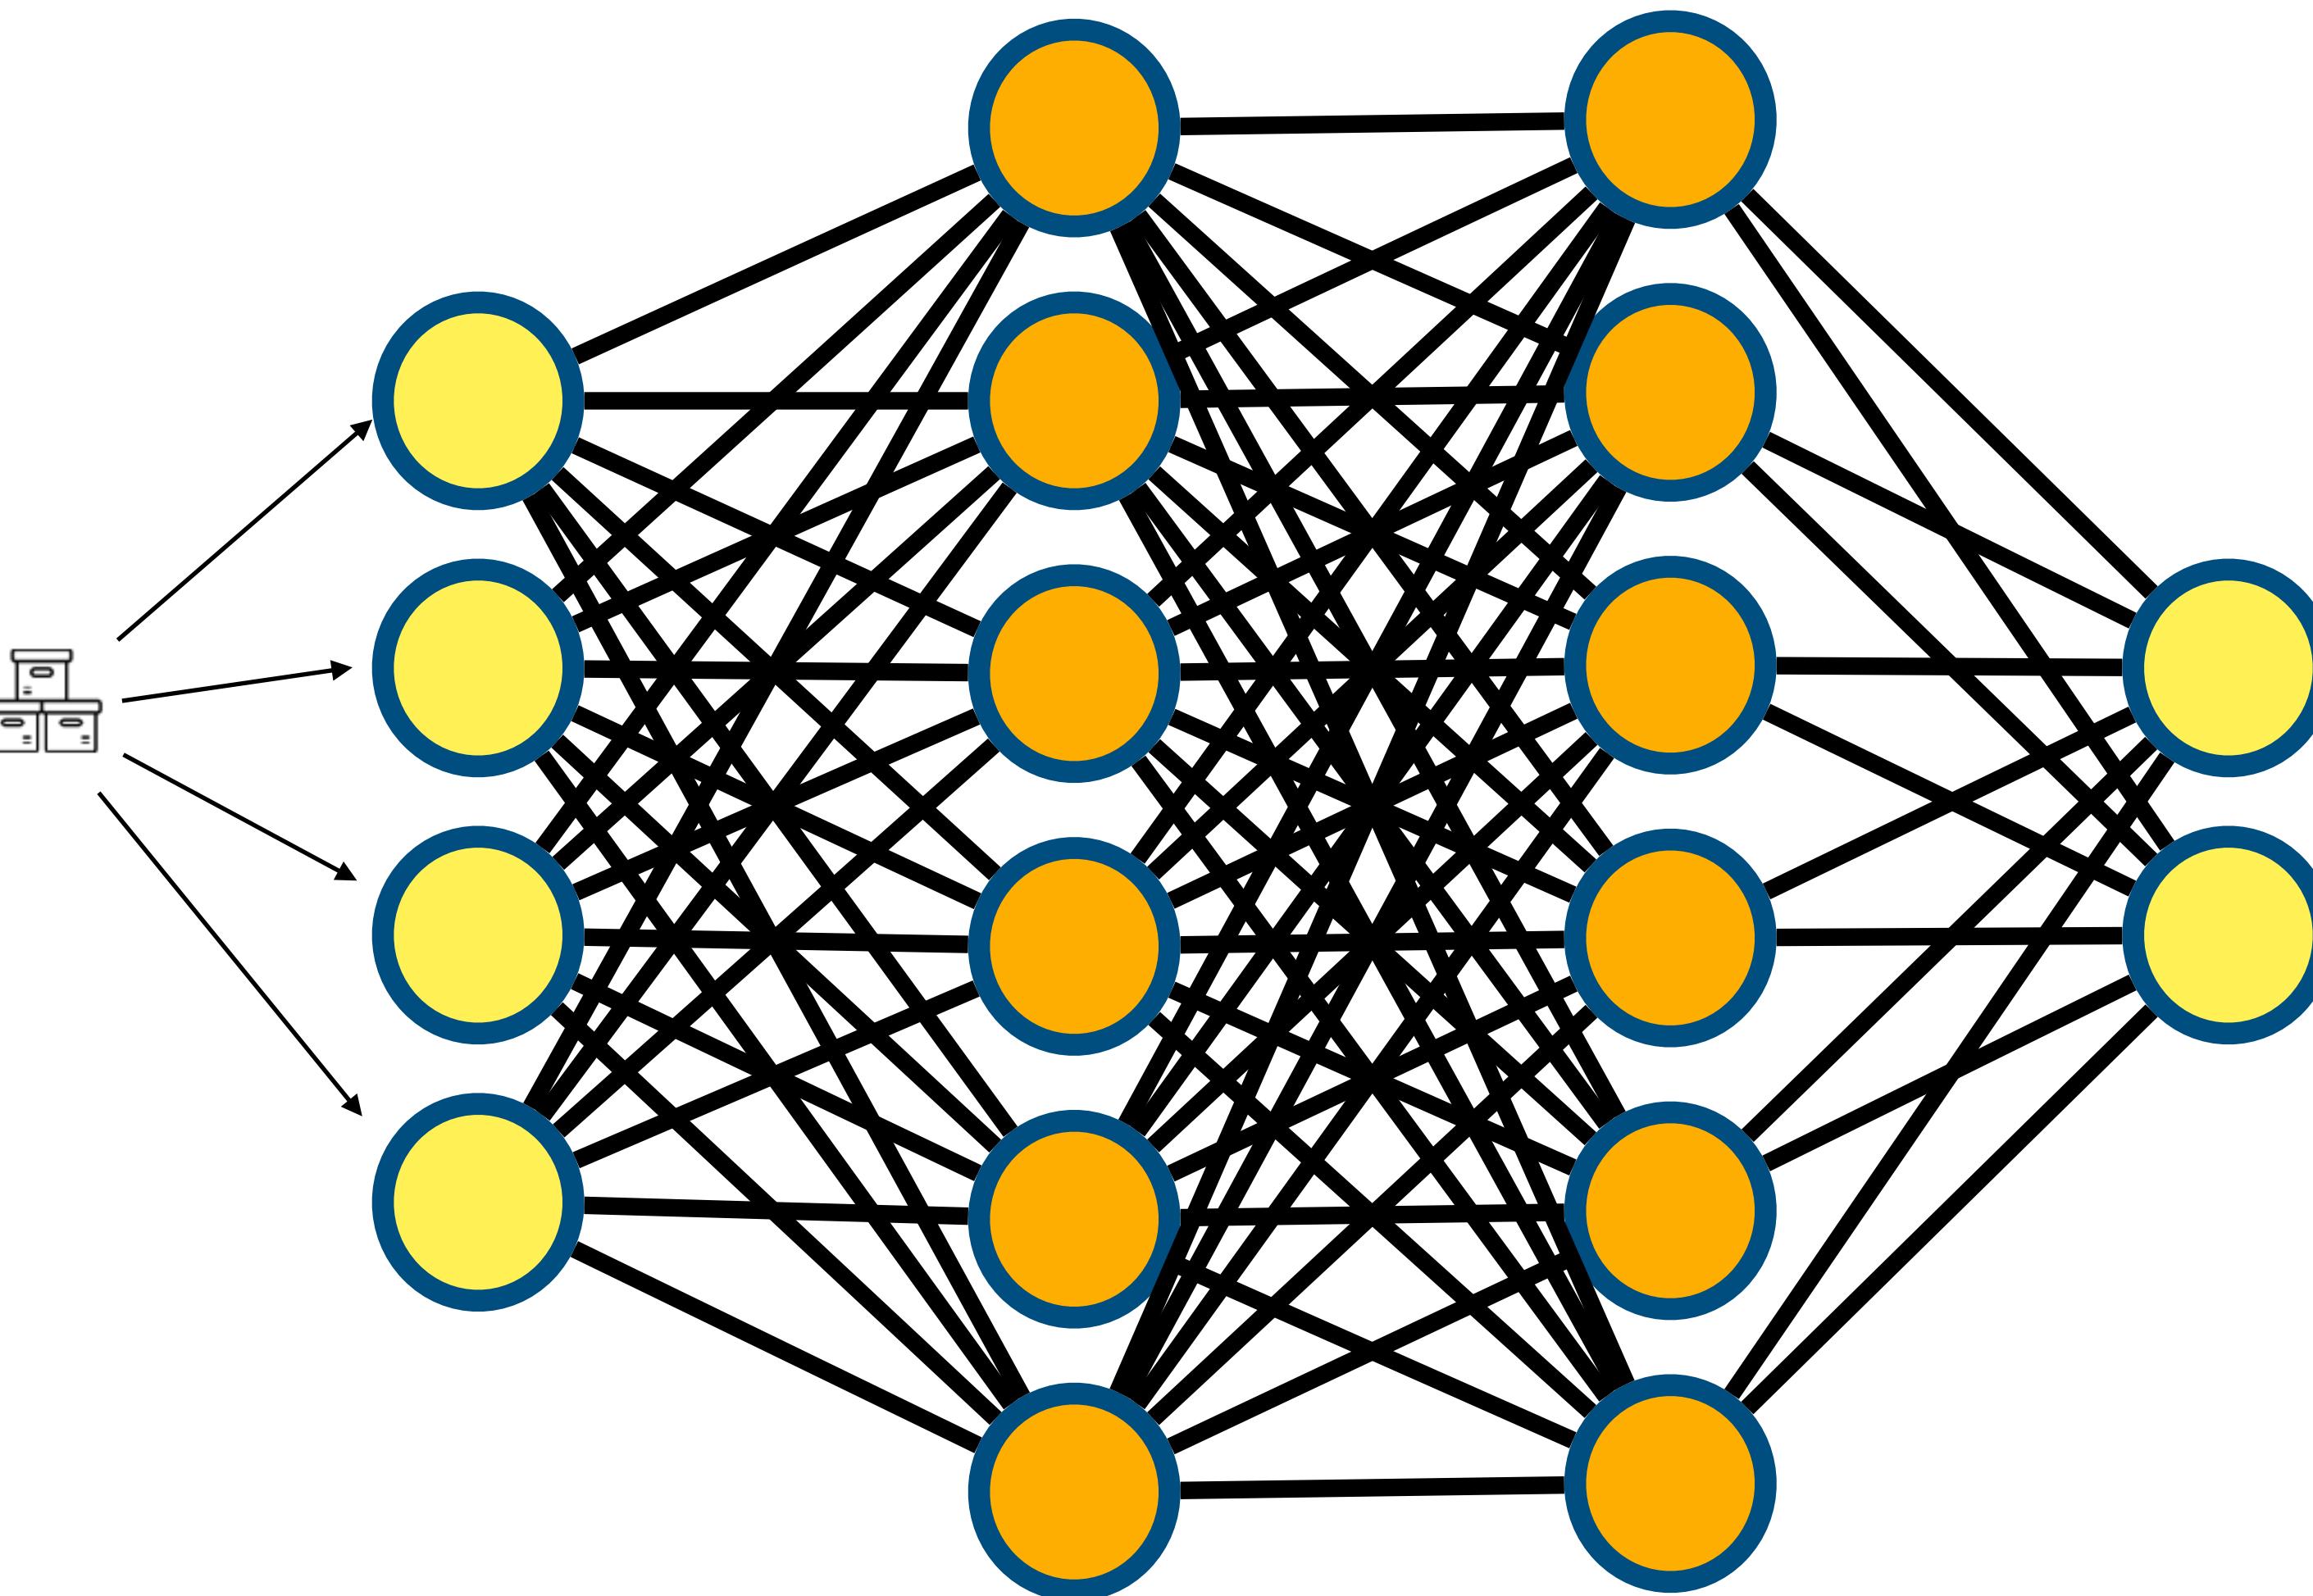
## Table of Data

Temperature (°C)	Humidity (%)	Dew Point (°C)	Cloud Cover (%)	Rain (Label: Y)
25	80	20	90	1(Rain)
30	50	15	40	0(No Rain)
22	85	18	95	1(Rain)
28	60	17	60	0(No Rain)
26	78	19	85	1(Rain)
31	72	16	35	0(No Rain)
24	90	21	98	1(Rain)
29	65	18	70	0(No Rain)
23	88	20	92	1(Rain)
27	58	17	50	0(No Rain)
30	80	20	90	1(Rain)
22	50	15	40	0(No Rain)
28	82	19	95	1(Rain)
26	60	17	60	0(No Rain)
31	78	19	85	1(Rain)
24	55	16	35	0(No Rain)
29	90	21	98	1(Rain)
23	72	18	70	0(No Rain)
27	88	20	92	1(Rain)
30	58	17	50	0(No Rain)
22	80	20	90	1(Rain)
28	50	15	40	0(No Rain)
26	72	18	95	1(Rain)
31	60	17	60	0(No Rain)
24	78	19	85	1(Rain)
29	55	16	35	0(No Rain)
23	90	21	98	1(Rain)
27	65	18	70	0(No Rain)
30	88	20	92	1(Rain)
22	58	17	50	0(No Rain)
28	80	20	90	1(Rain)
26	50	15	40	0(No Rain)
31	82	18	95	1(Rain)
24	60	17	60	0(No Rain)
29	78	19	85	1(Rain)
23	55	16	35	0(No Rain)
27	90	21	98	1(Rain)
30	65	18	70	0(No Rain)
22	88	20	92	1(Rain)
28	28	17	50	0(No Rain)
26	80	20	90	1(Rain)
31	50	15	40	0(No Rain)
24	85	18	95	1(Rain)
29	60	17	60	0(No Rain)
23	78	19	85	1(Rain)
27	55	16	35	0(No Rain)
30	90	21	98	1(Rain)
22	65	18	70	0(No Rain)
28	88	20	92	1(Rain)
26	78	17	50	0(No Rain)
31	80	20	90	1(Rain)
24	52	15	40	0(No Rain)
29	85	18	95	1(Rain)
23	60	17	60	0(No Rain)
27	78	19	85	1(Rain)
23	55	16	35	0(No Rain)
23	65	18	70	0(No Rain)
23	88	20	92	1(Rain)



## Table of Data

Temperature (°C)	Humidity (%)	Dew Point (°C)	Cloud Cover (%)	Rain (Label: Y)
25	80	20	90	1(Rain)
30	50	15	40	0(No Rain)
22	85	18	95	1(Rain)
28	60	17	60	0(No Rain)
26	78	19	85	1(Rain)
31	72	16	35	1(Rain)
24	90	21	98	1(Rain)
29	65	18	70	0(No Rain)
23	88	20	92	1(Rain)
27	58	17	50	0(No Rain)
30	80	20	90	1(Rain)
22	50	15	40	0(No Rain)
28	52	19	95	1(Rain)
26	60	17	60	0(No Rain)
31	78	19	85	1(Rain)
24	55	16	35	0(No Rain)
29	90	21	98	1(Rain)
23	72	18	70	1(Rain)
27	88	20	92	1(Rain)
30	58	17	50	0(No Rain)
22	80	20	90	1(Rain)
28	50	15	40	0(No Rain)
26	72	18	95	1(Rain)
31	60	17	60	0(No Rain)
24	78	19	85	1(Rain)
29	55	16	35	0(No Rain)
23	90	21	98	1(Rain)
27	65	18	70	0(No Rain)
30	88	20	92	1(Rain)
22	58	17	50	0(No Rain)
28	80	20	90	1(Rain)
26	50	15	40	0(No Rain)
31	85	18	95	1(Rain)
24	60	17	60	0(No Rain)
29	72	19	85	1(Rain)
23	55	16	35	0(No Rain)
27	90	21	98	1(Rain)
30	65	18	70	0(No Rain)
22	88	20	92	1(Rain)
28	52	17	50	0(No Rain)
26	80	20	90	1(Rain)
31	50	15	40	0(No Rain)
24	85	18	95	1(Rain)
29	60	17	60	0(No Rain)
23	78	19	85	1(Rain)
27	55	16	35	0(No Rain)
30	90	21	98	1(Rain)
22	65	18	70	0(No Rain)
28	88	20	92	1(Rain)
26	58	17	50	0(No Rain)
31	80	20	90	1(Rain)
24	52	15	40	0(No Rain)
29	85	18	95	1(Rain)
23	60	17	60	0(No Rain)
27	78	19	85	1(Rain)
23	55	16	35	0(No Rain)
23	72	21	98	1(Rain)
23	65	18	70	0(No Rain)
23	88	20	92	1(Rain)





### Example Batch of Data (10 Observations)

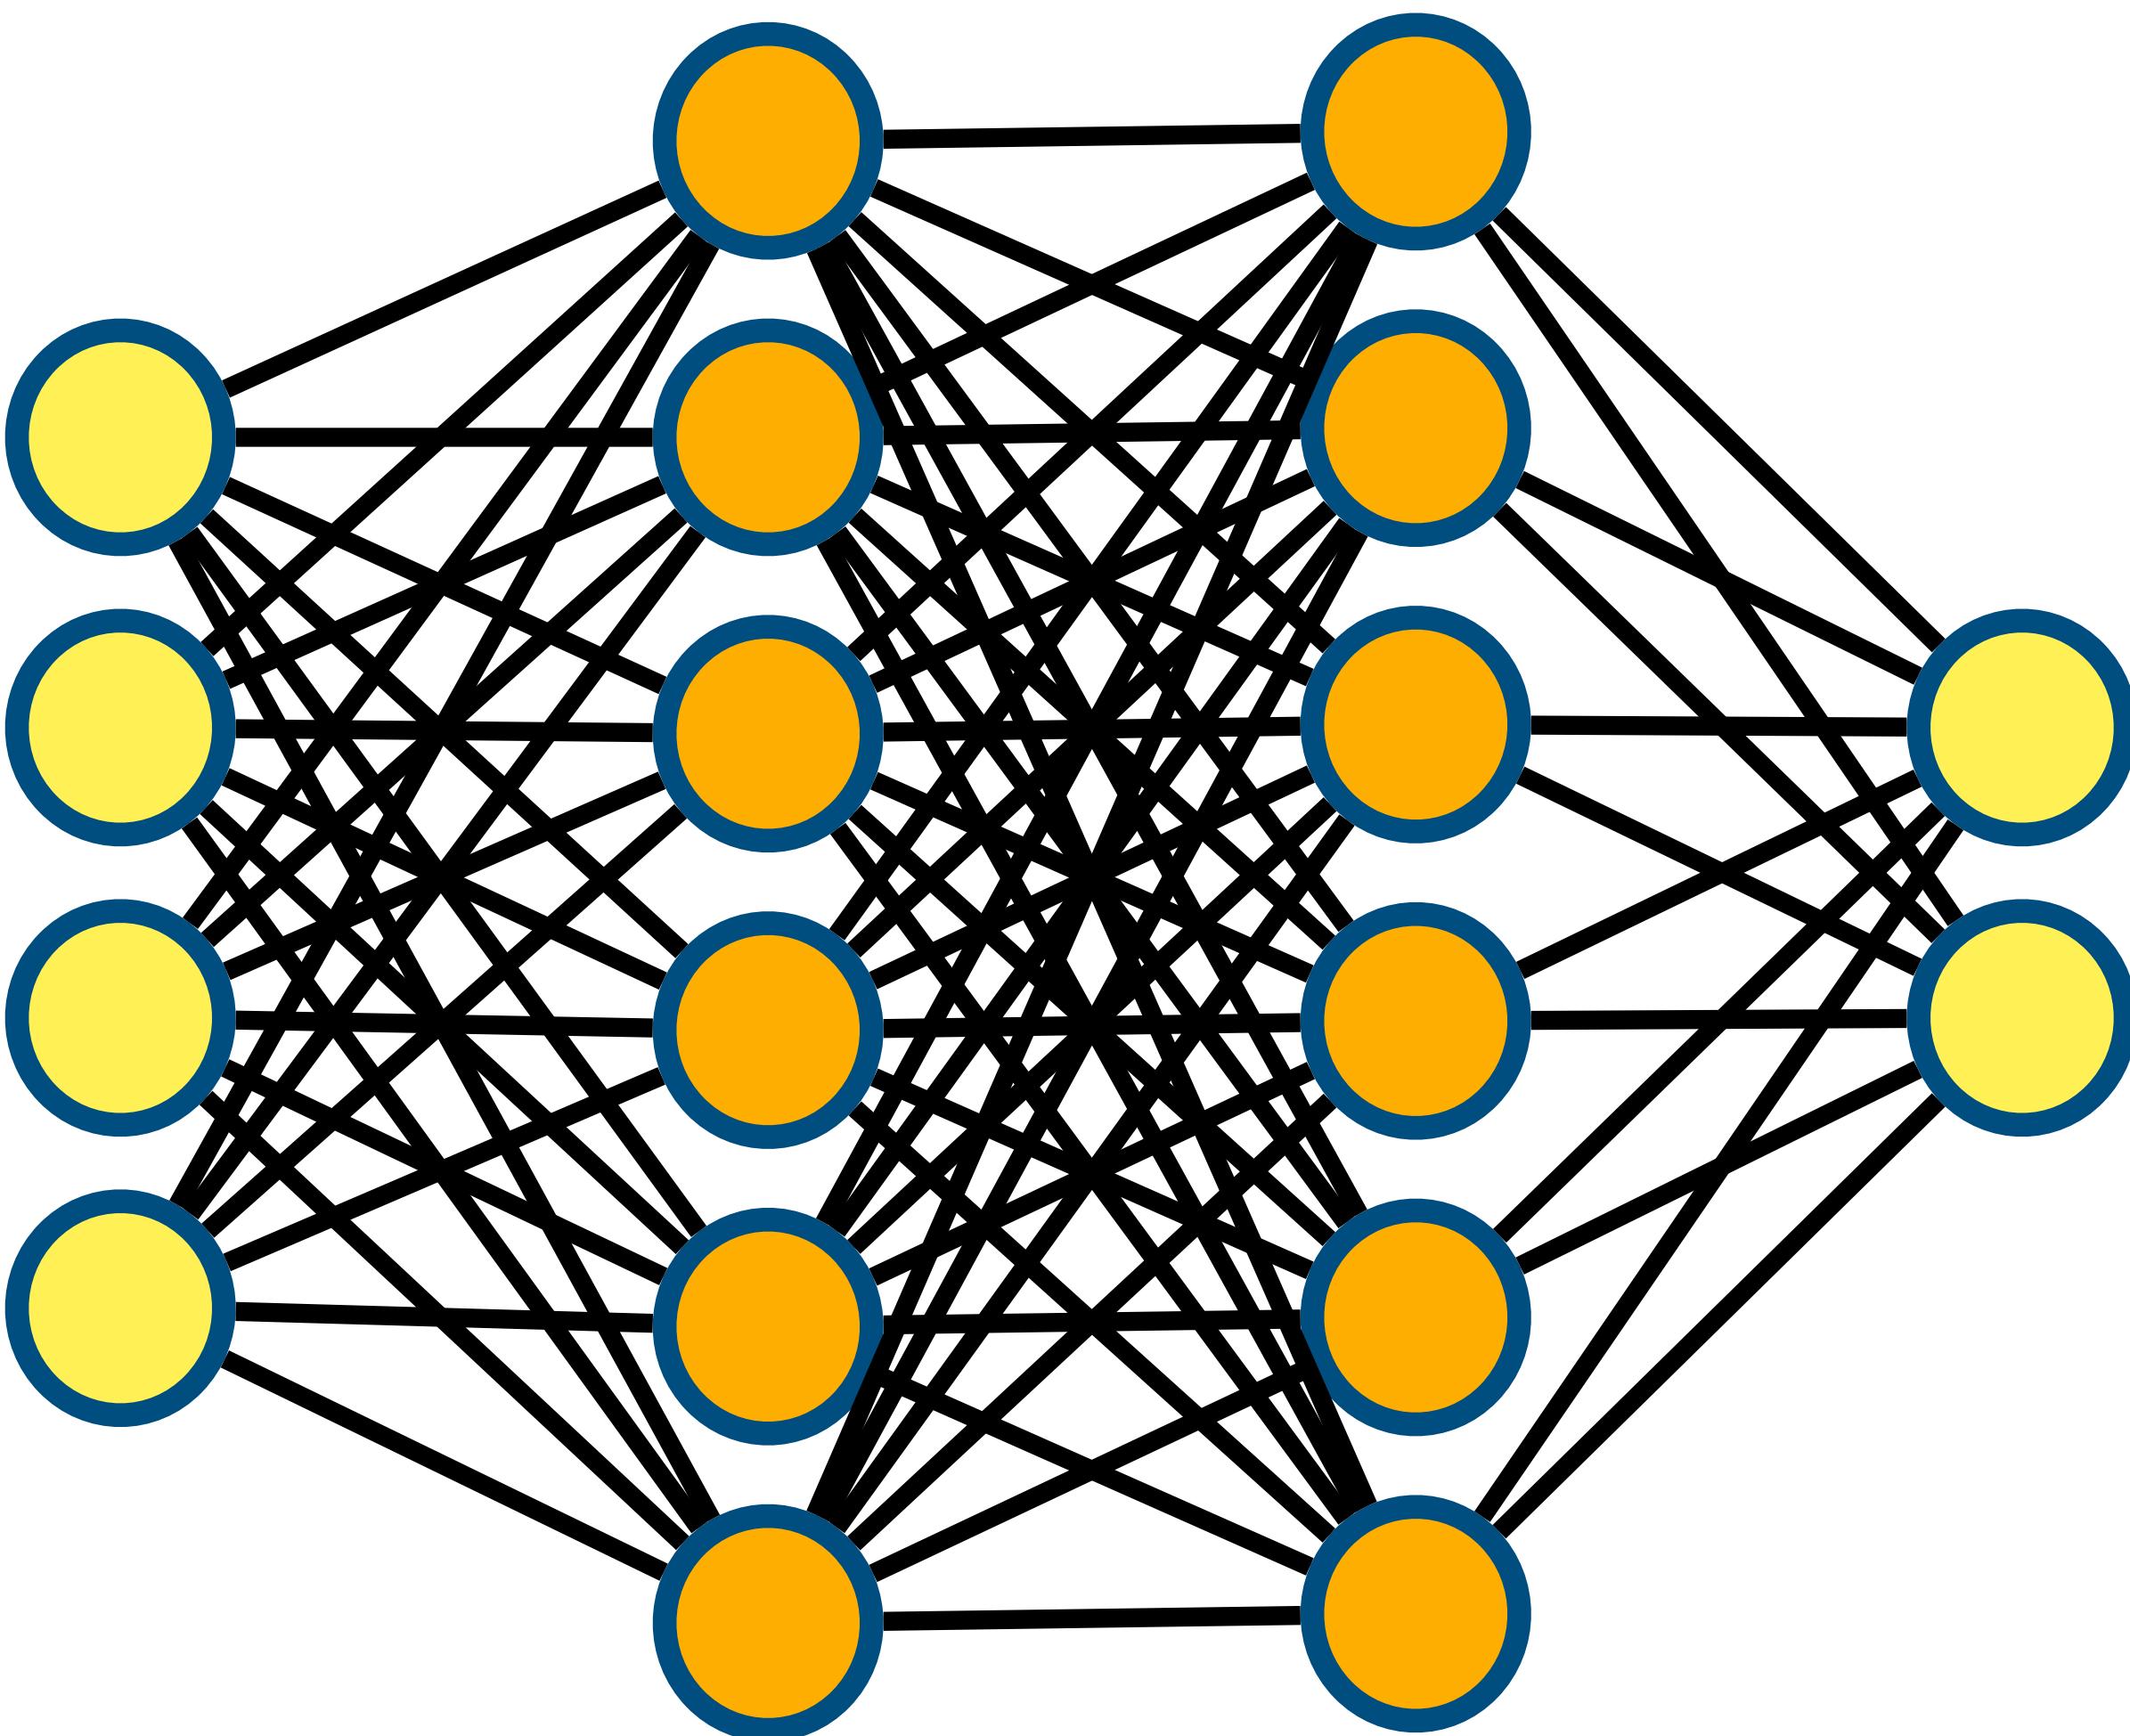
Temperature (°C)	Humidity (%)	Dew Point	Cloud Cover	Rain (Label: Y)
25	80	20	90	1 (Rain)
30	50	15	40	0 (No Rain)
22	85	18	95	1 (Rain)
28	60	17	60	0 (No Rain)
26	78	19	85	1 (Rain)
31	55	16	35	0 (No Rain)
24	90	21	98	1 (Rain)
29	65	18	70	0 (No Rain)
23	88	20	92	1 (Rain)
27	58	17	50	0 (No Rain)

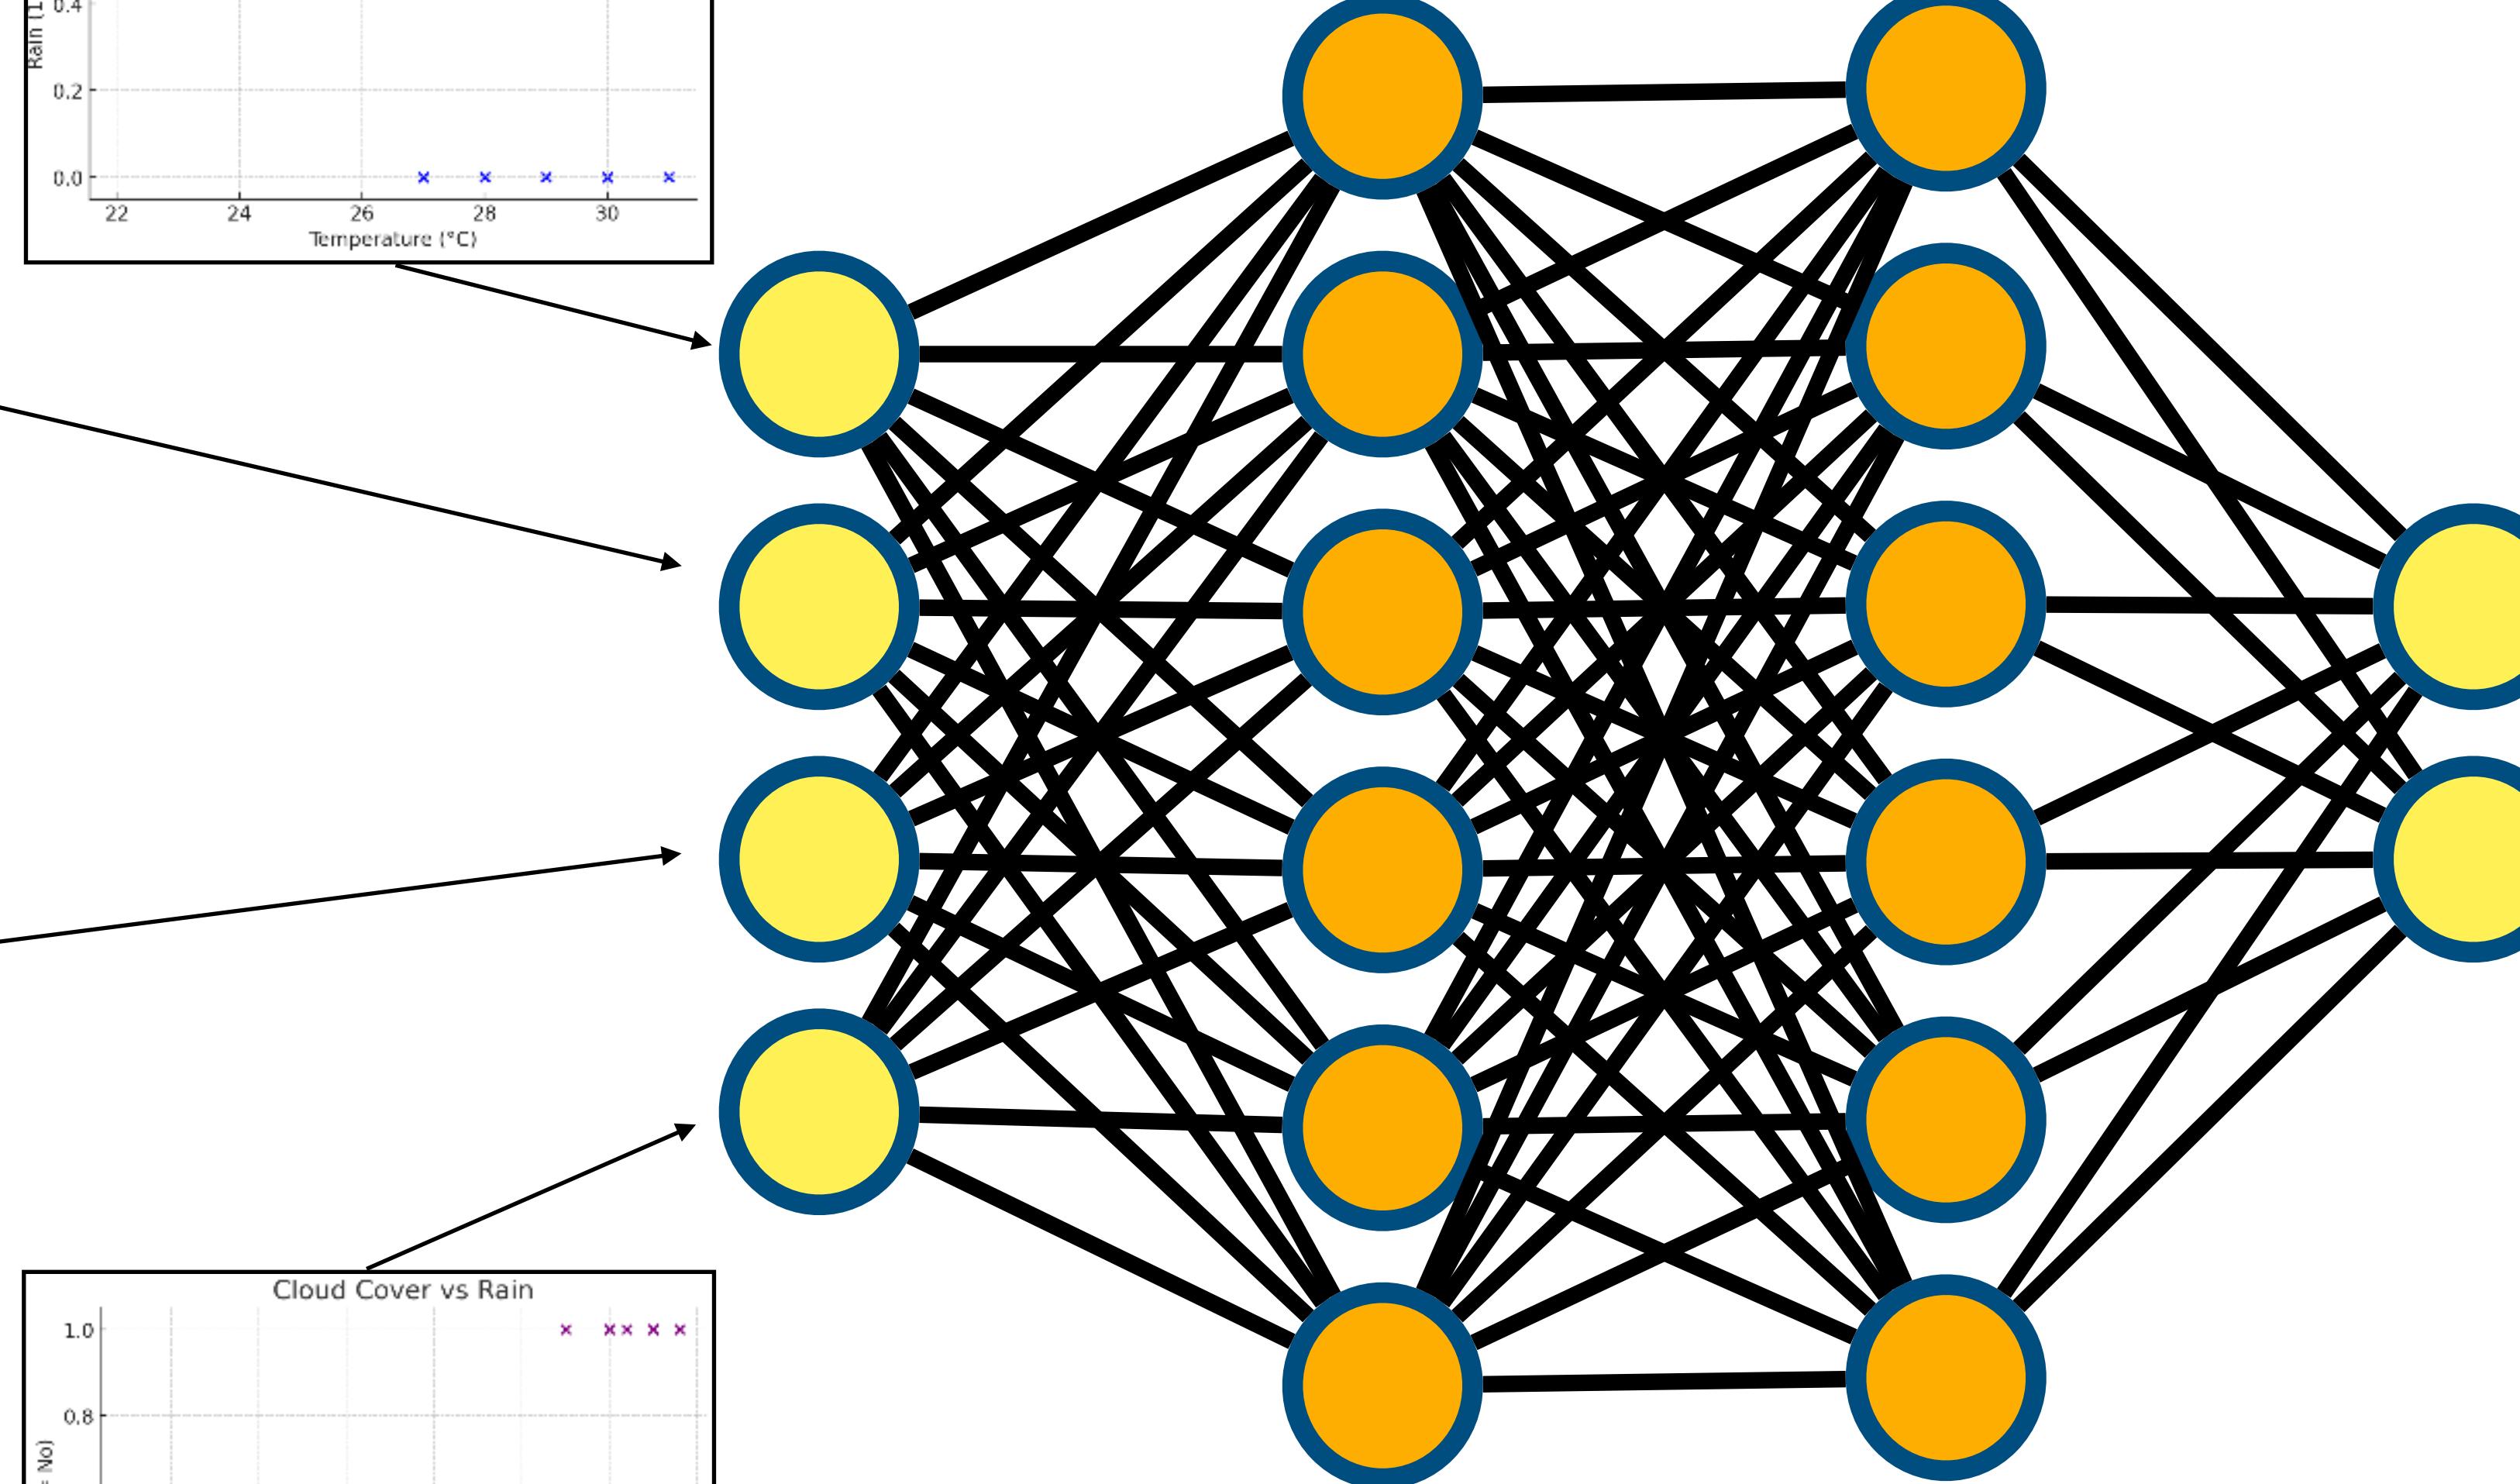
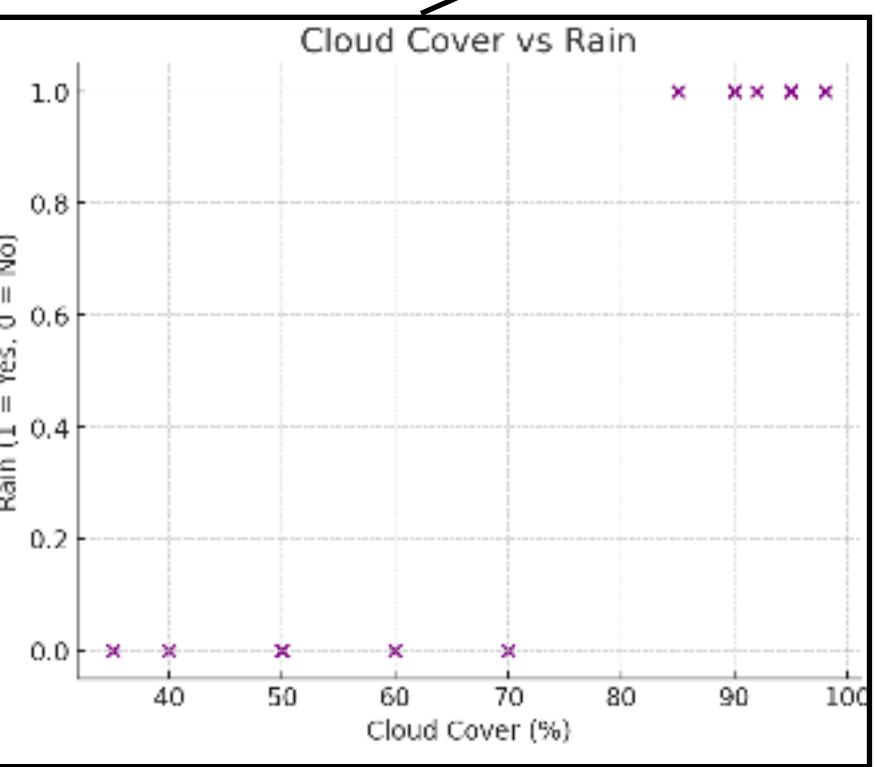
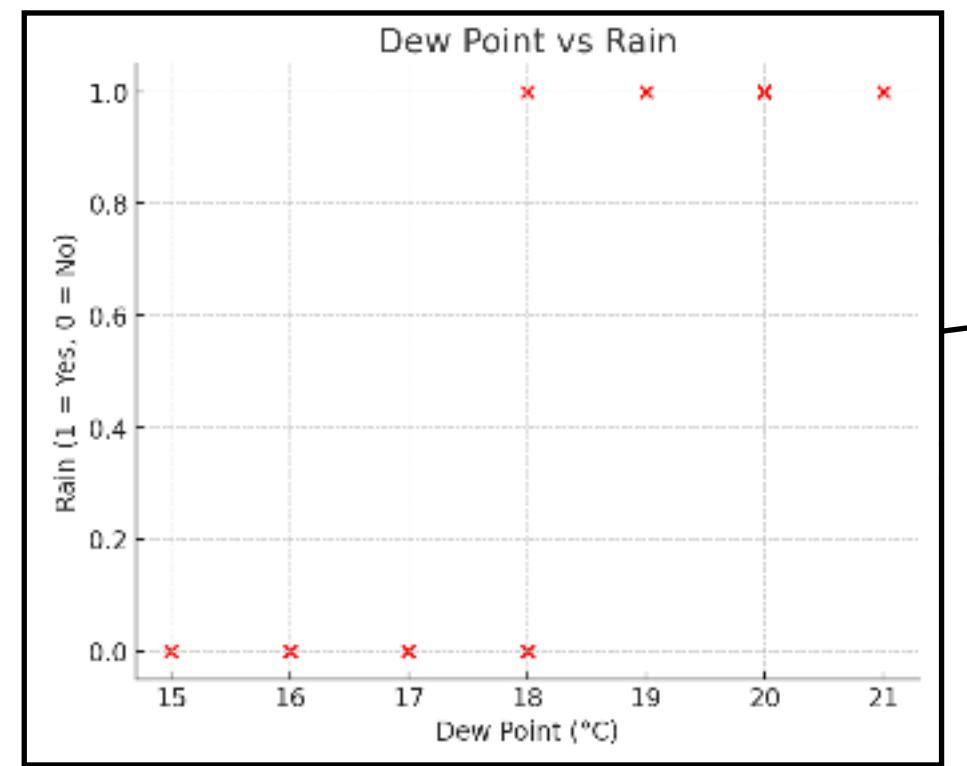
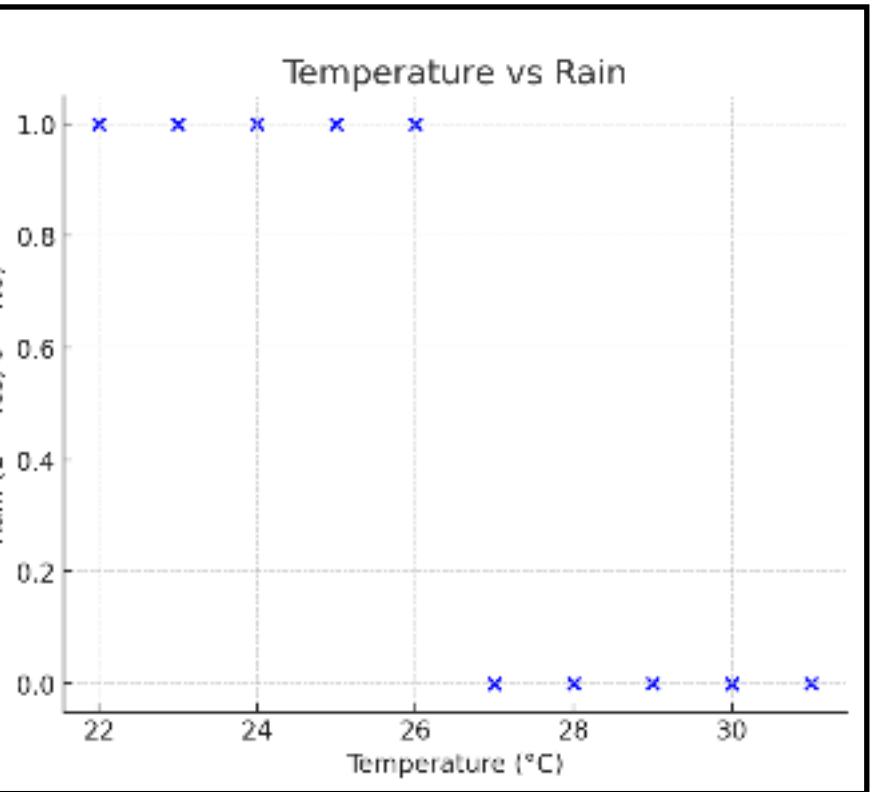
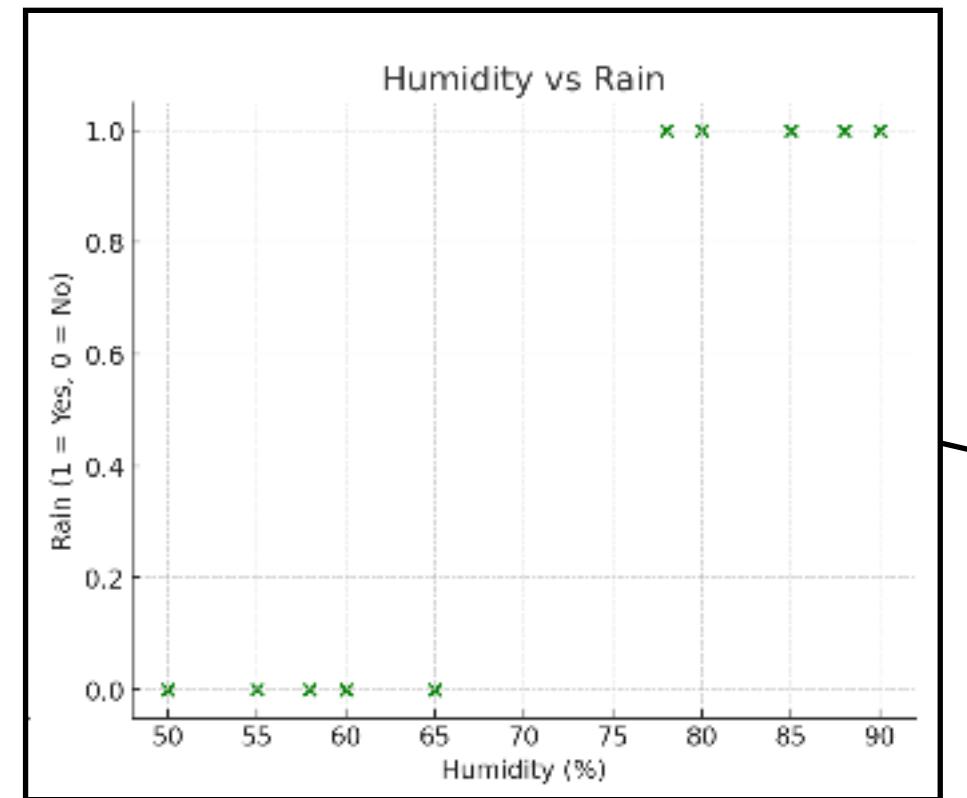
Temp: [25, 30, 22, 28, 26, 31, 24, 29, 23, 27]

Humidity [80, 50, 85, 60, 78, 55, 90, 65, 88, 58]

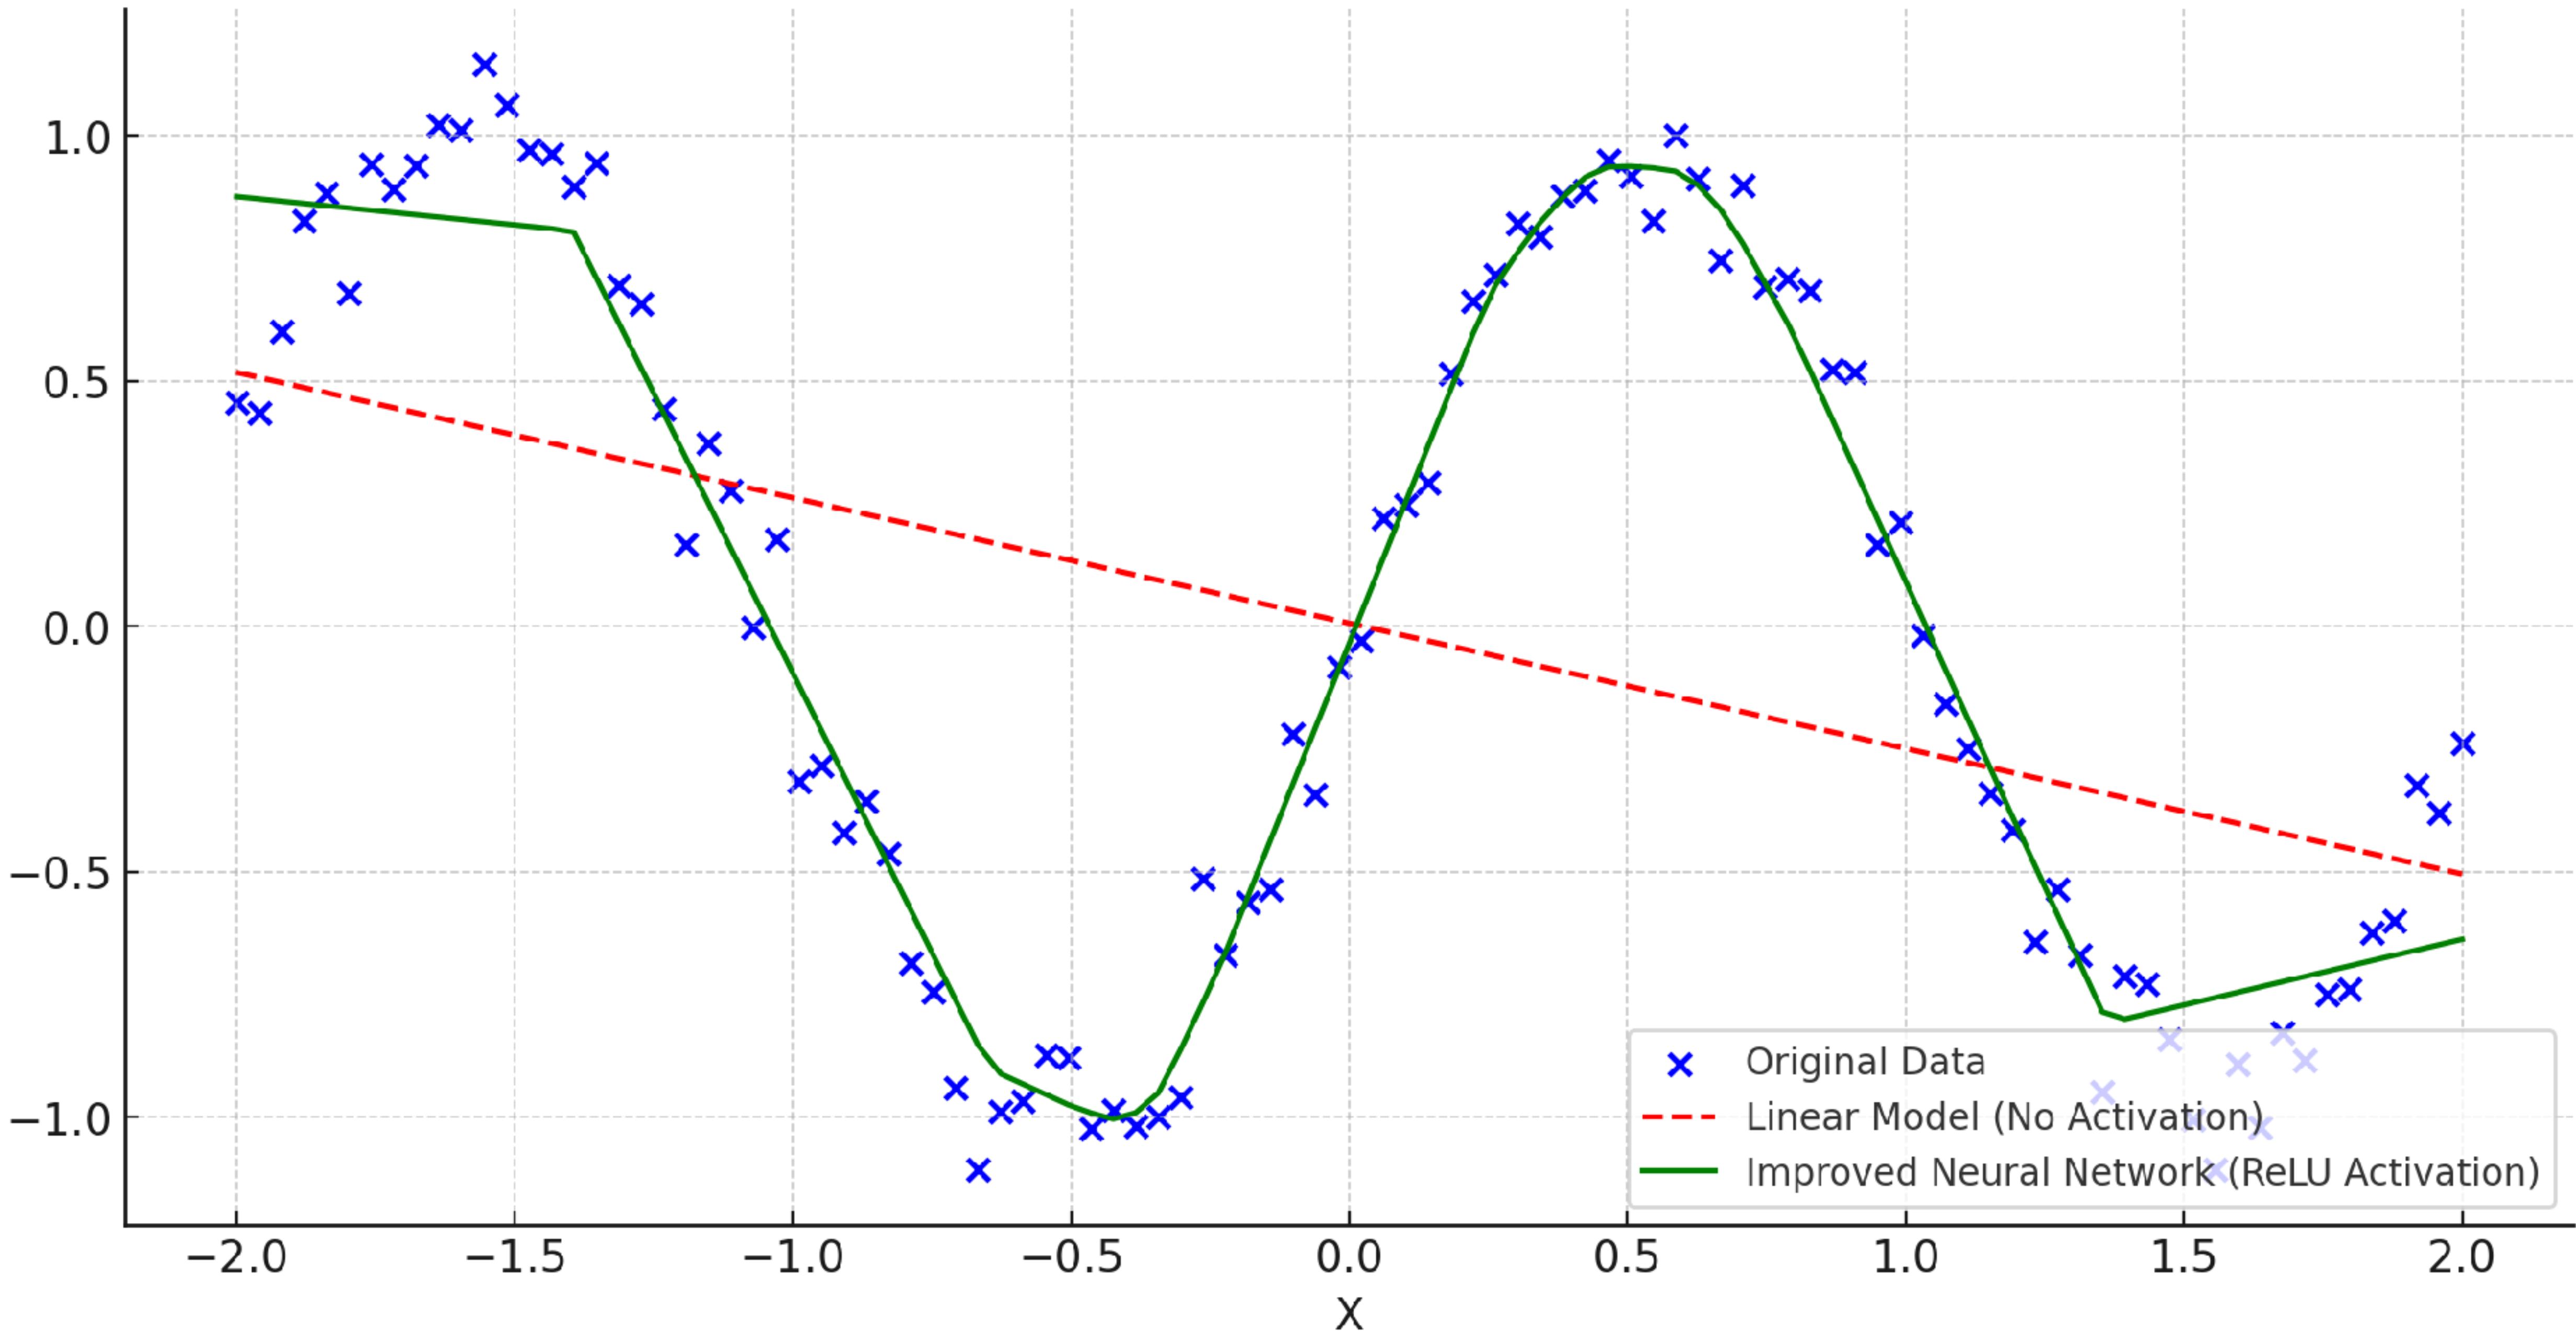
Dew Point [20, 15, 18, 17, 19, 16, 21, 18, 20, 17]

Cloud Cover [90, 40, 95, 60, 85, 35, 98, 70, 92, 50]

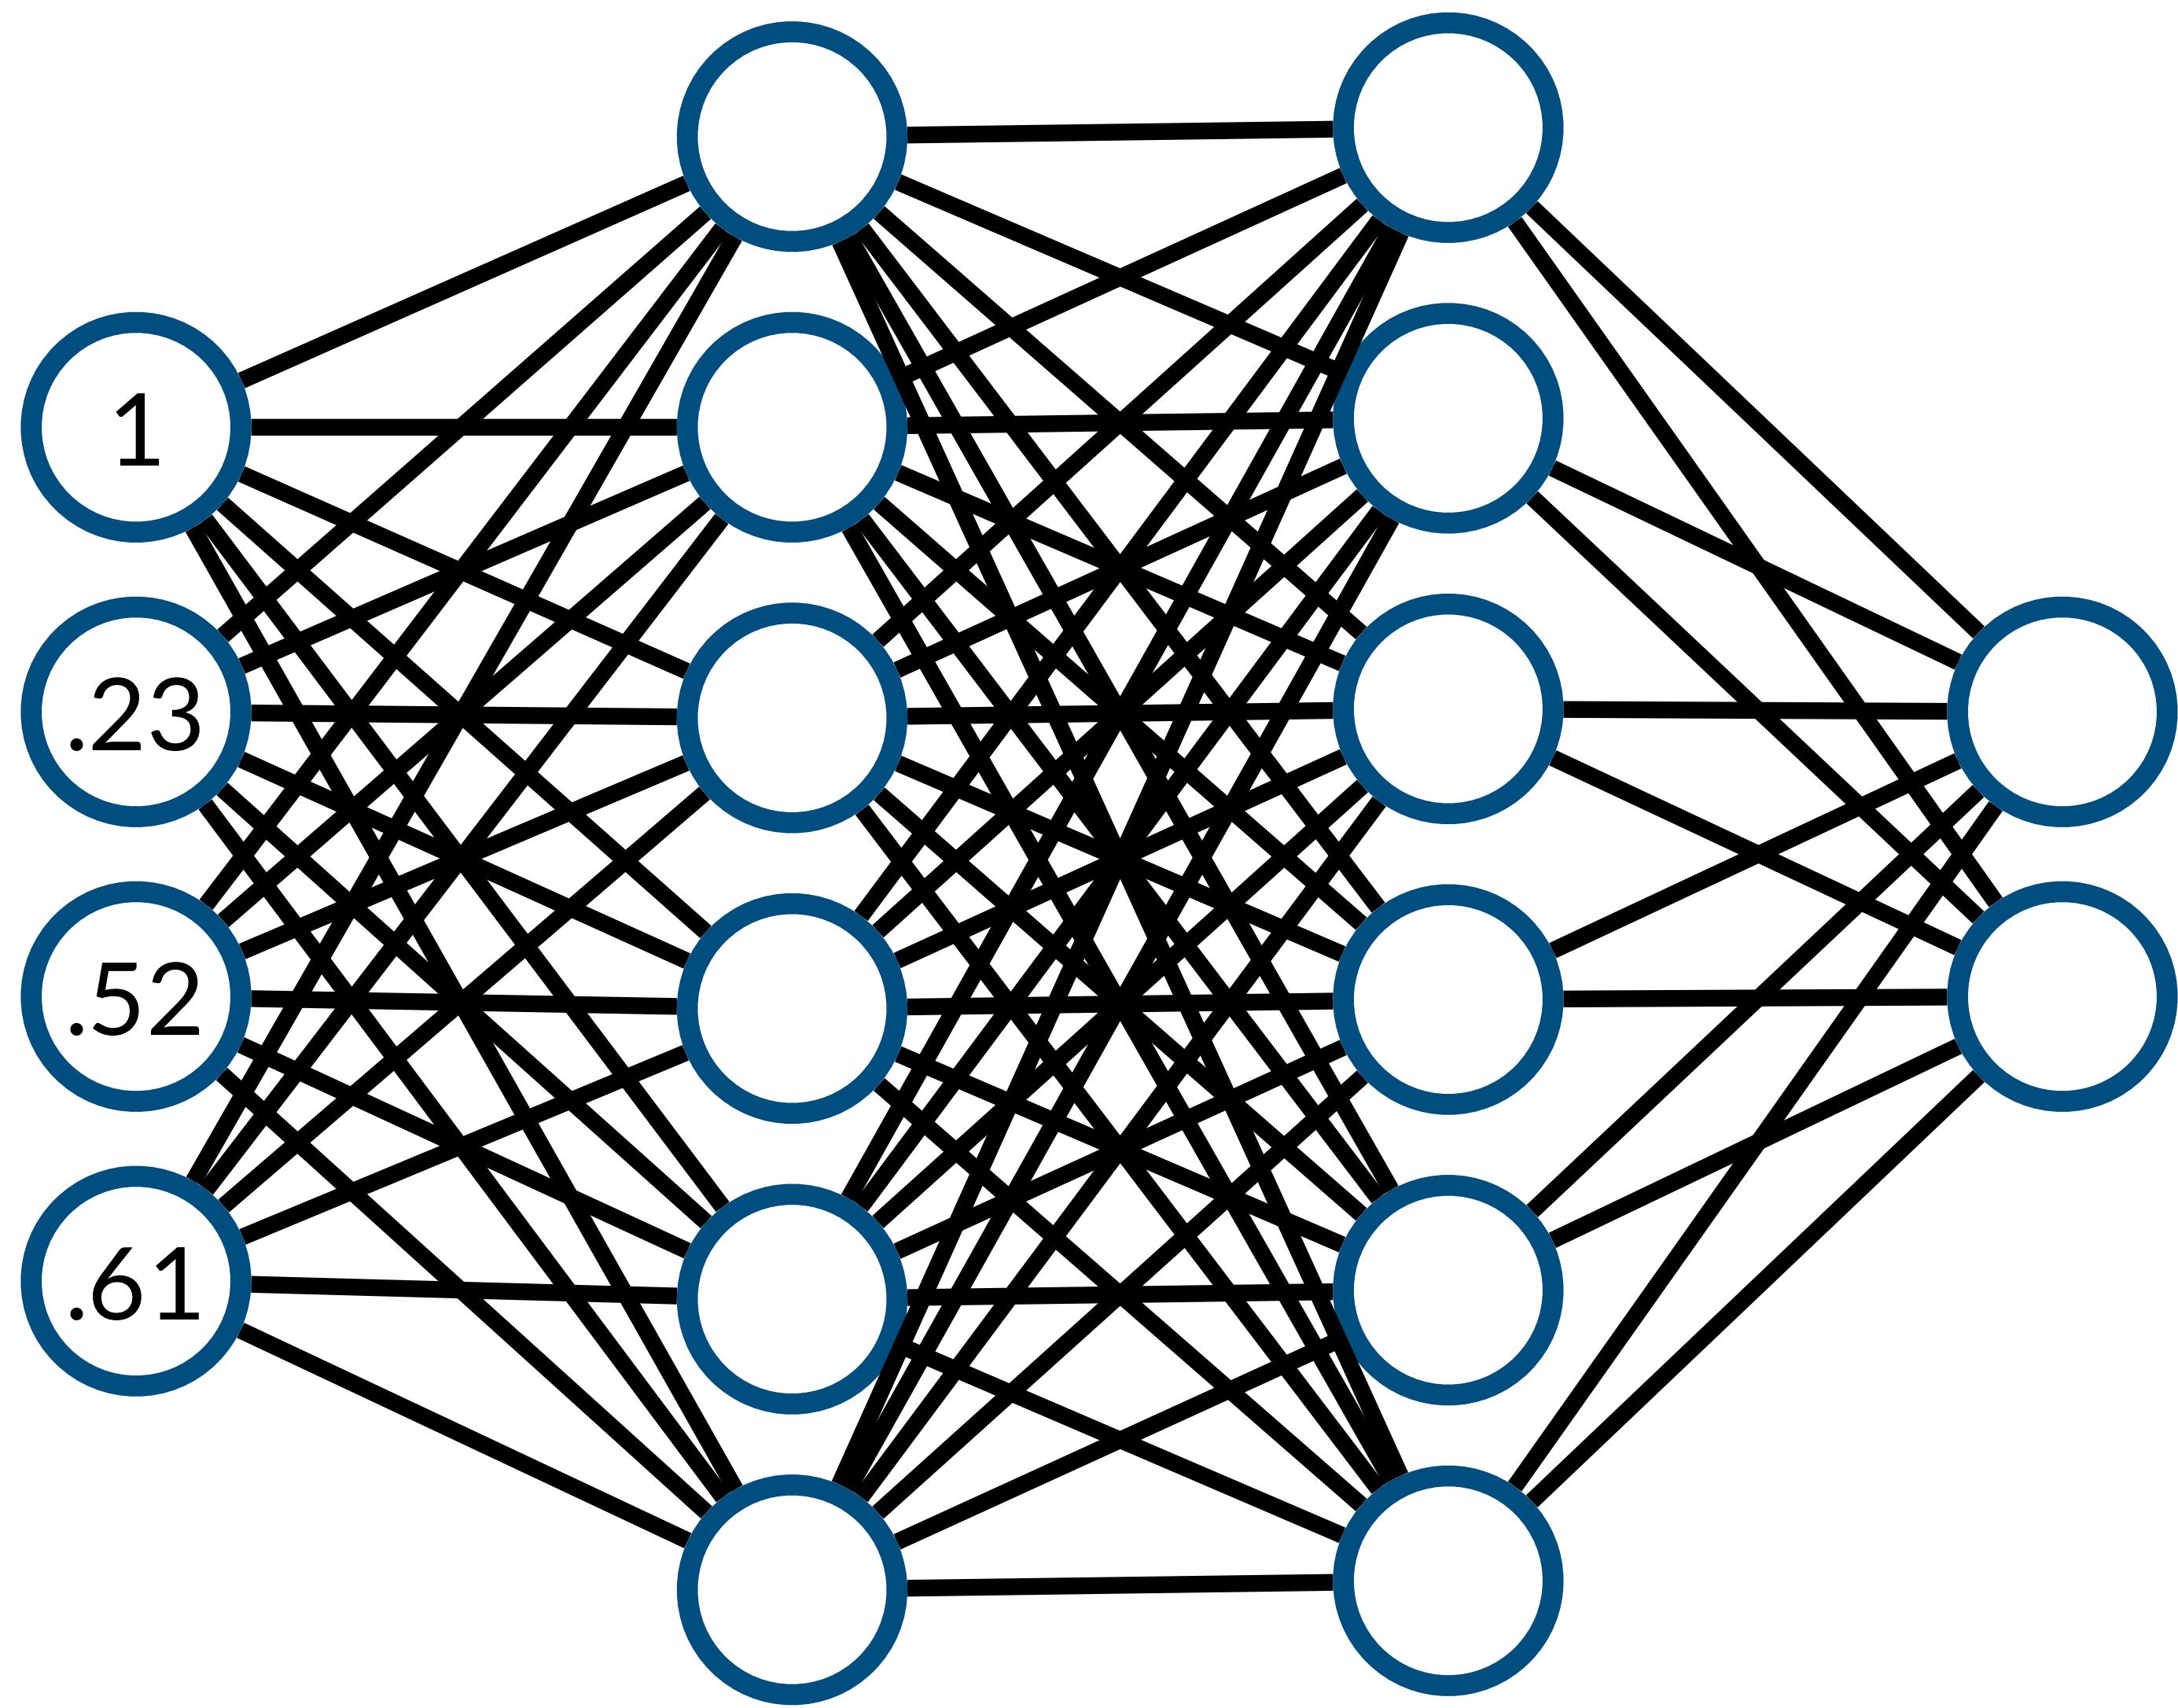


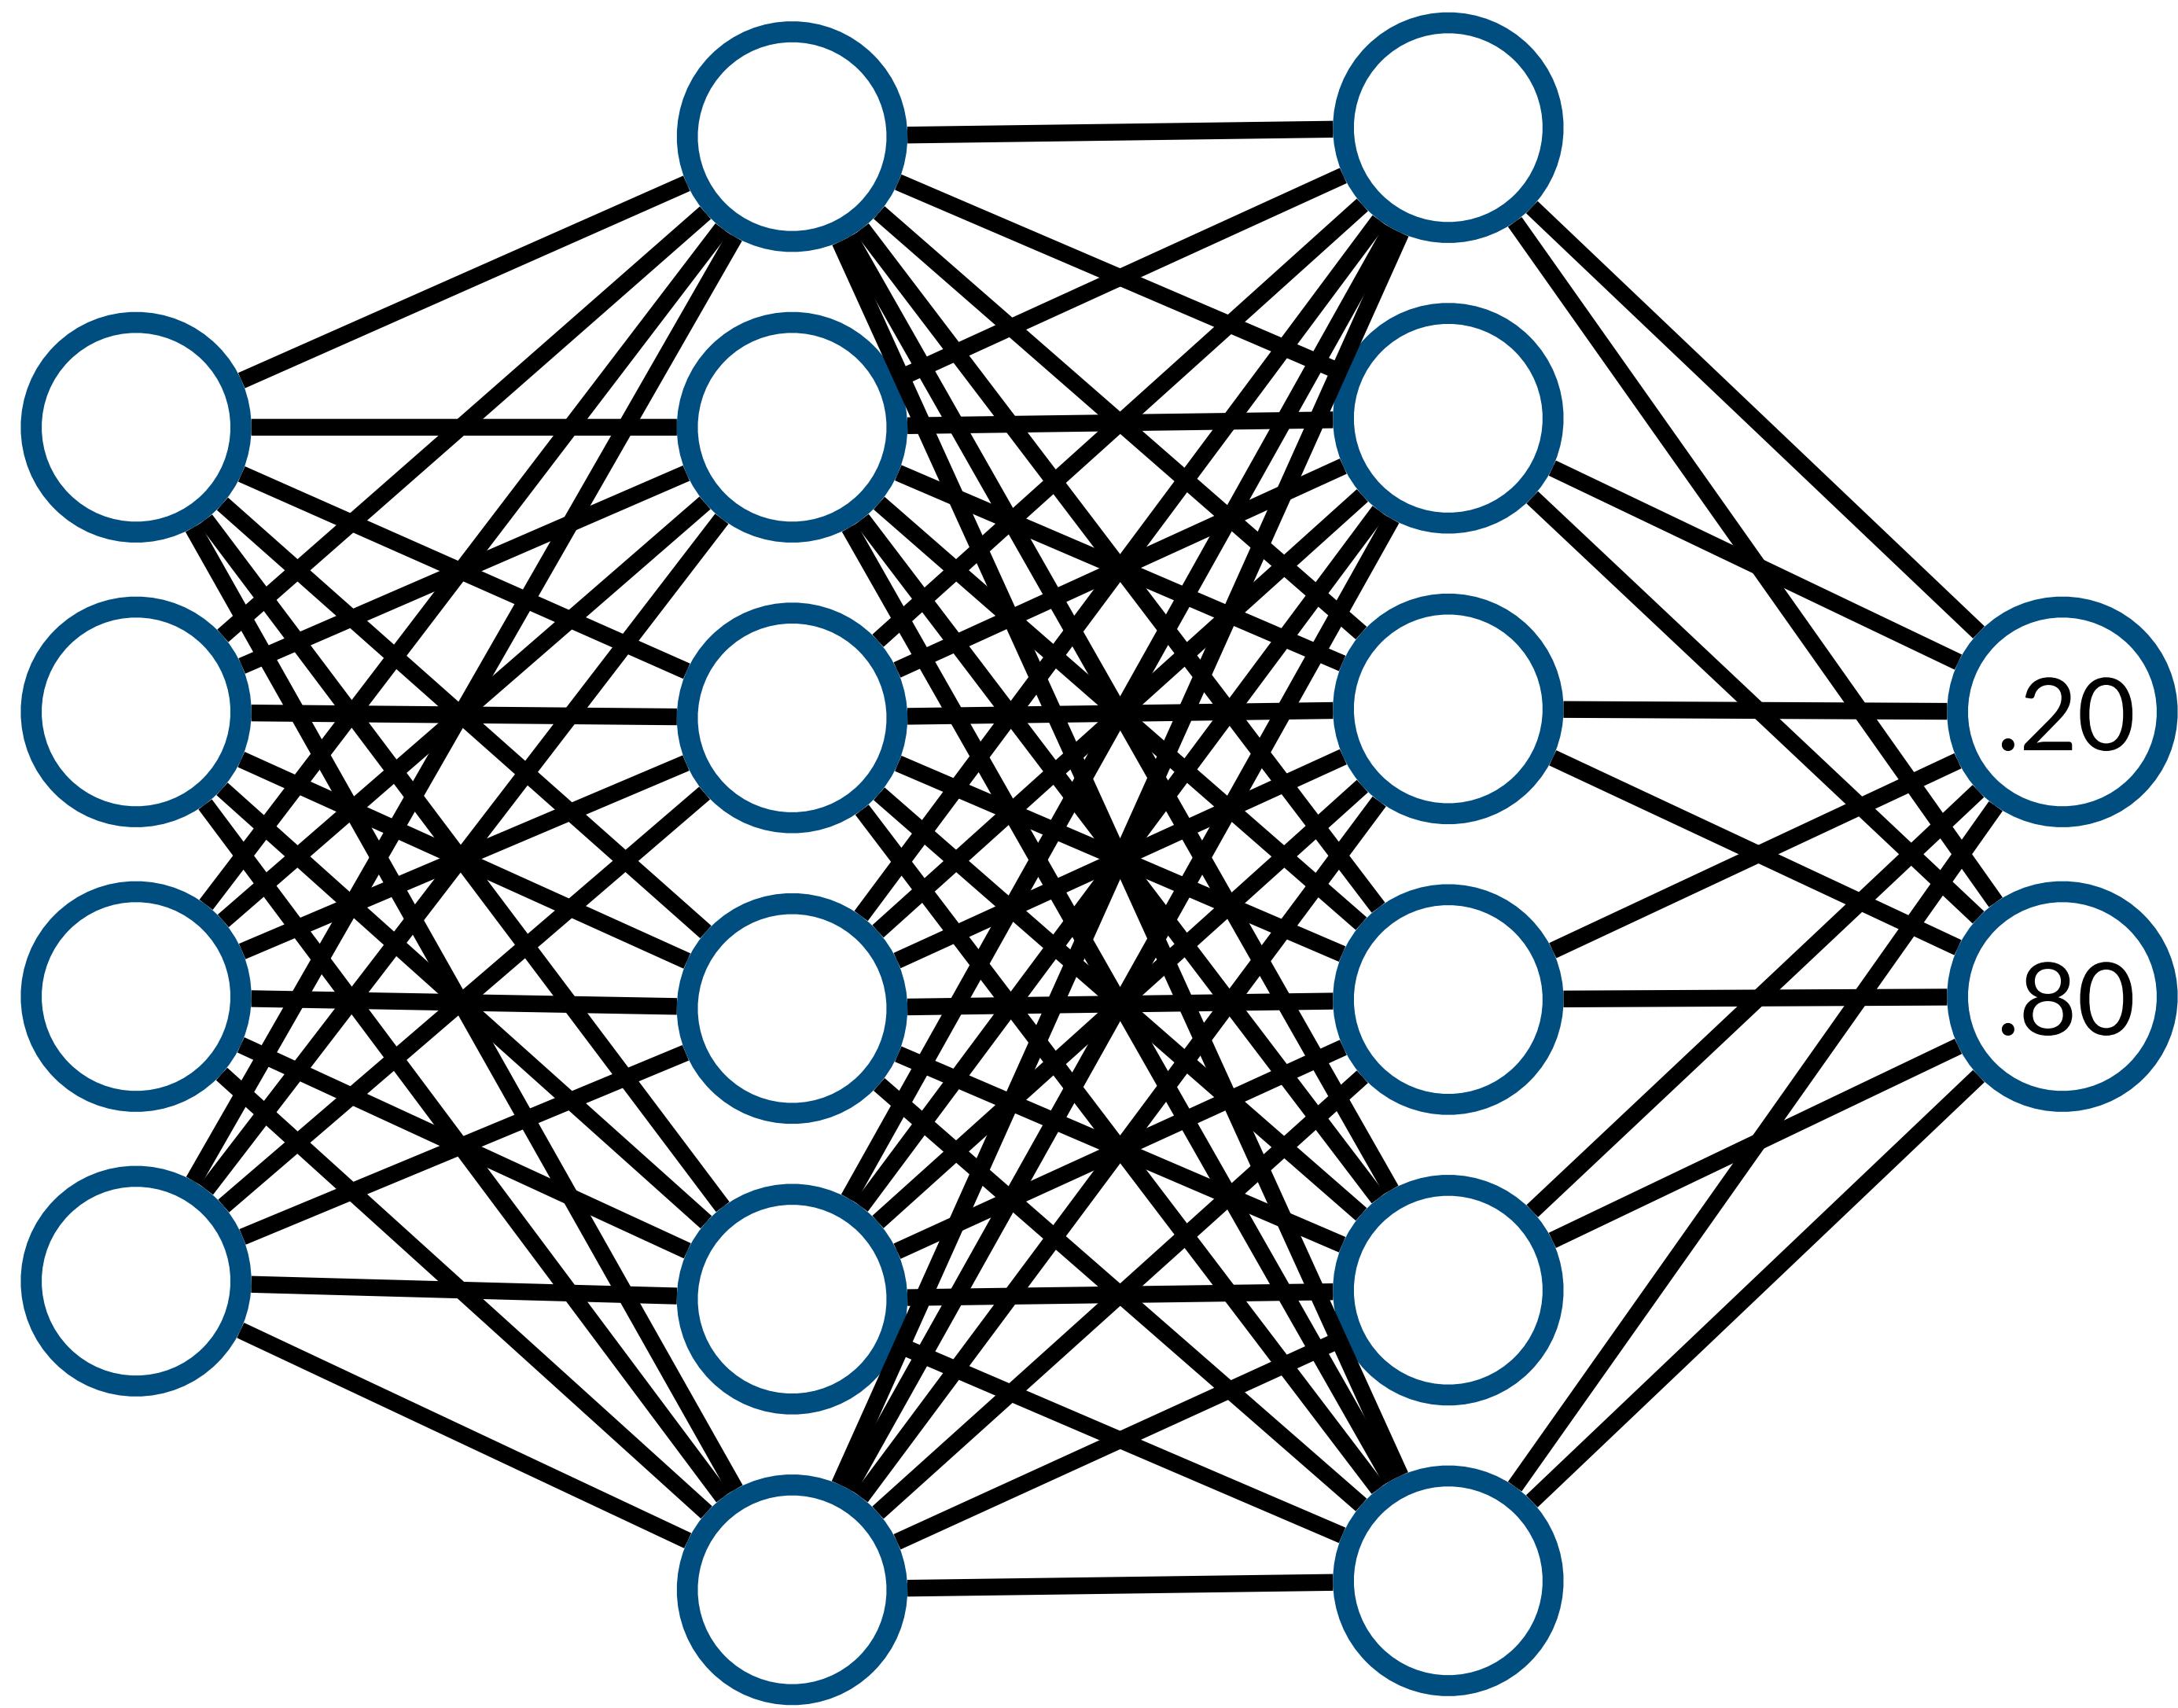


# Improved Non-Linear Fit with Neural Network (ReLU Activation)



# Activation Function

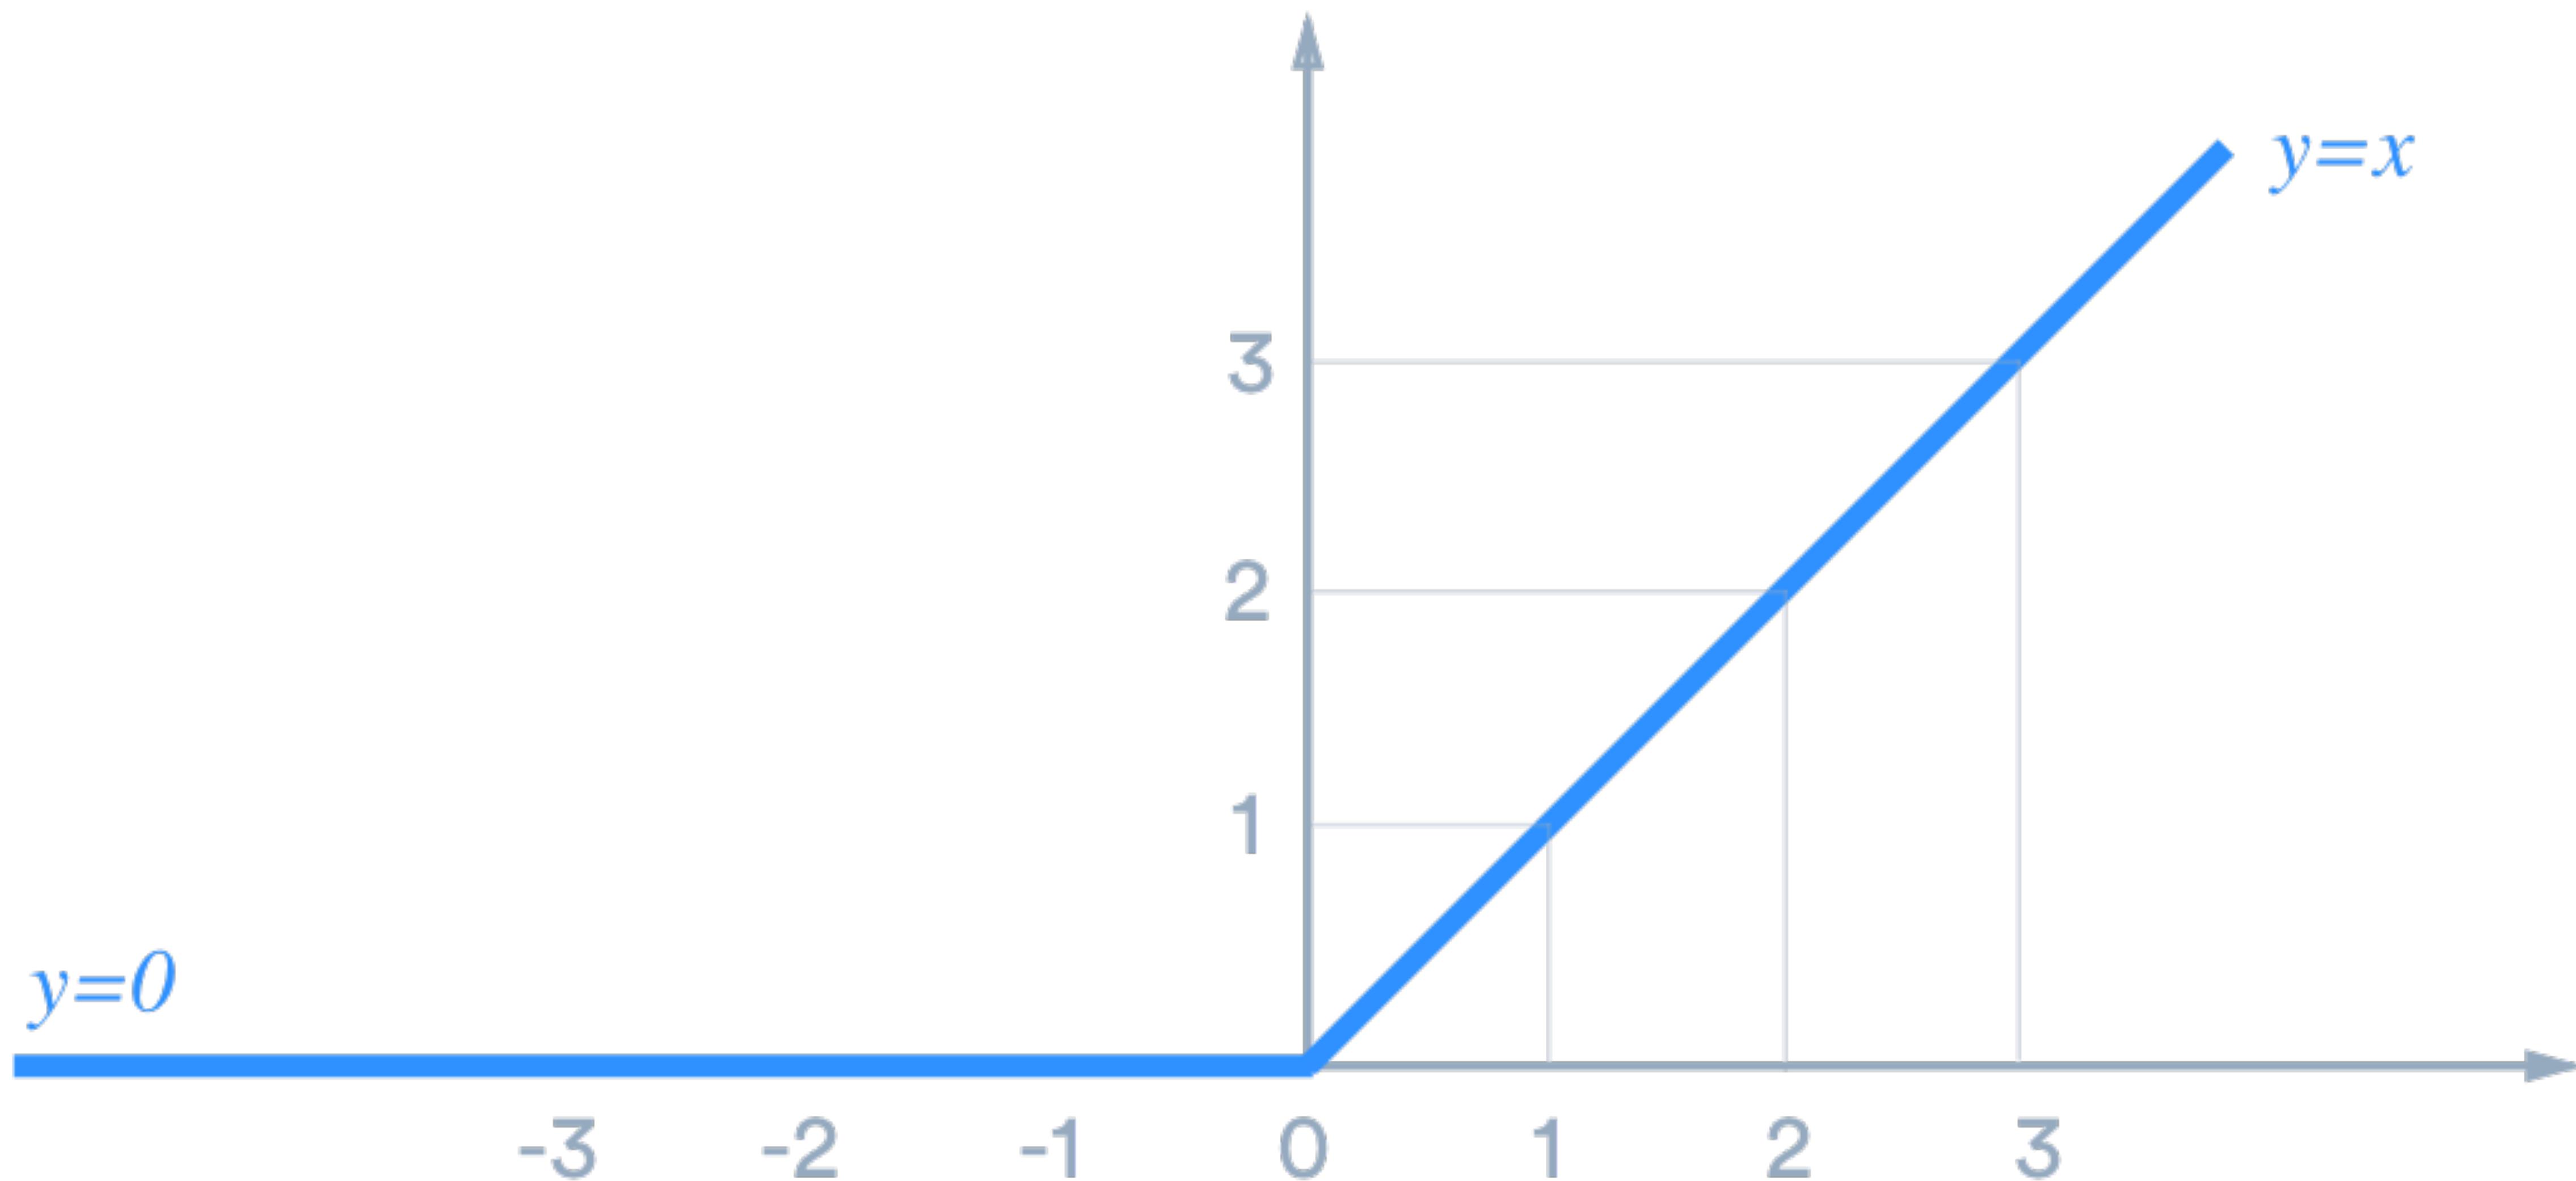




# Activation Function

- **Definition** - A non-linear function which decides if the output of a neuron should be propagated or allowed to continue forward
- Without activation functions, neural networks will not be able to solve non-linear problems, without it, it is the same a linear regression
- Functions include:
  - Sigmoid (final - candidate)
  - tanh
  - relu
  - selu
  - leaky relu
  - softmax (final)

# Rectified Linear Unit



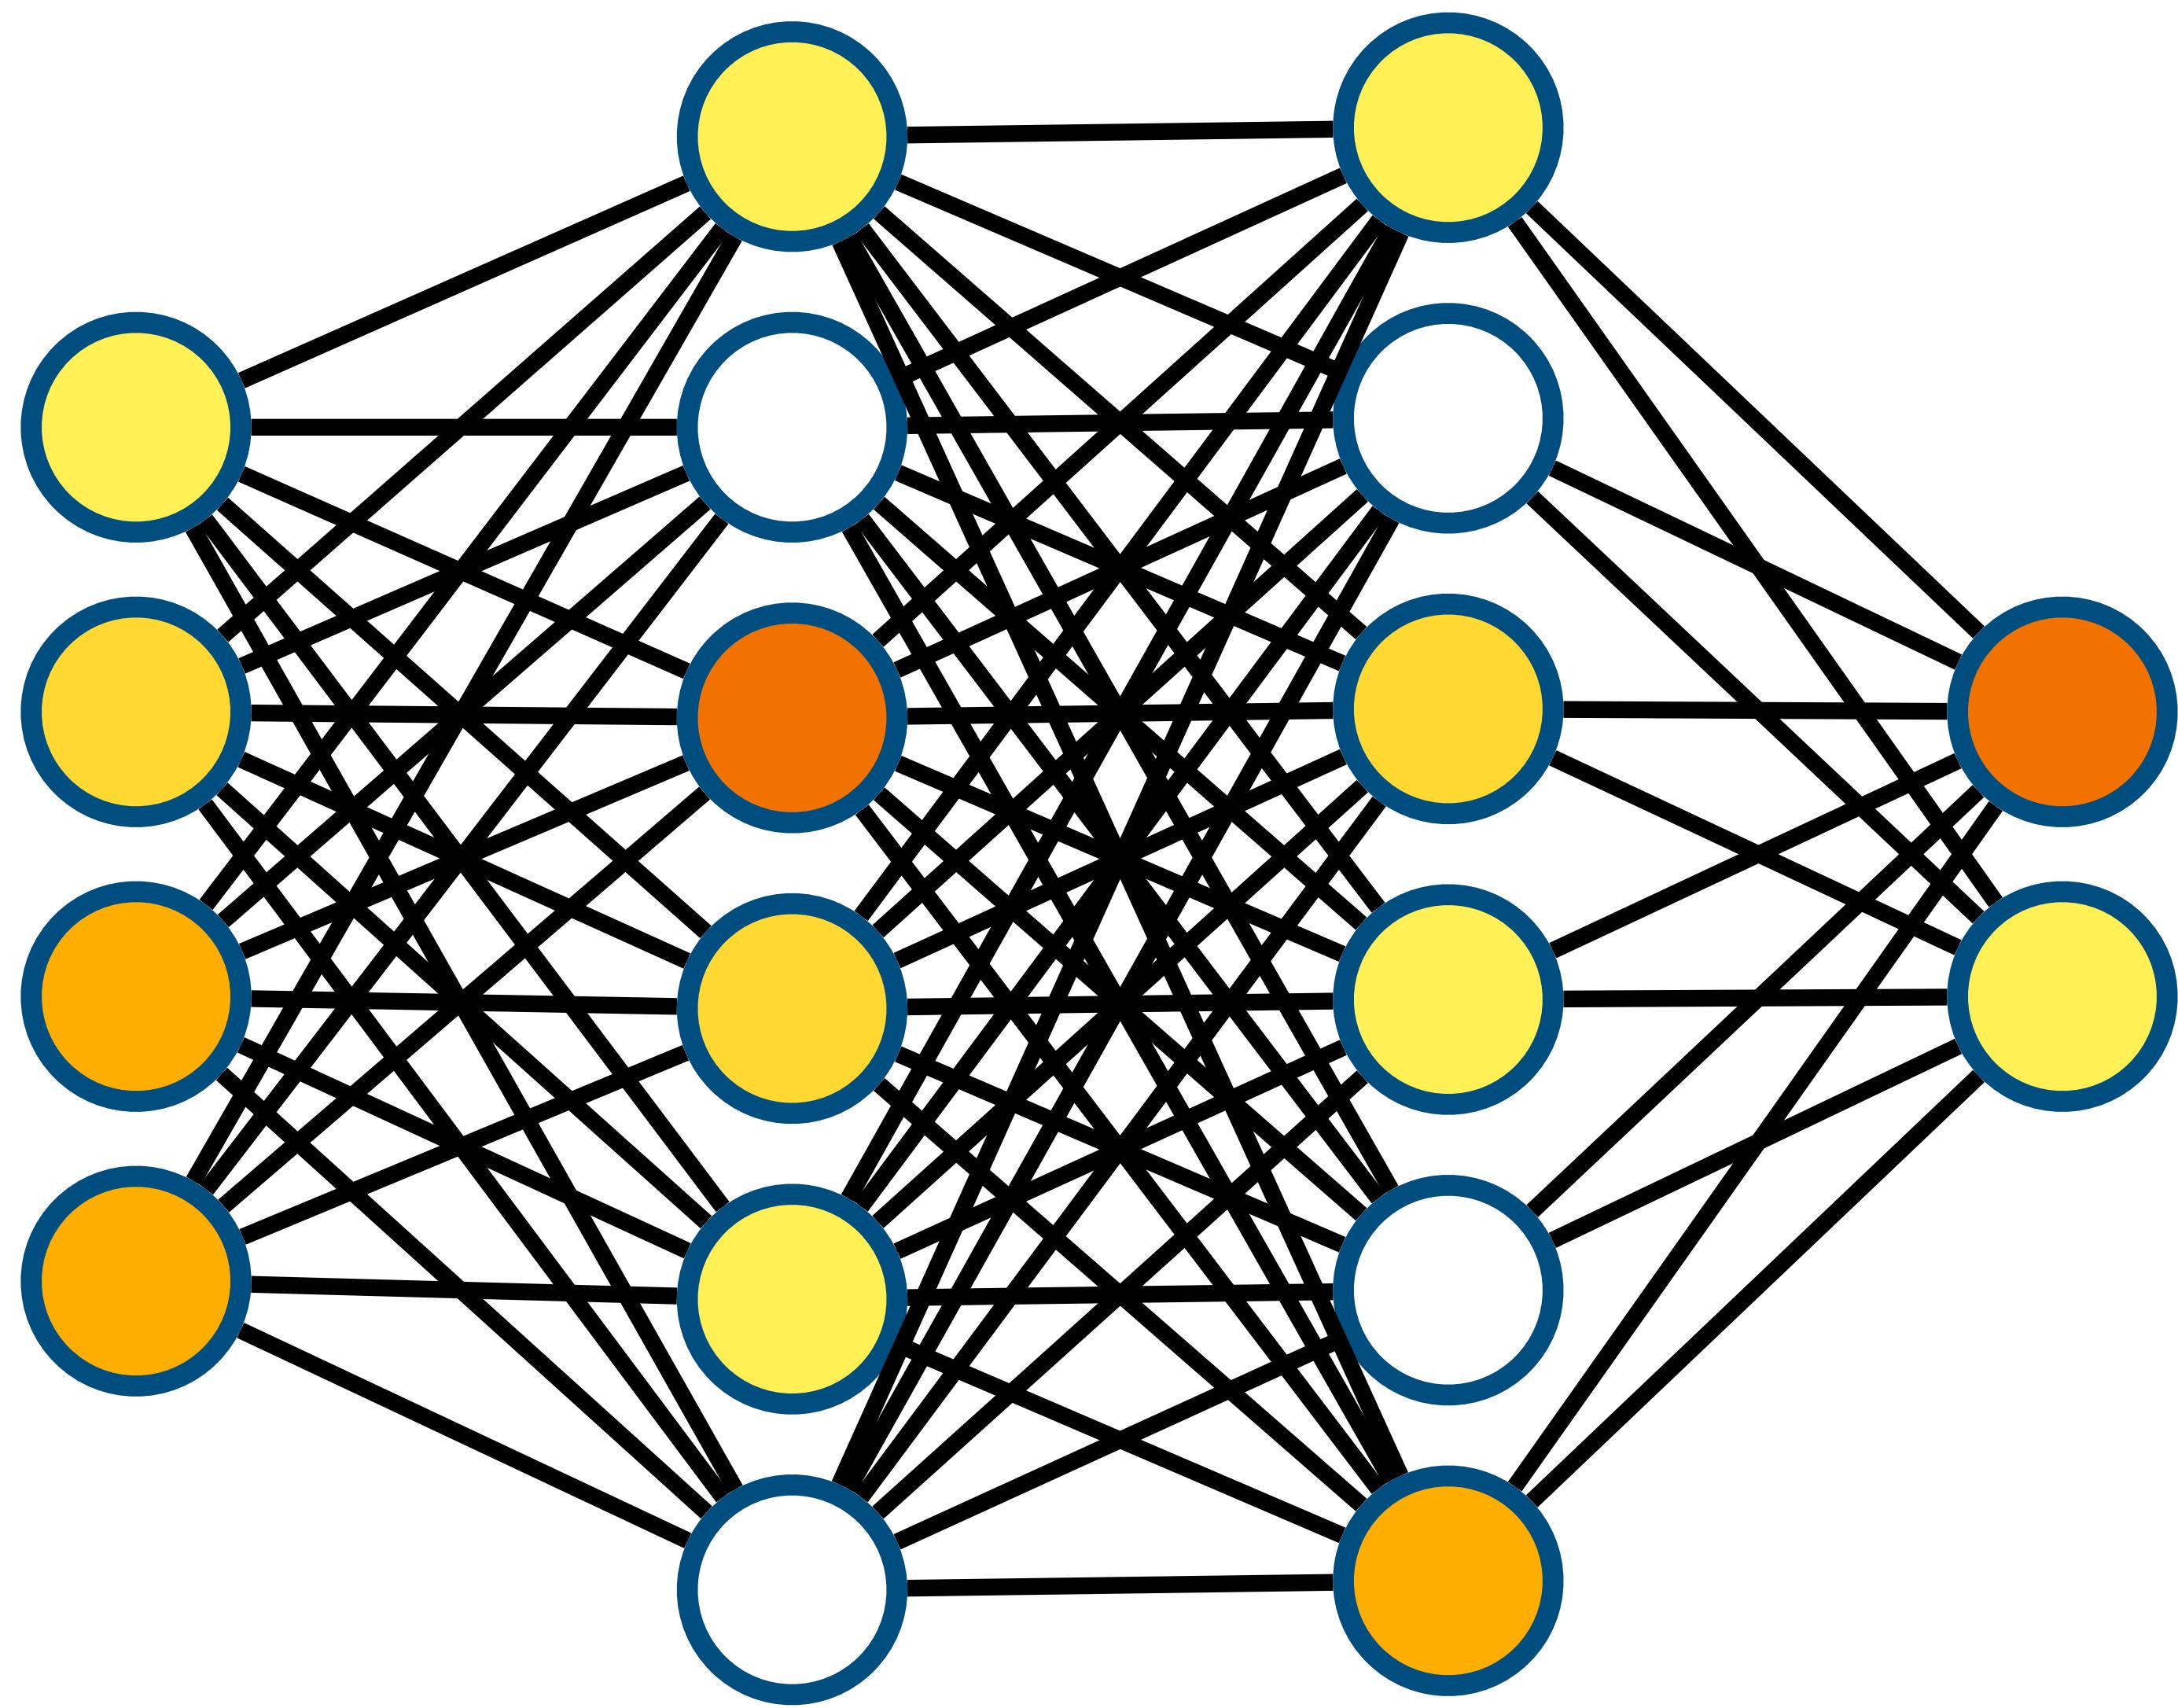


Image 1: Boot



Image 2: Pullover



Image 3: Trouser



Image 4: Trouser



Image 5: Shirt



Image 6: Trouser

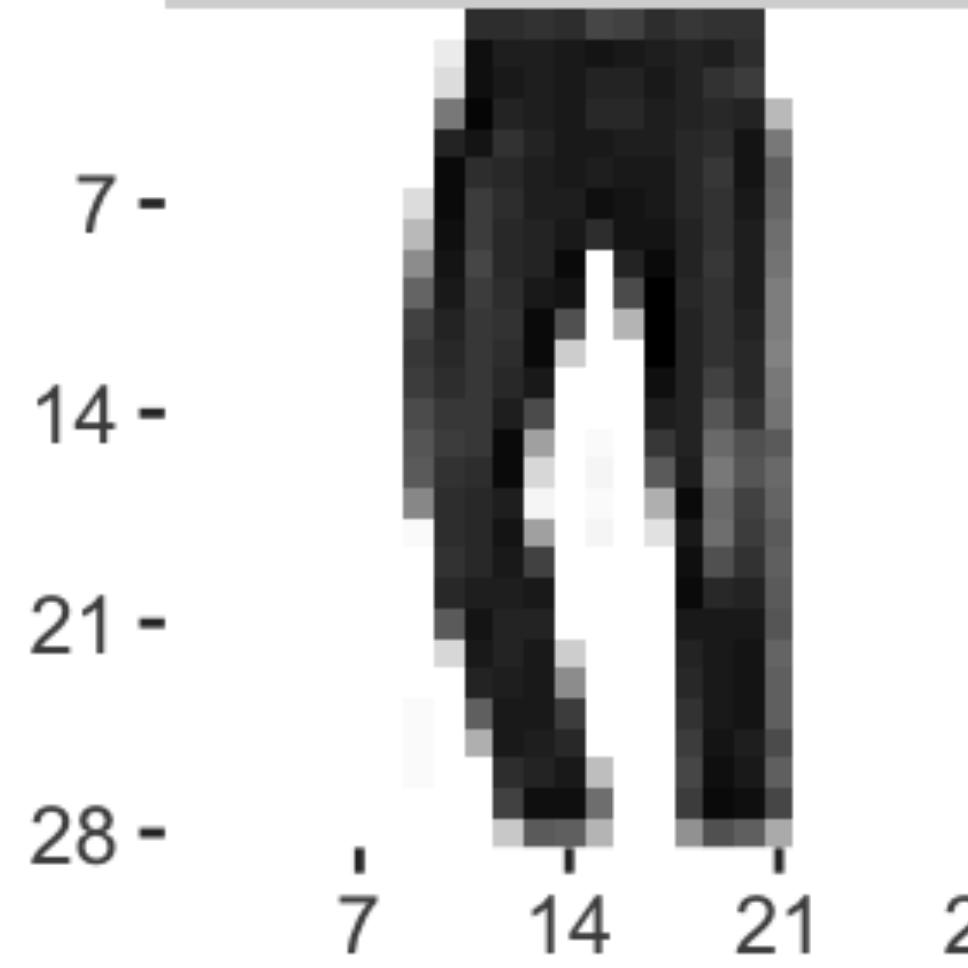


Image 7: Coat

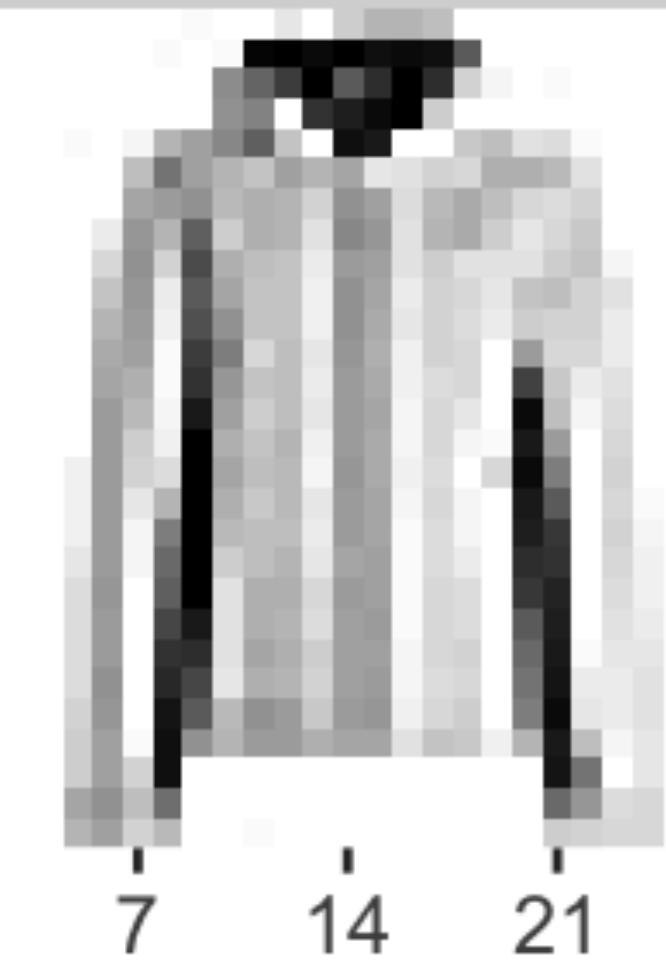


Image 8: Shirt

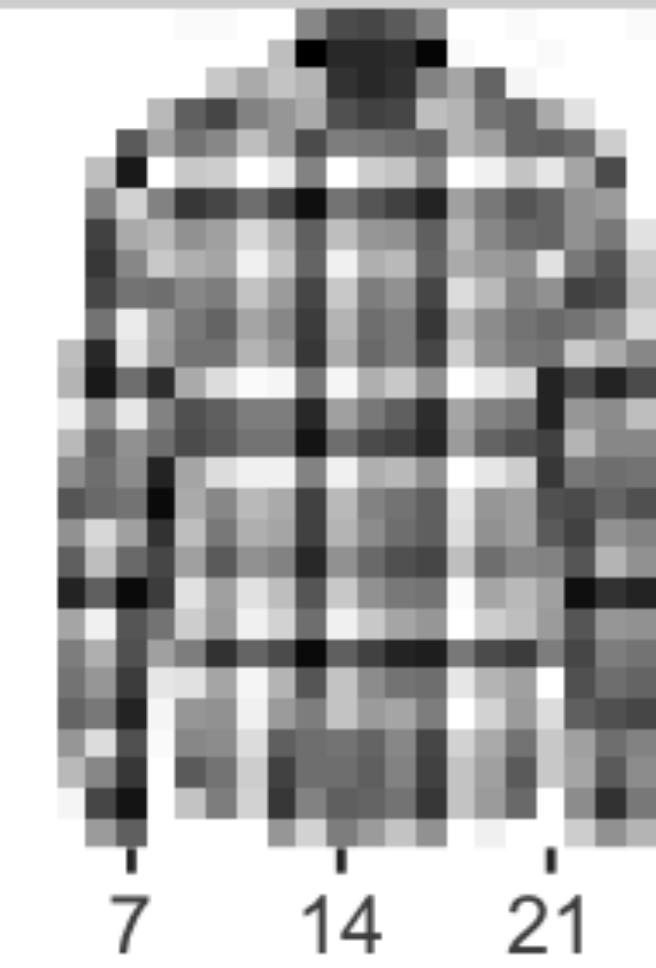


Image 9: Sandal

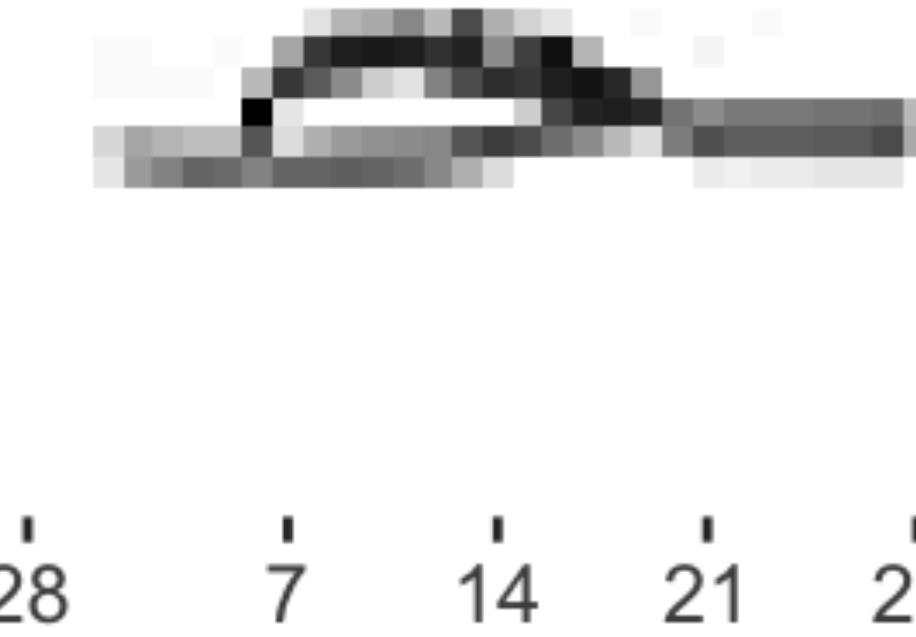


Image 10: Sneaker

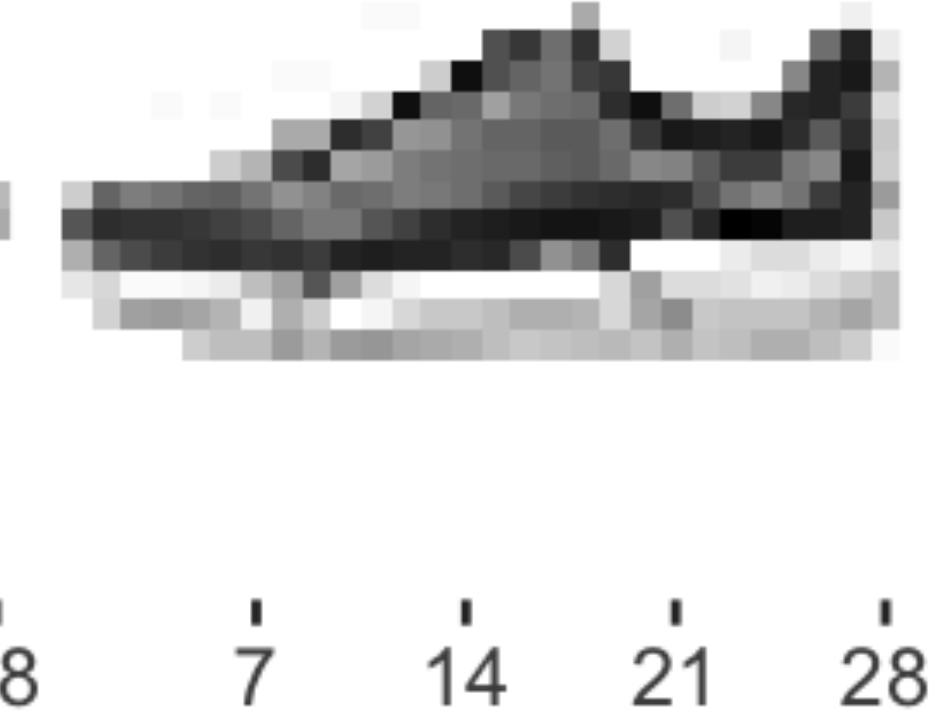


Image 1: Boot



Image 2: Pullover



Image 3: Trouser



Image 4: Trouser

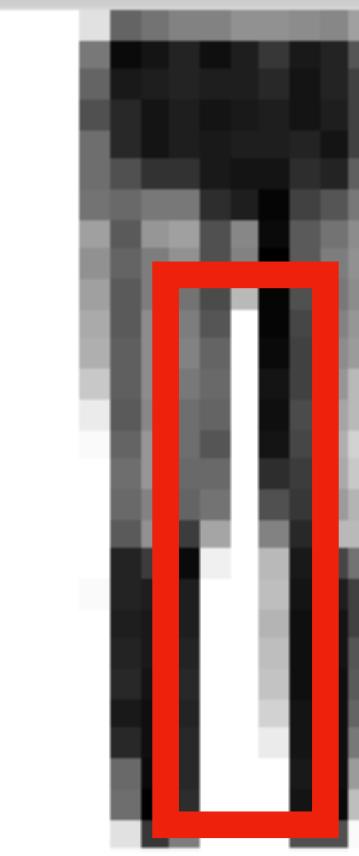


Image 5: Shirt



Image 6: Trouser

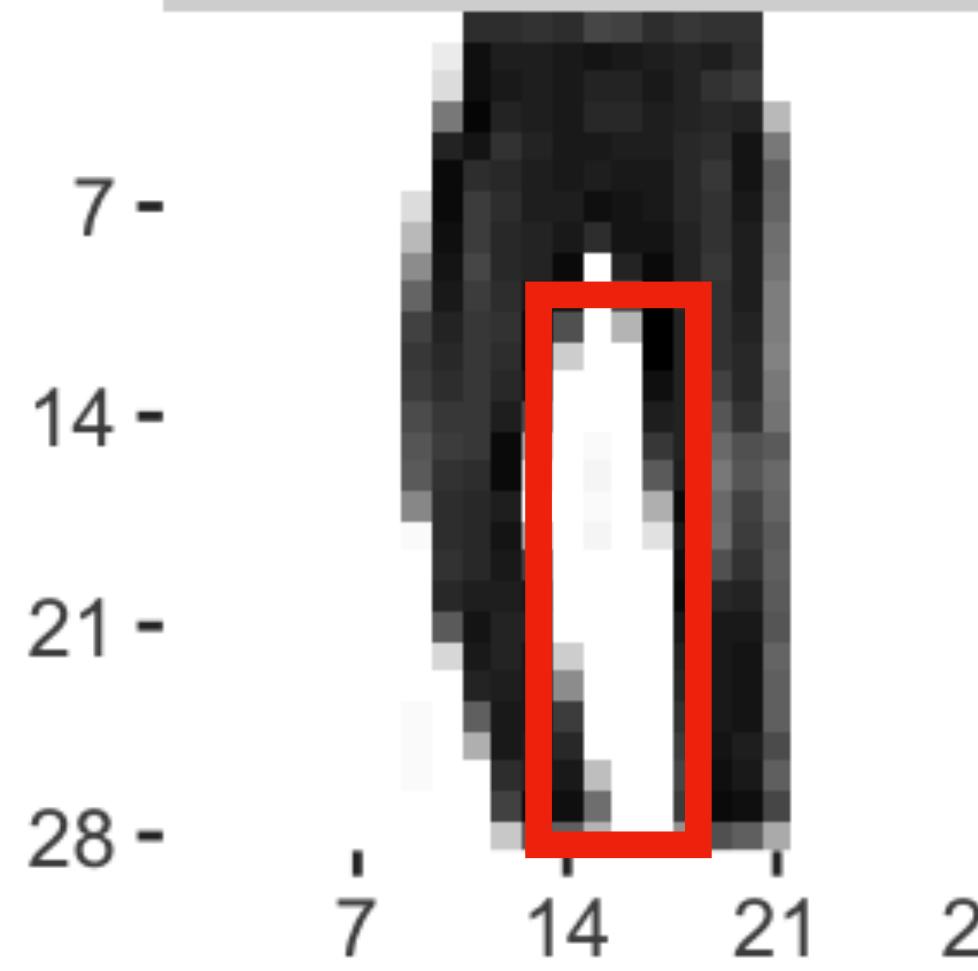


Image 7: Coat

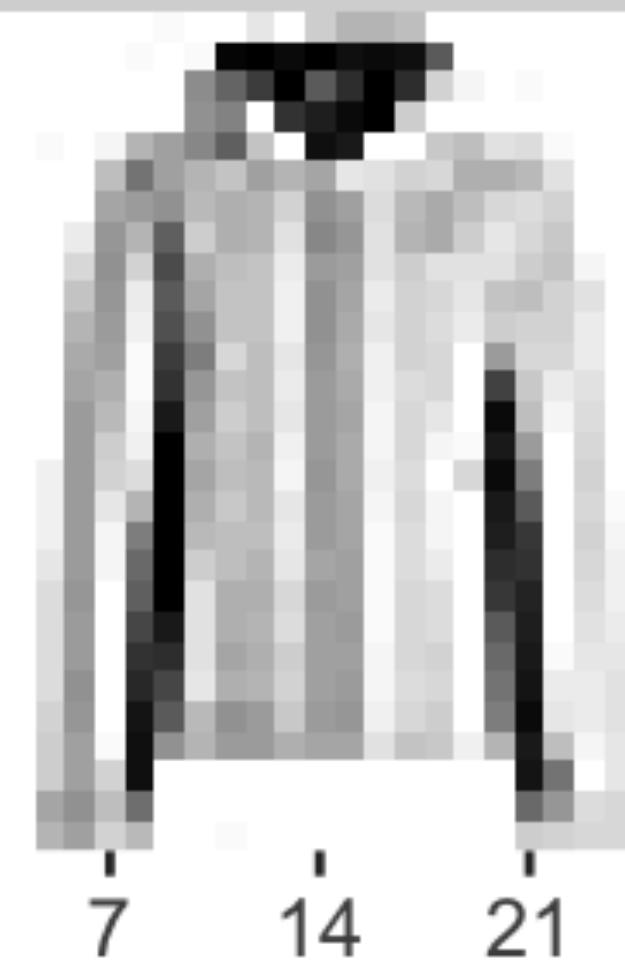


Image 8: Shirt

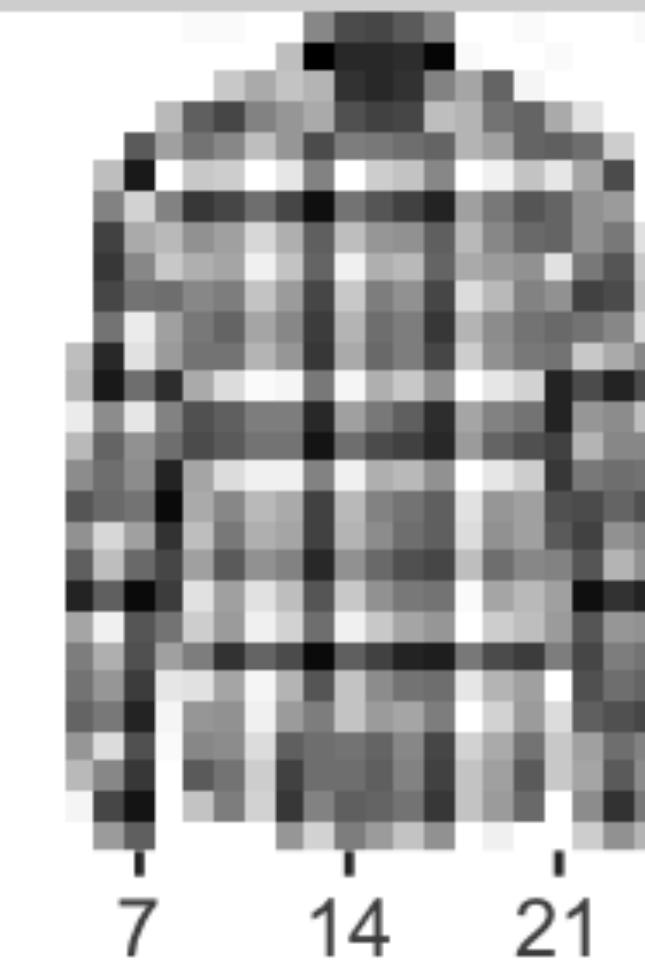


Image 9: Sandal

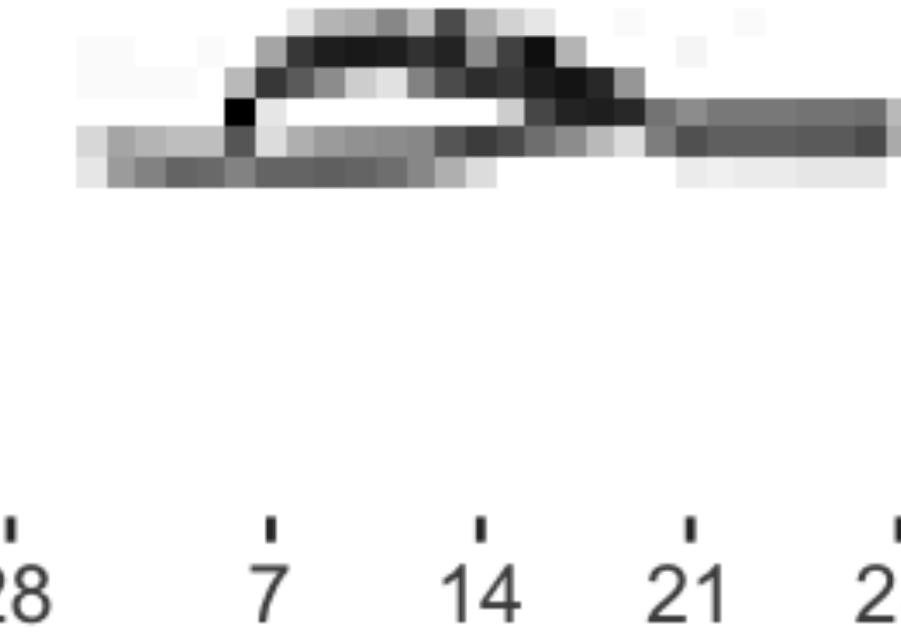
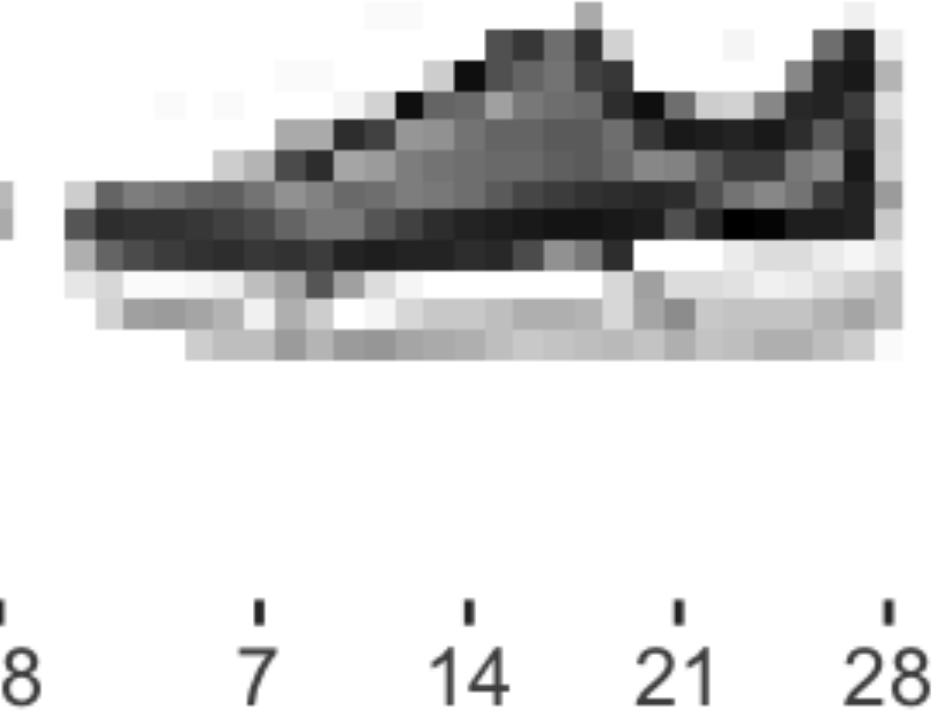
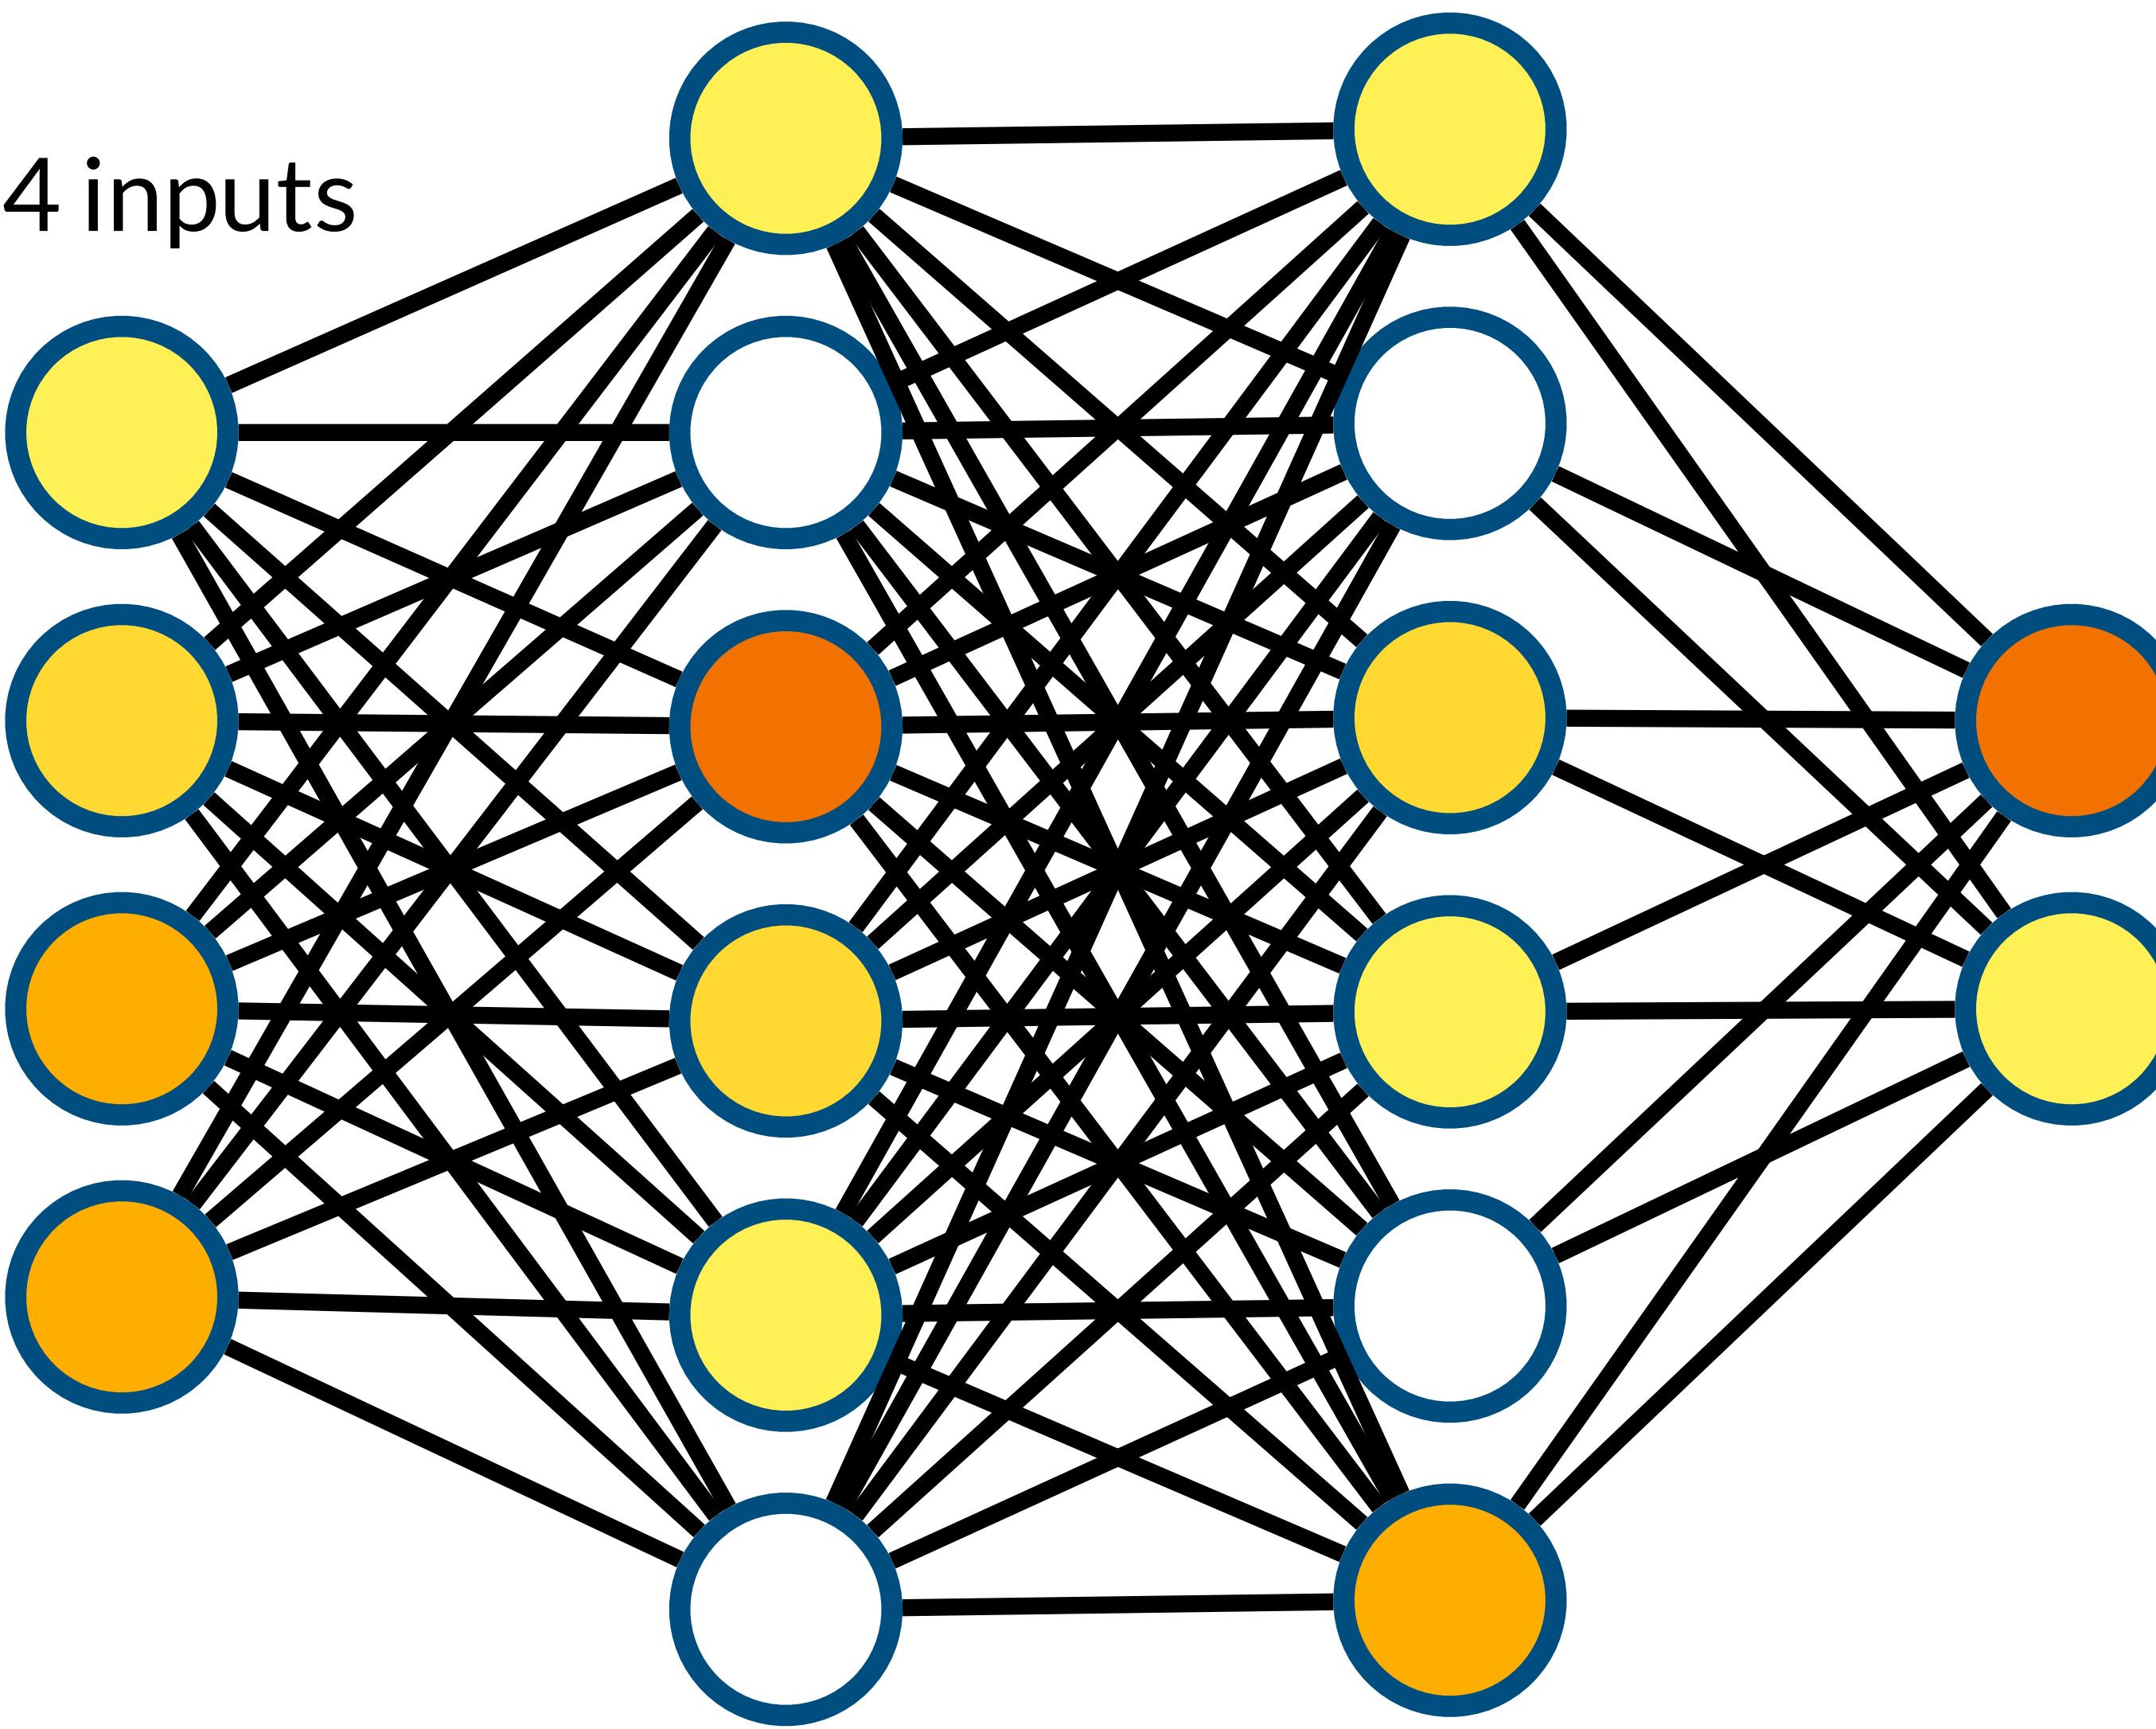


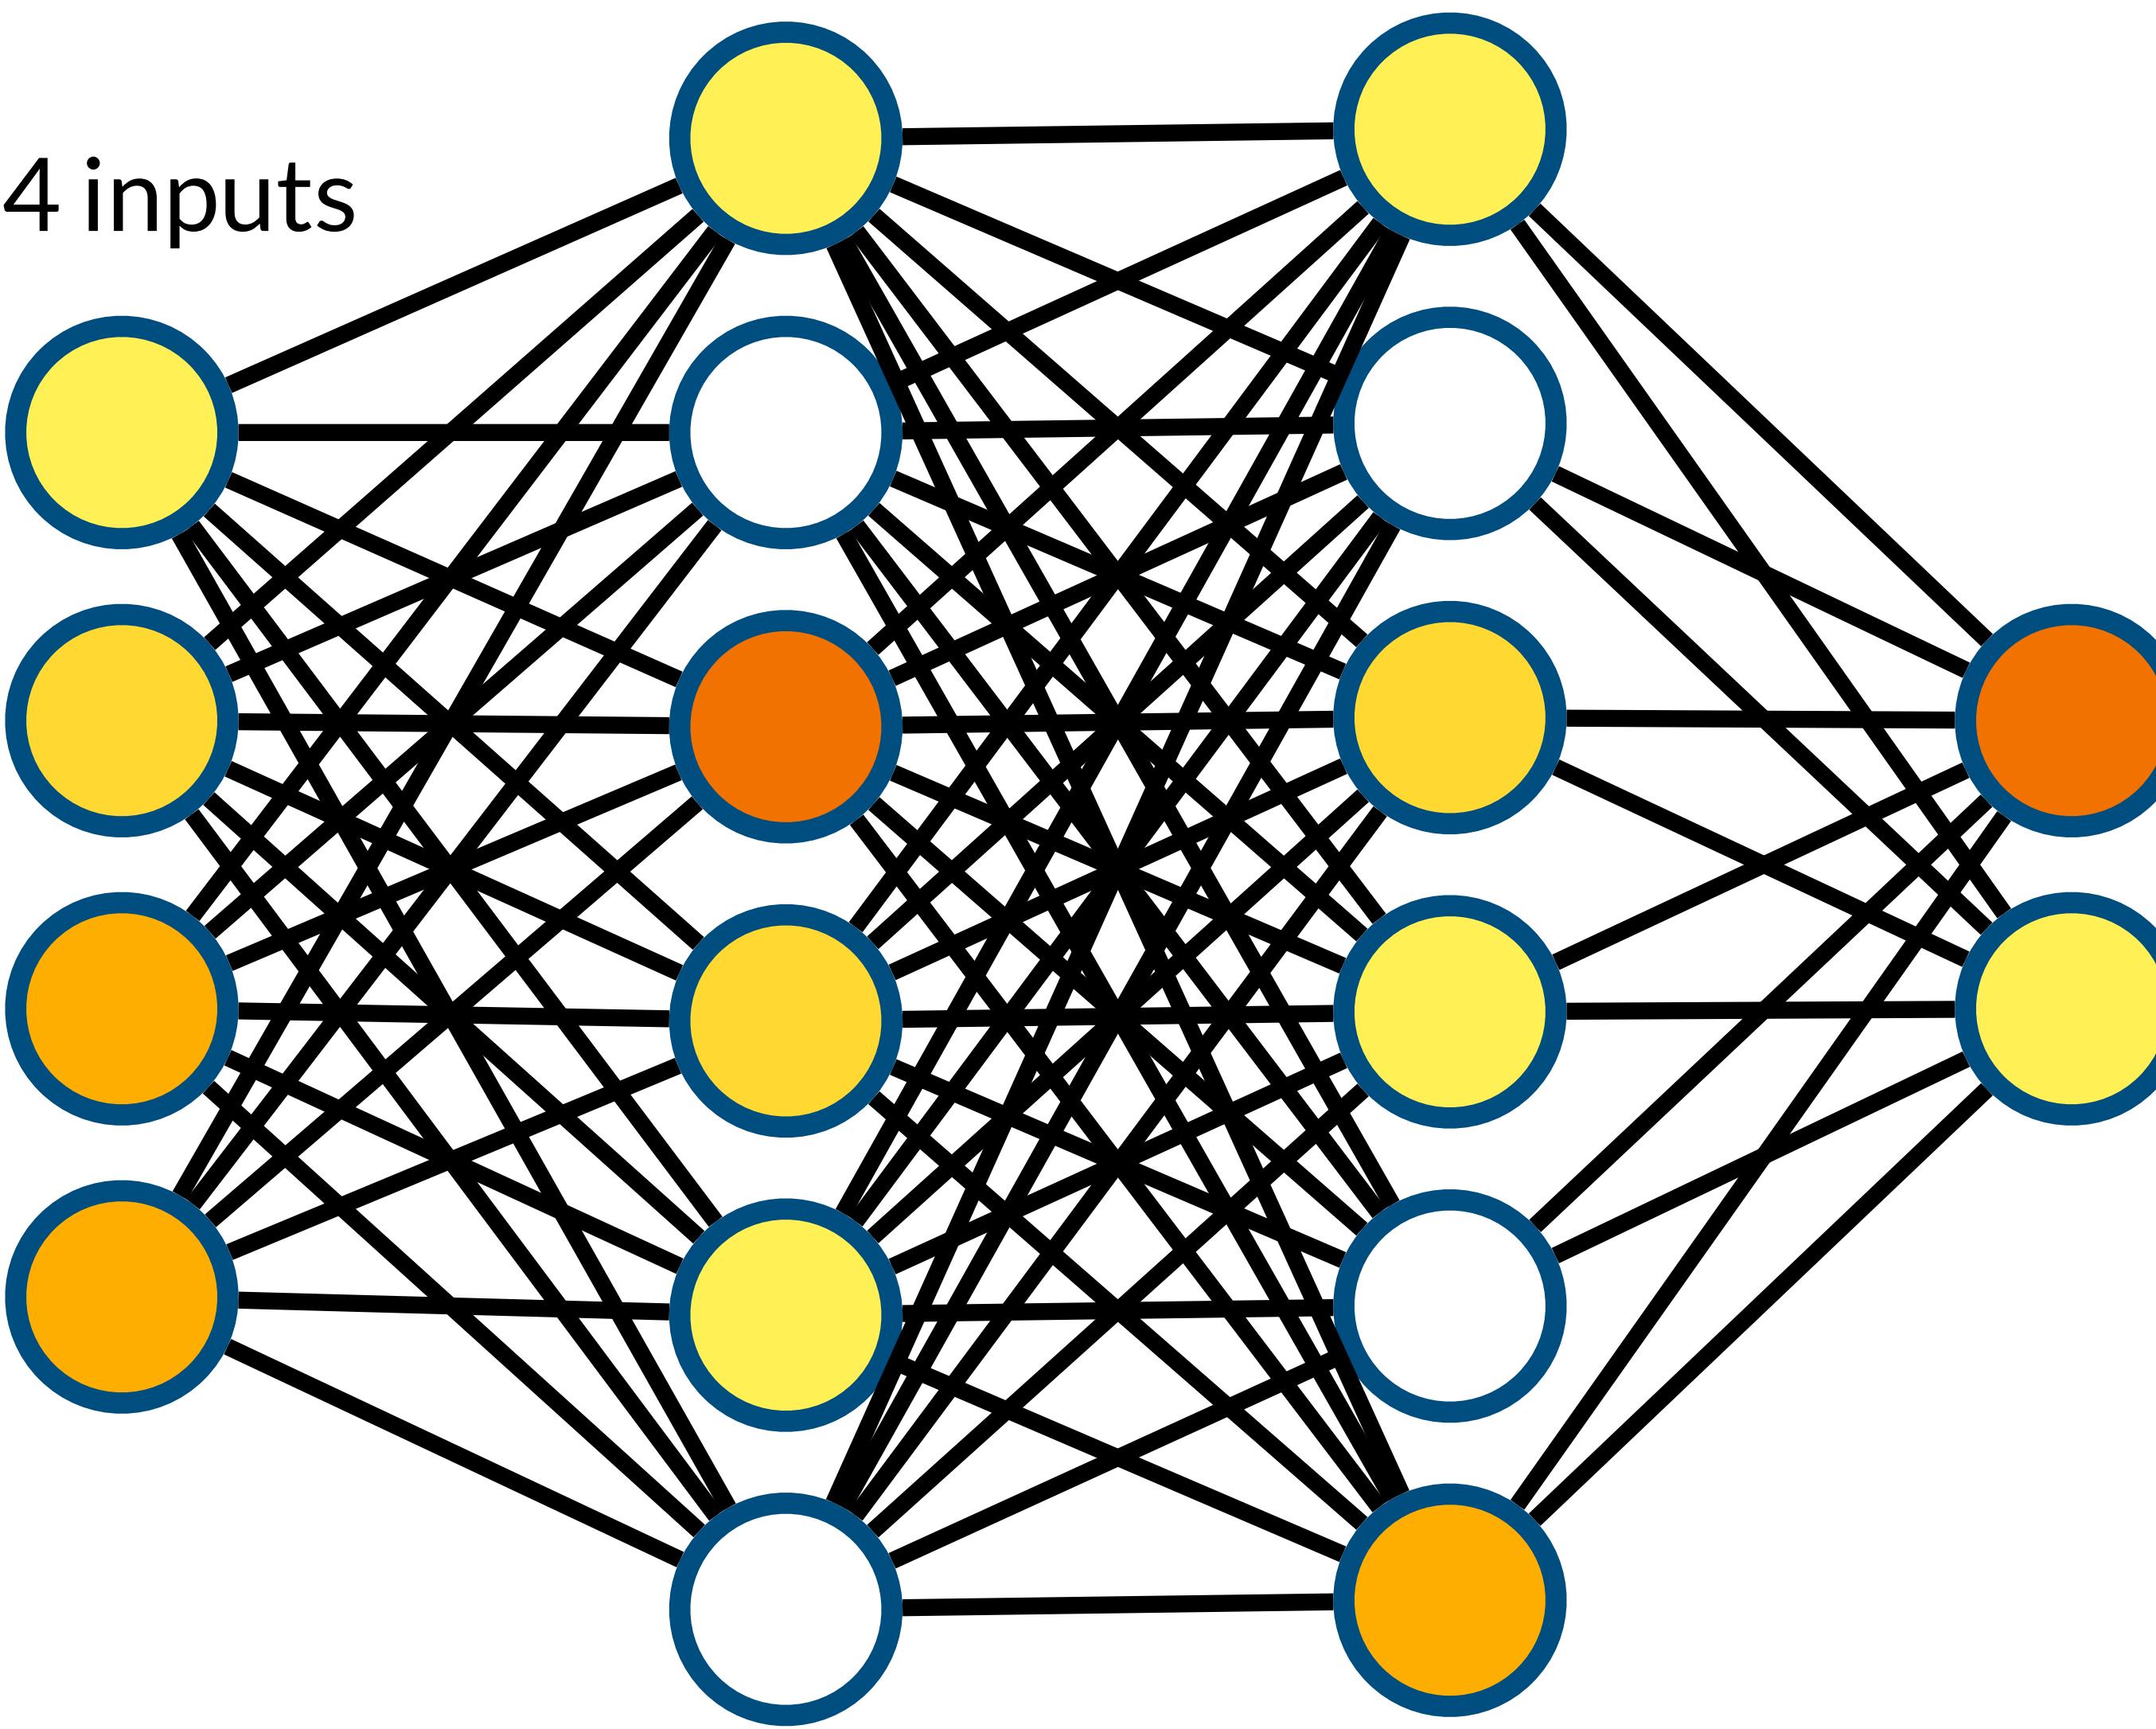
Image 10: Sneaker

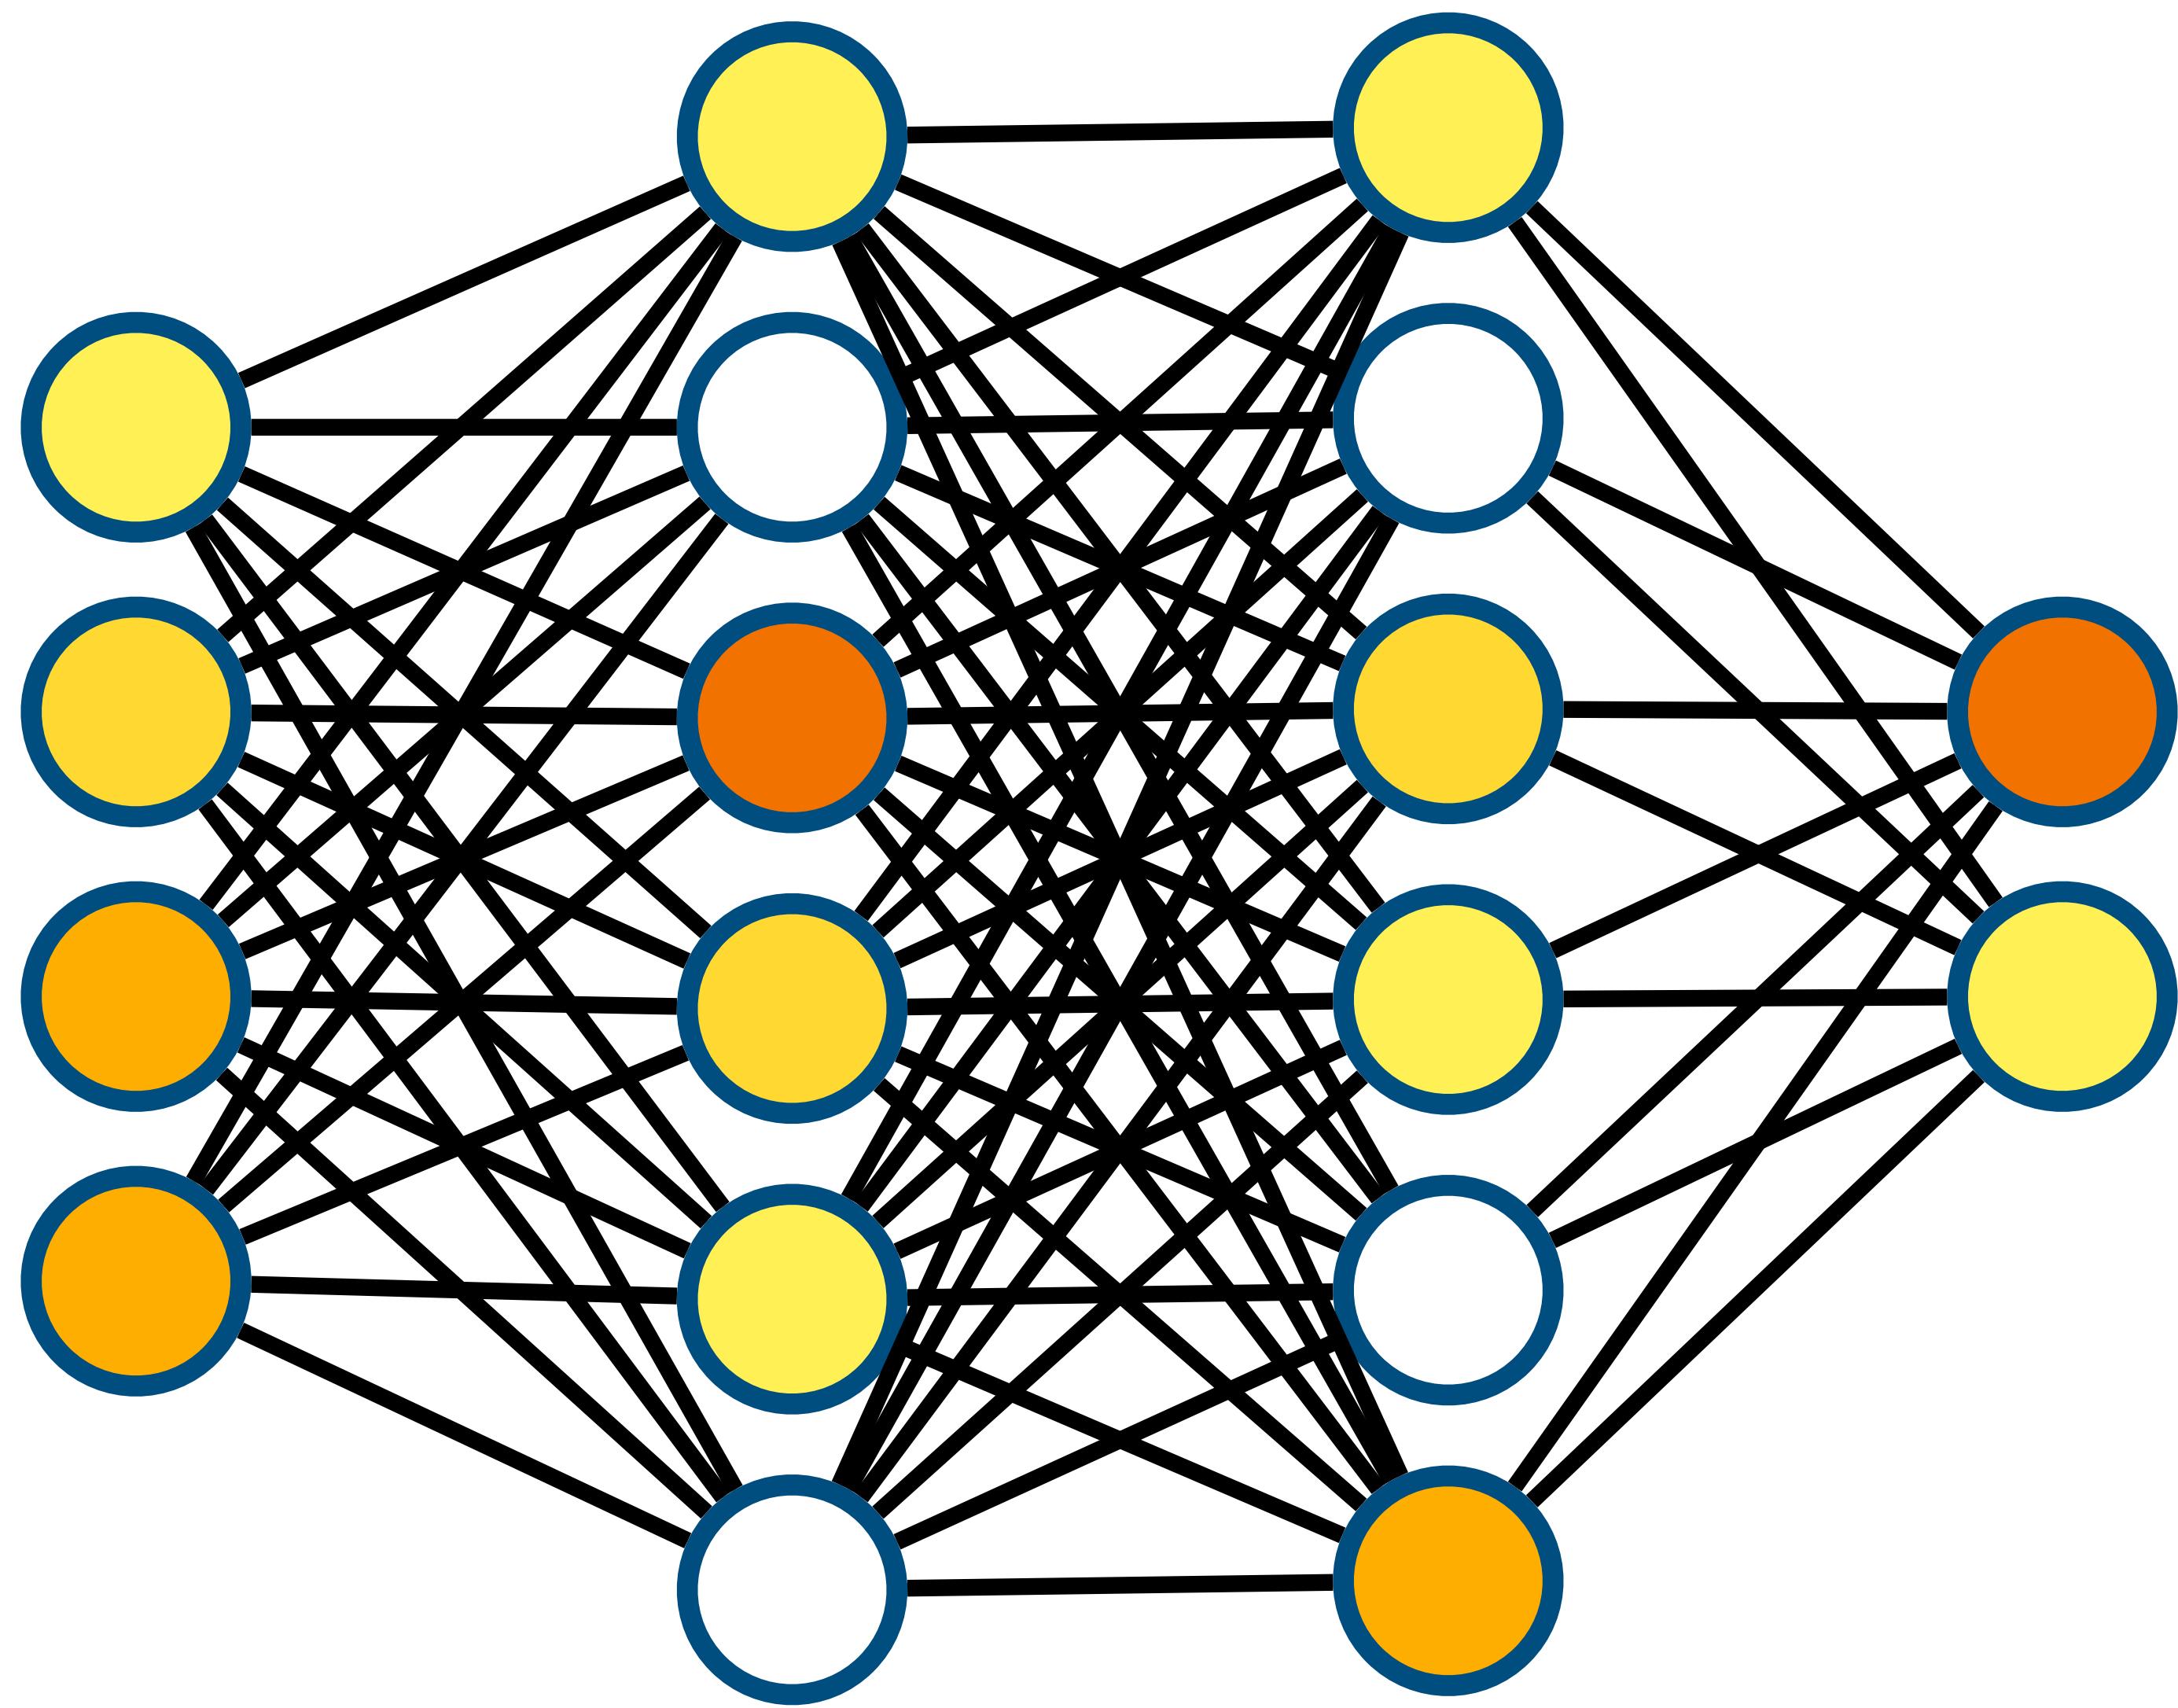


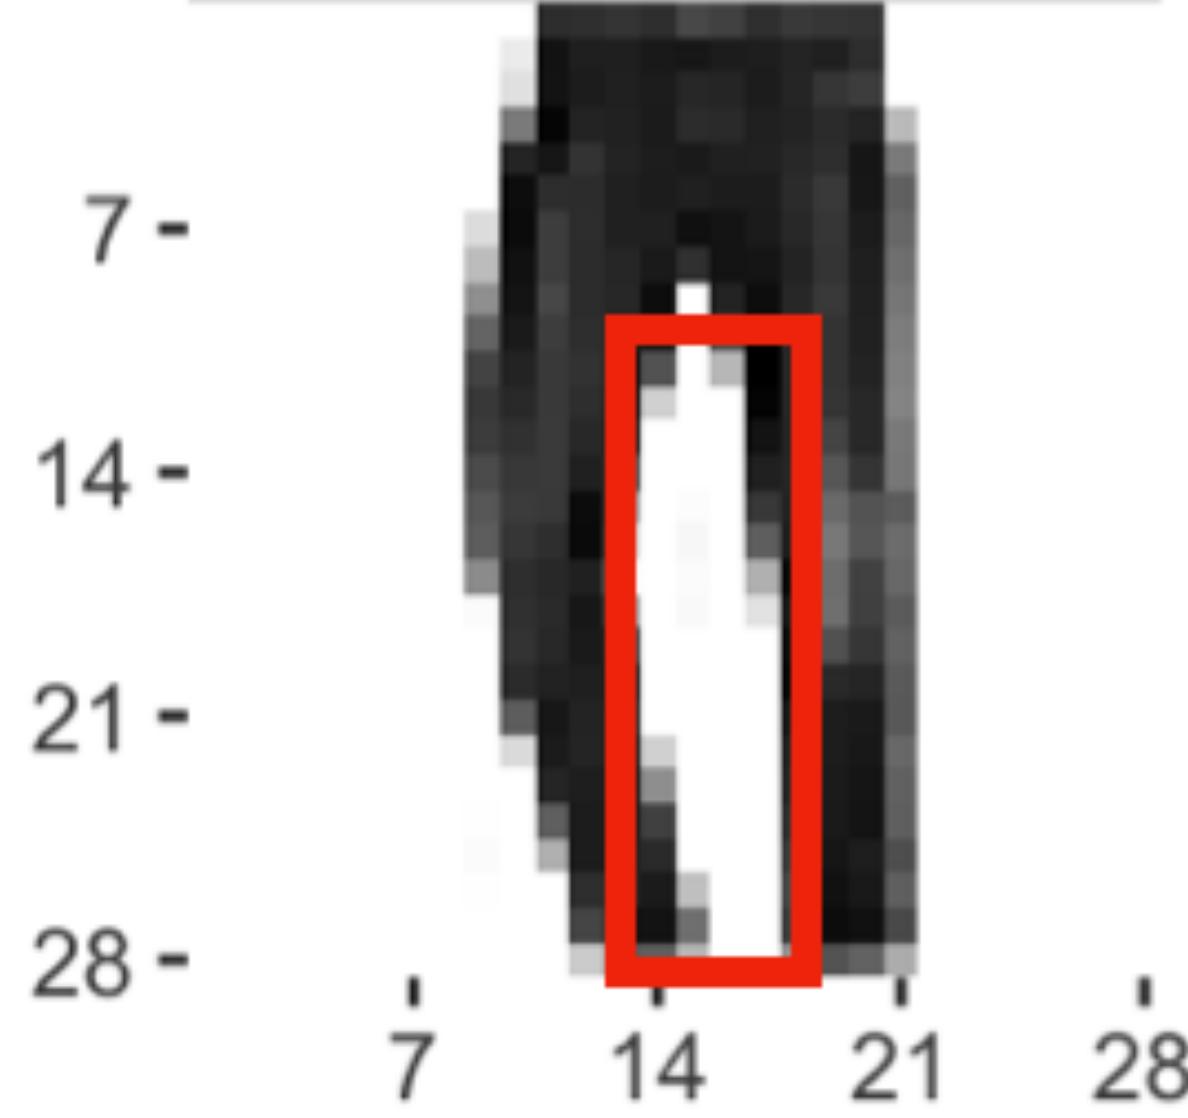
784 inputs



784 inputs



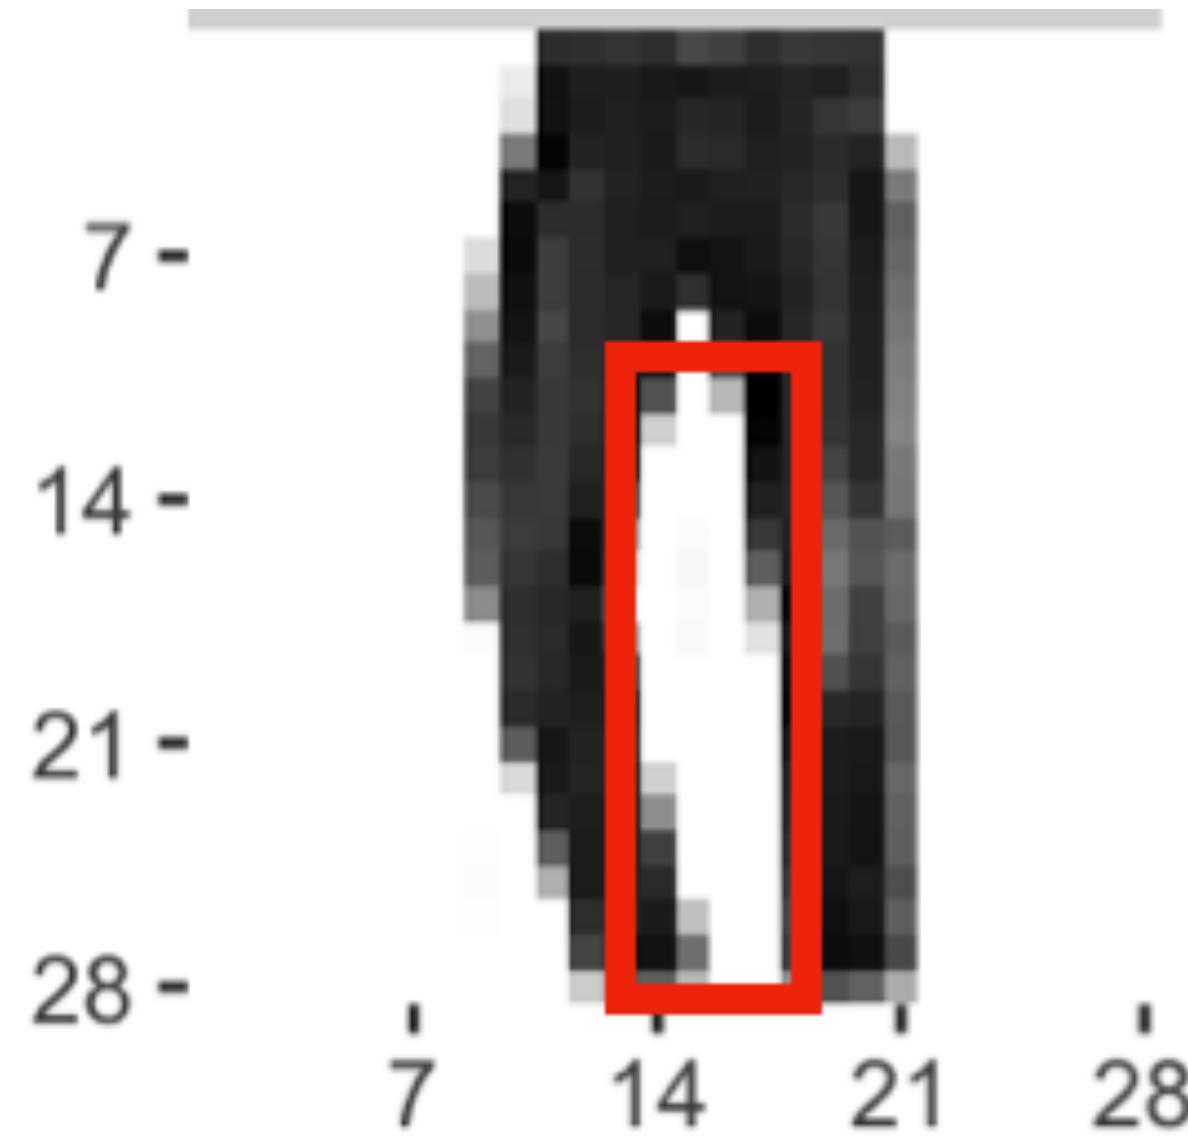




$\text{values} = (a_1, a_2, a_3 \dots a_{784})$

$\text{weighted values} = (w_1 * a_1 + w_2 * a_2, \dots, w_{784} * a_{784})$

## The bias is a correctional shift



values =  $(a_1, a_2, a_3 \dots a_{784})$

weighted values =  $(w_1 * a_1 + w_2 * a_2, \dots, w_{784} * a_{784} - \text{bias})$

through an activation function =  $f(w_1 * a_1 + w_2 * a_2, \dots, w_{784} * a_{784} - \text{bias})$

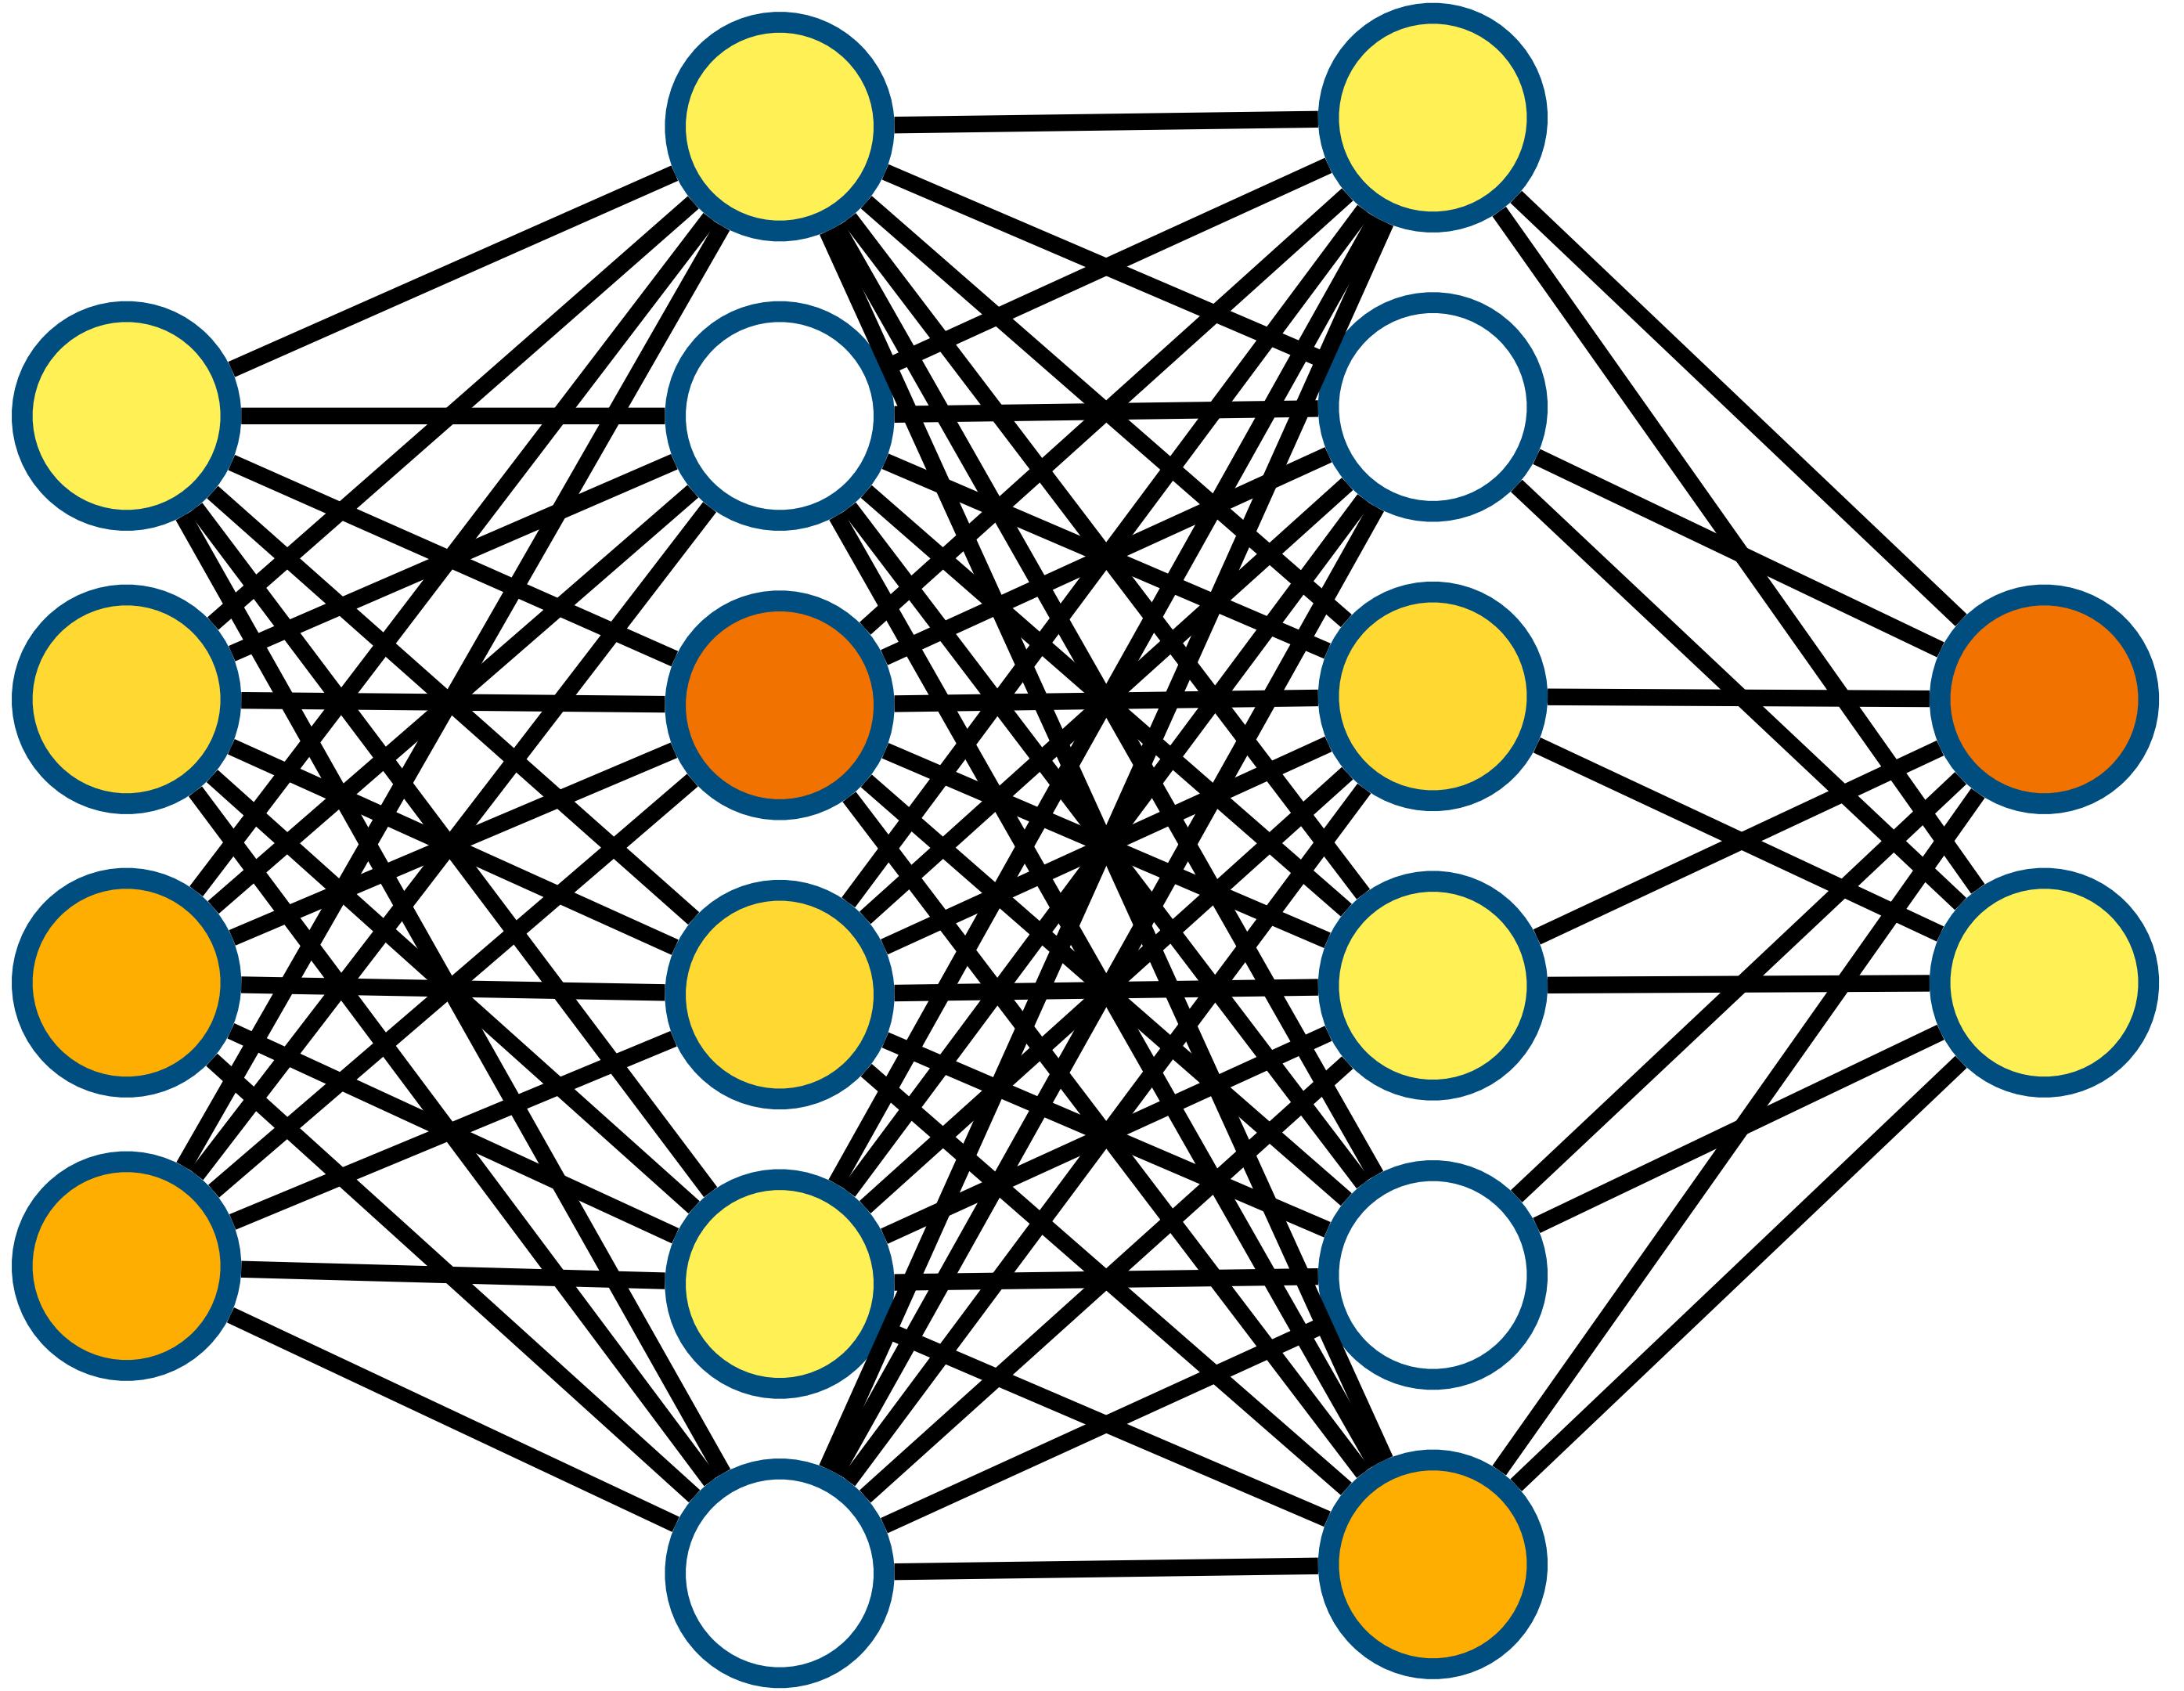
$$f(w1 * a1 + w2 * a2, ..., w784 * a784 - bias)$$

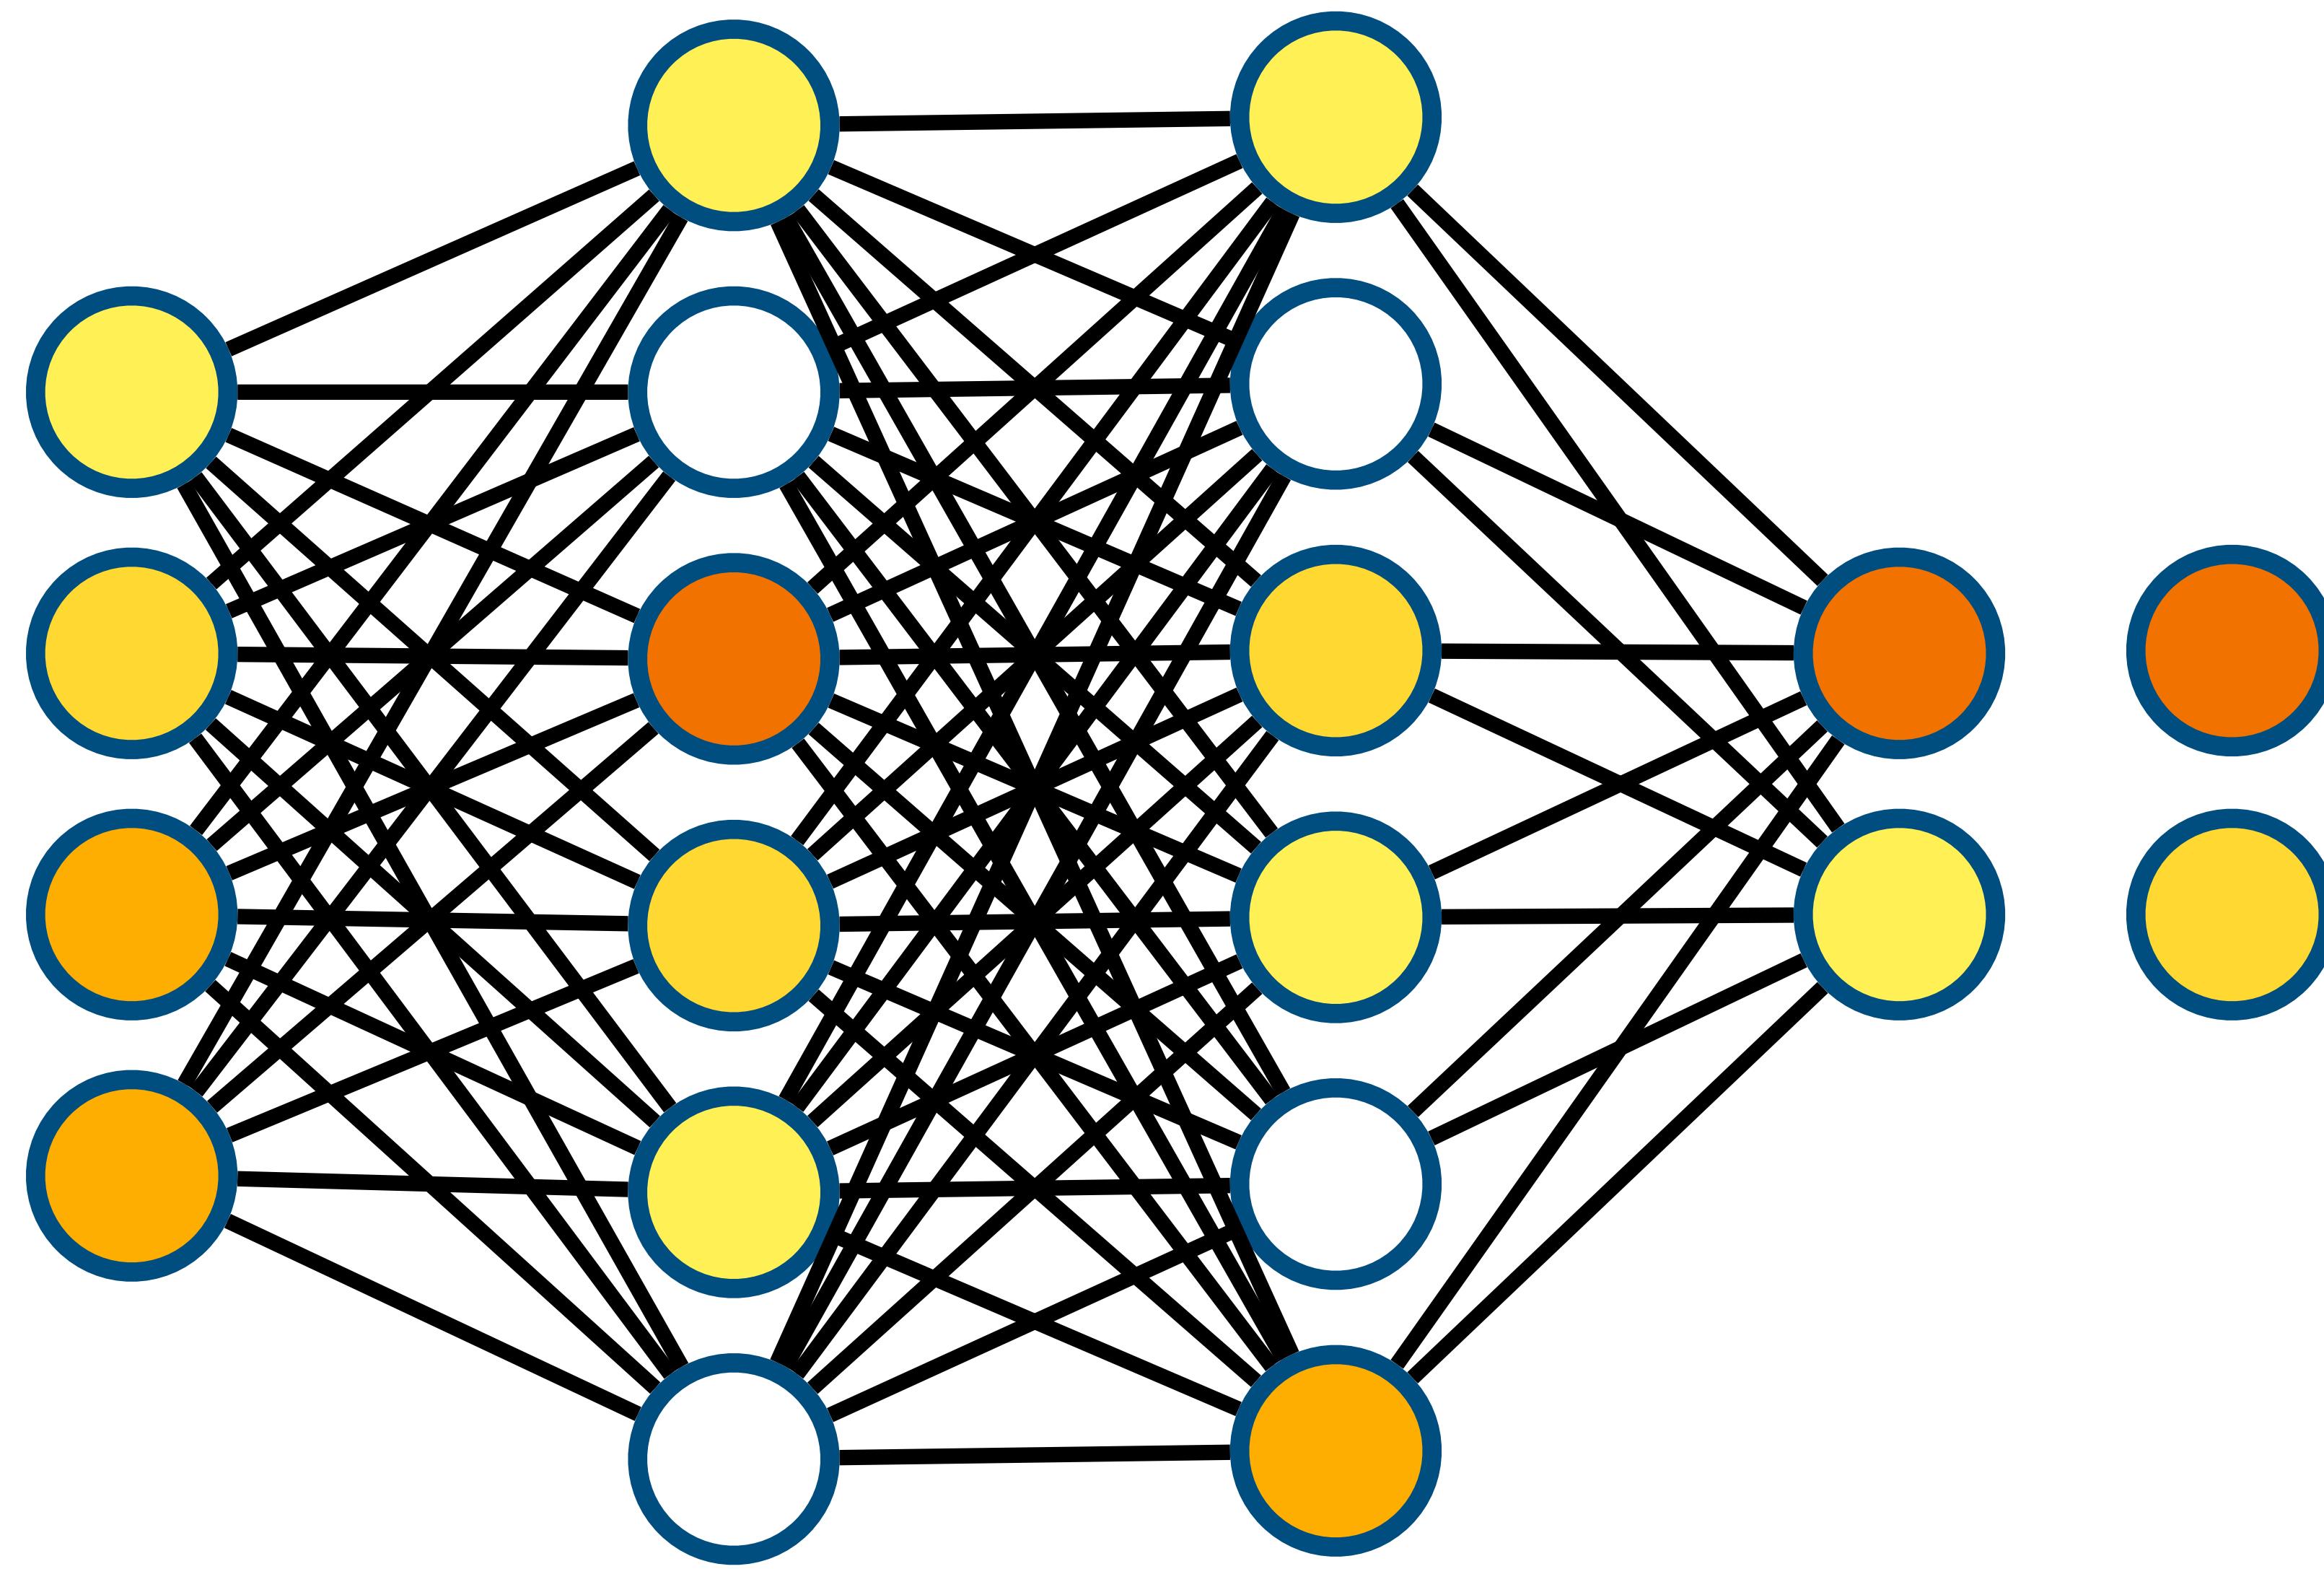
$$f(\mathbf{W}+\boldsymbol{b})$$

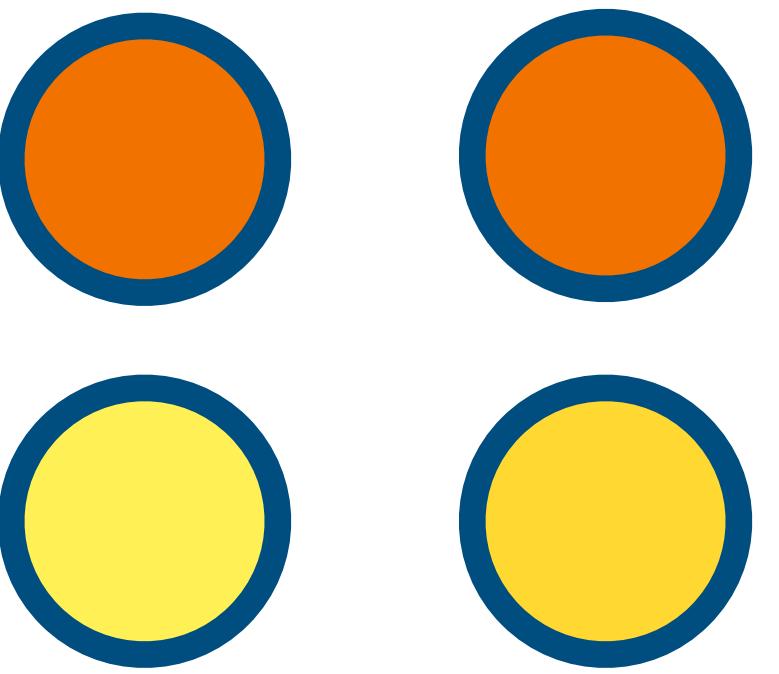
$$ReLU(w1 * a1 + w2 * a2, ..., w784 * a784 - bias)$$

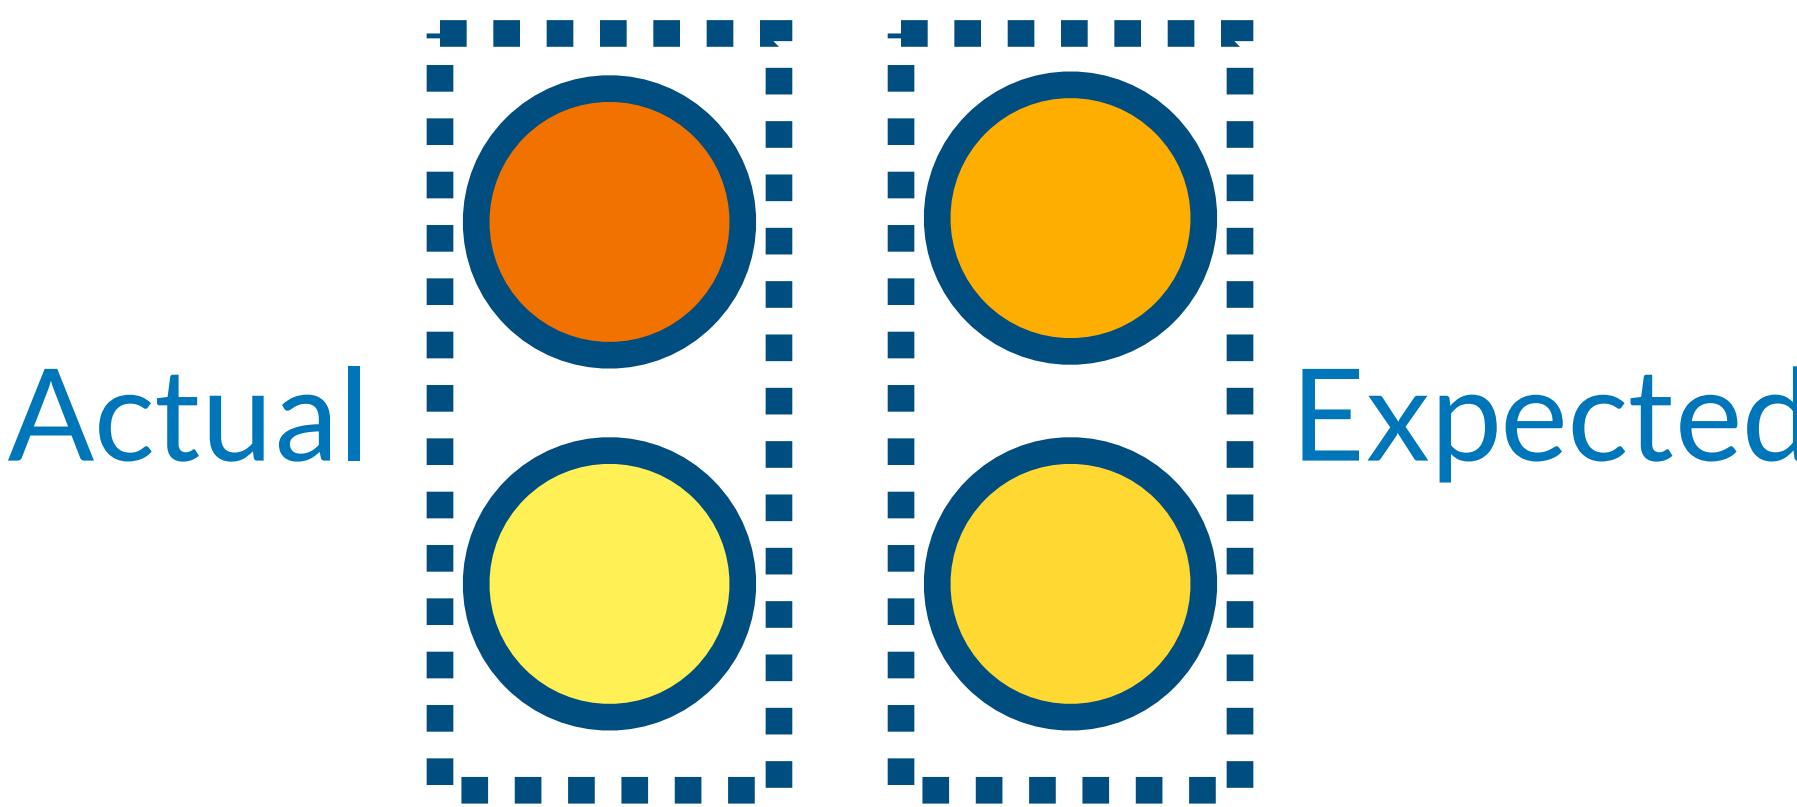
$$ReLU(\mathbf{W} + \mathbf{b})$$

# Cost Function



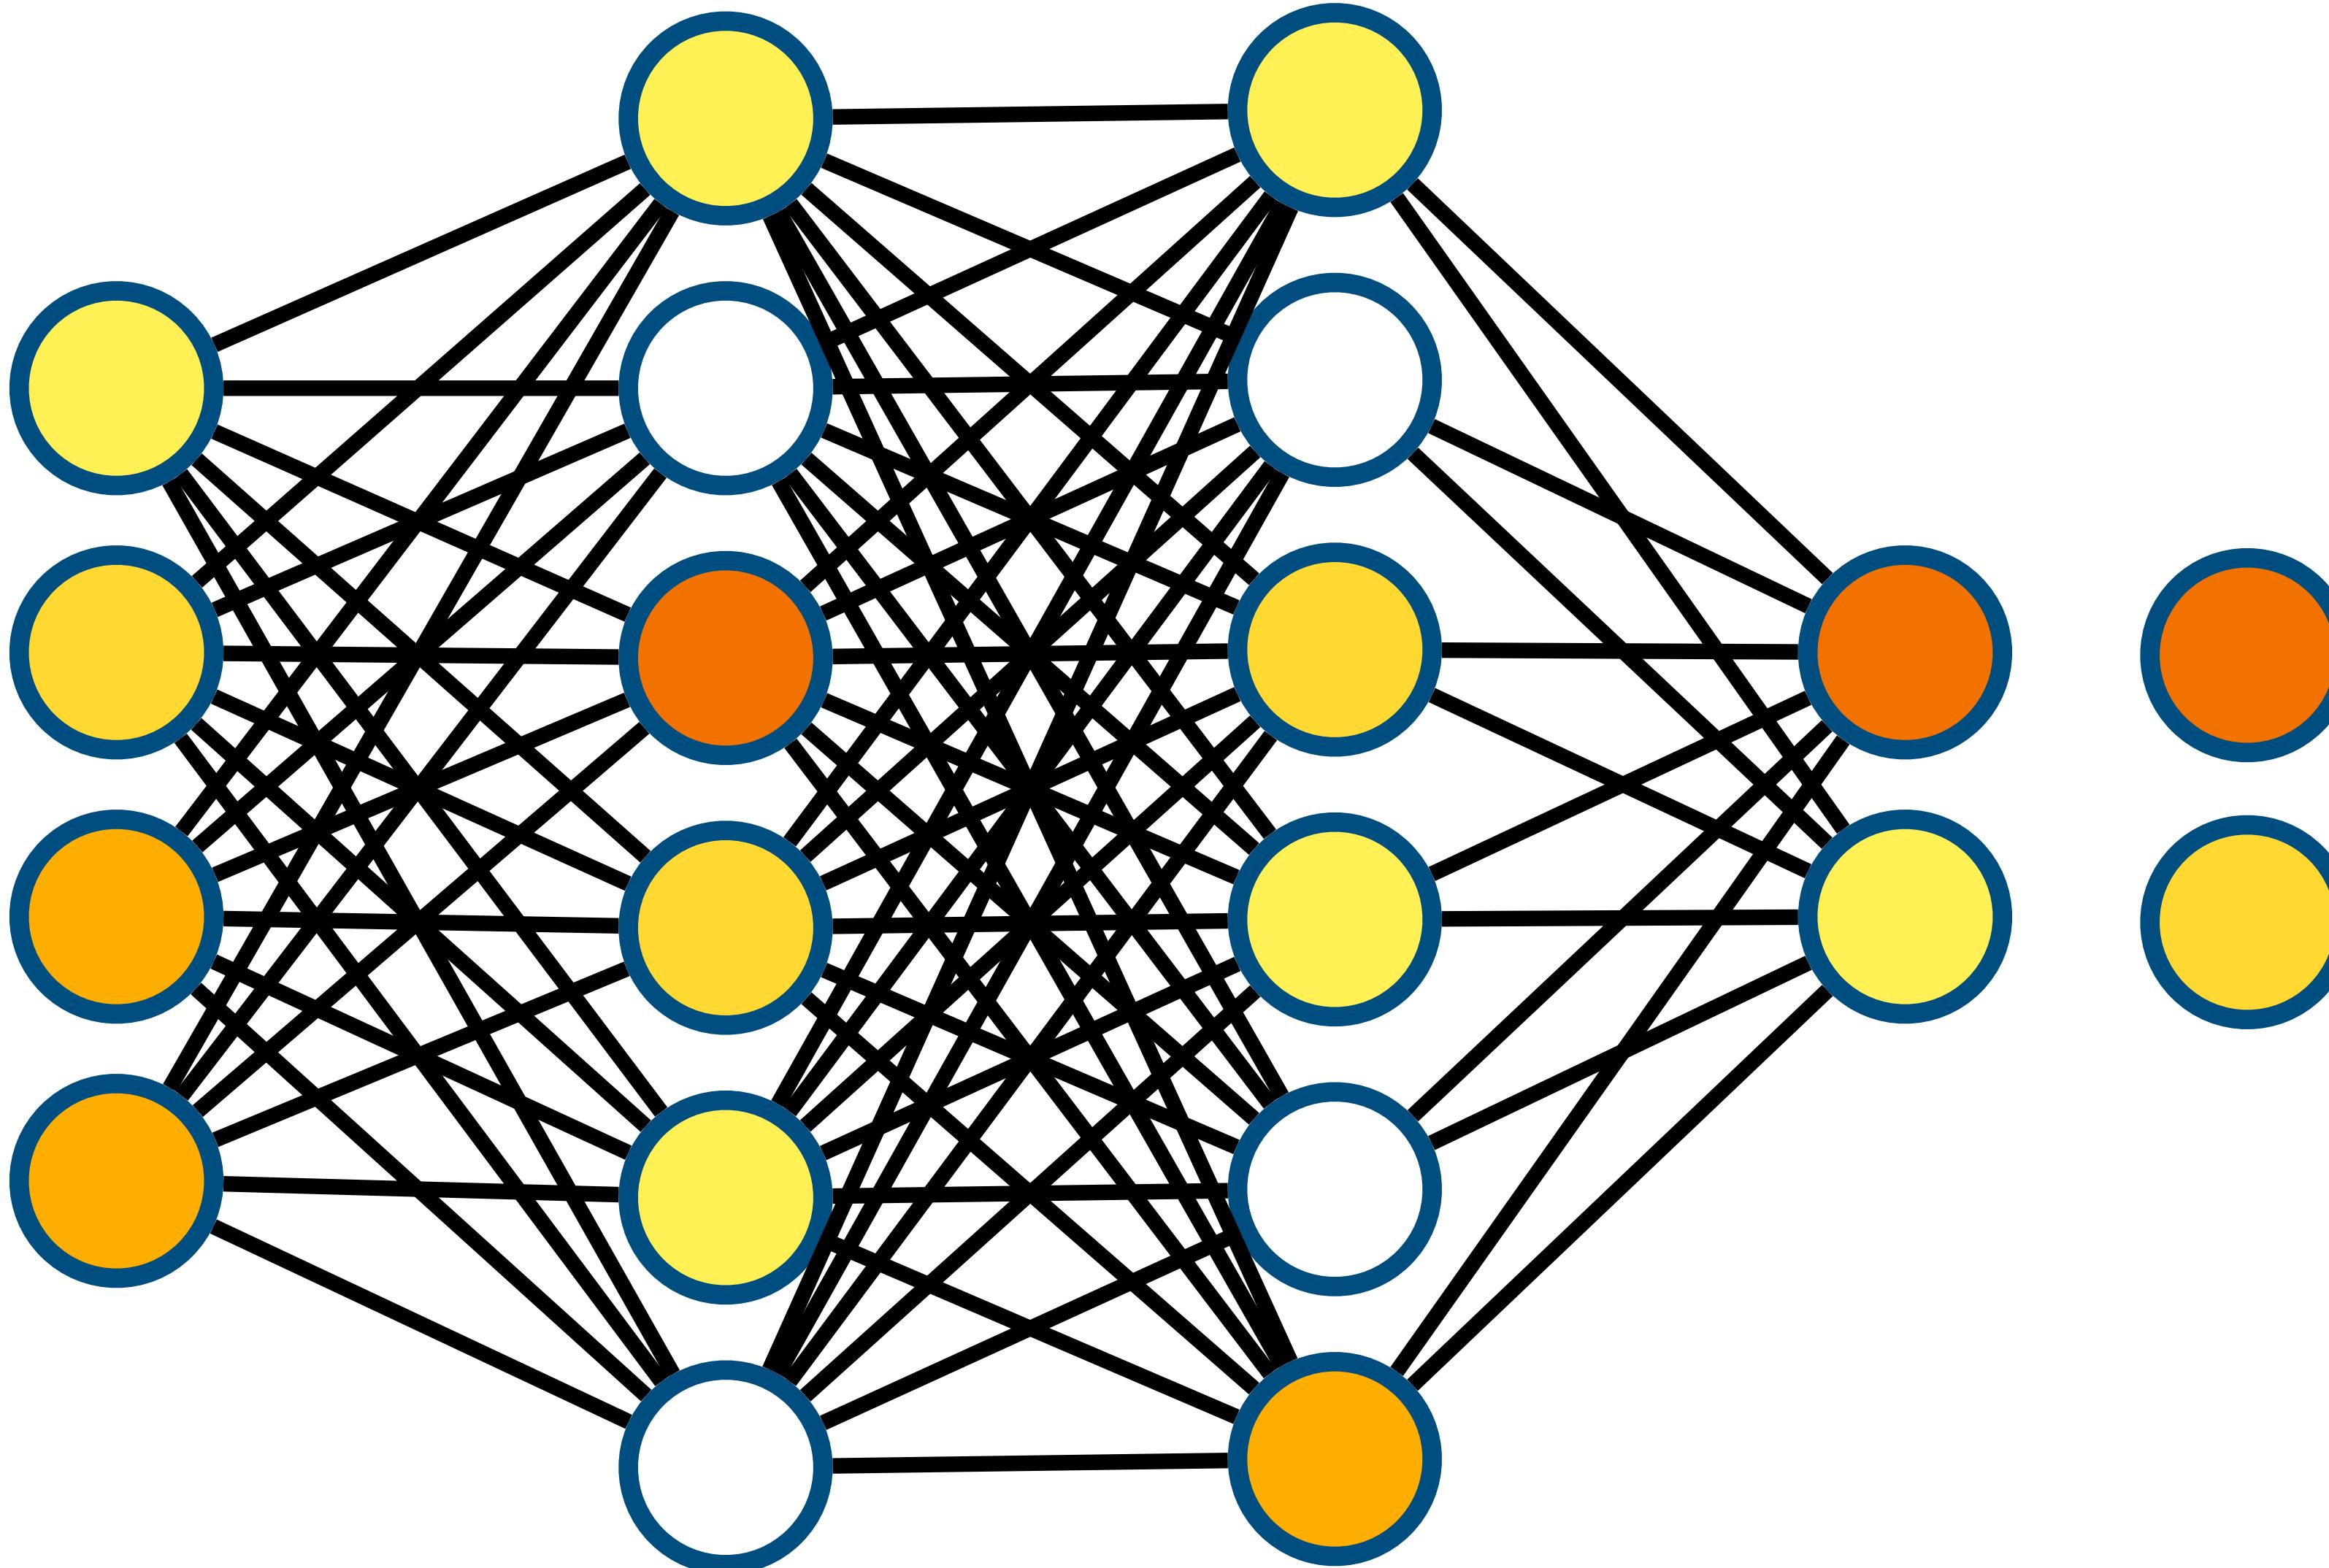




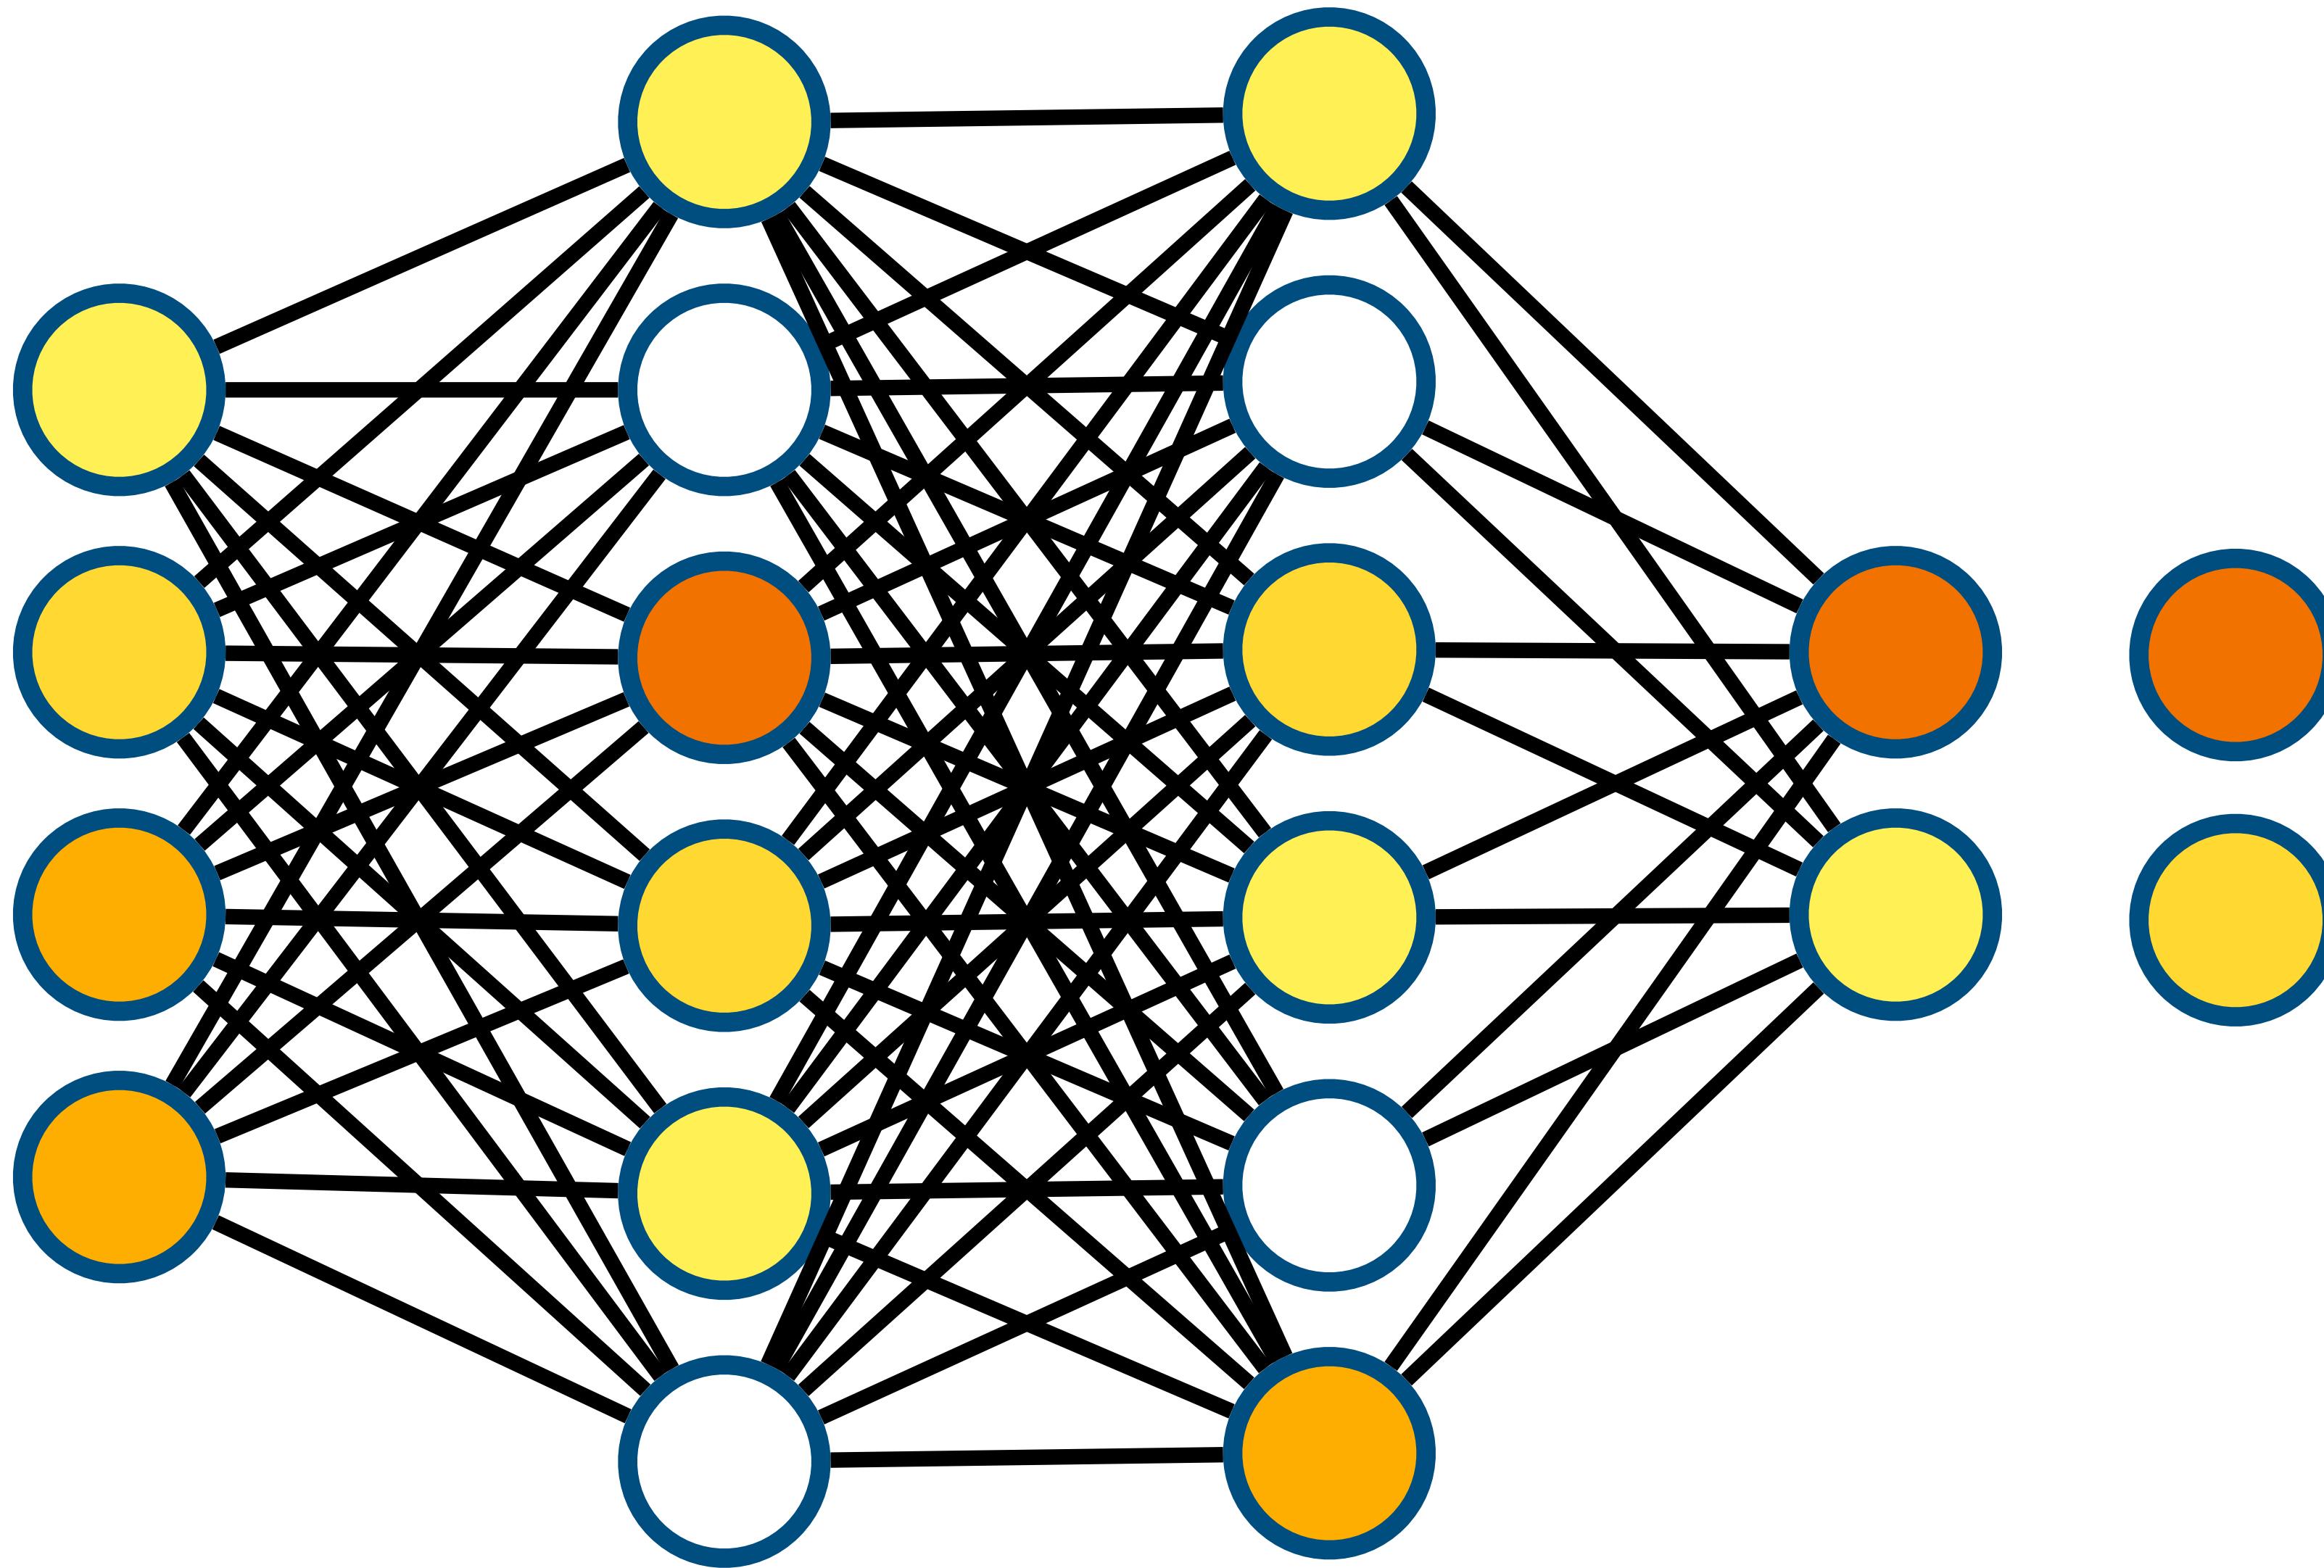


The "cost" is what is the difference:  $(\text{actual} - \text{expected})^2$

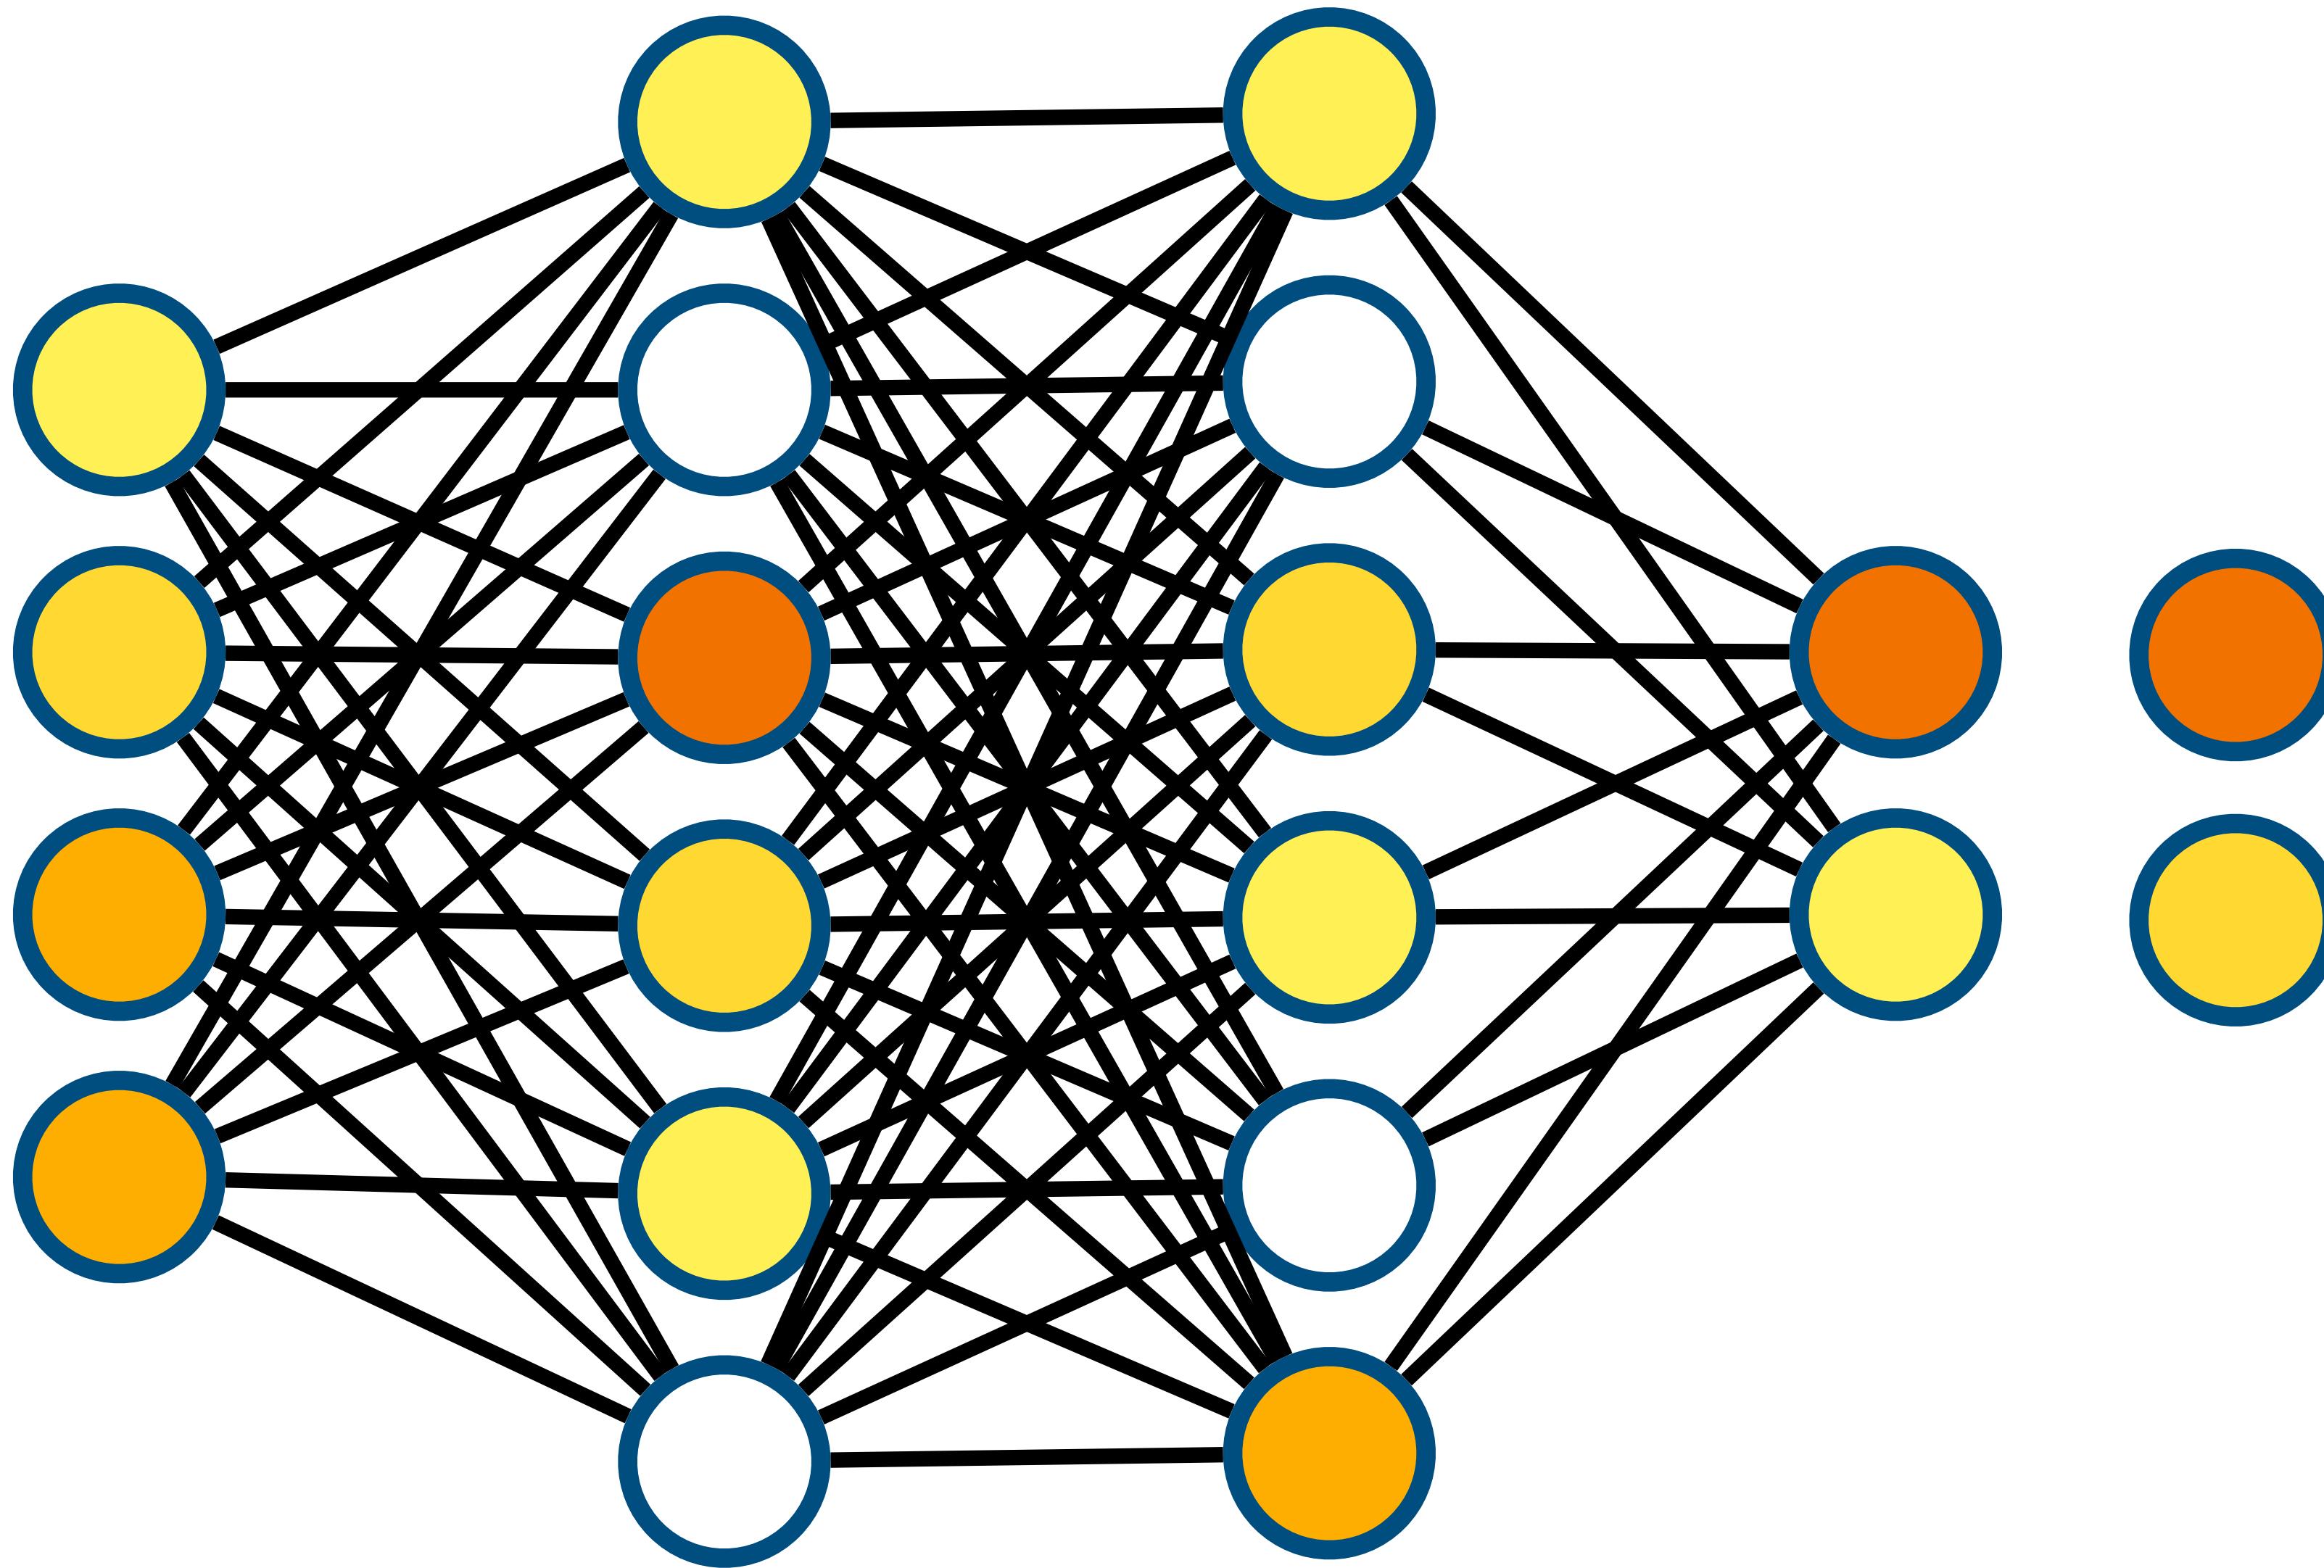
# Backpropagation



Goal: How can we minimize the cost function?



Go back and "tweak" each of the neurons and measure the response



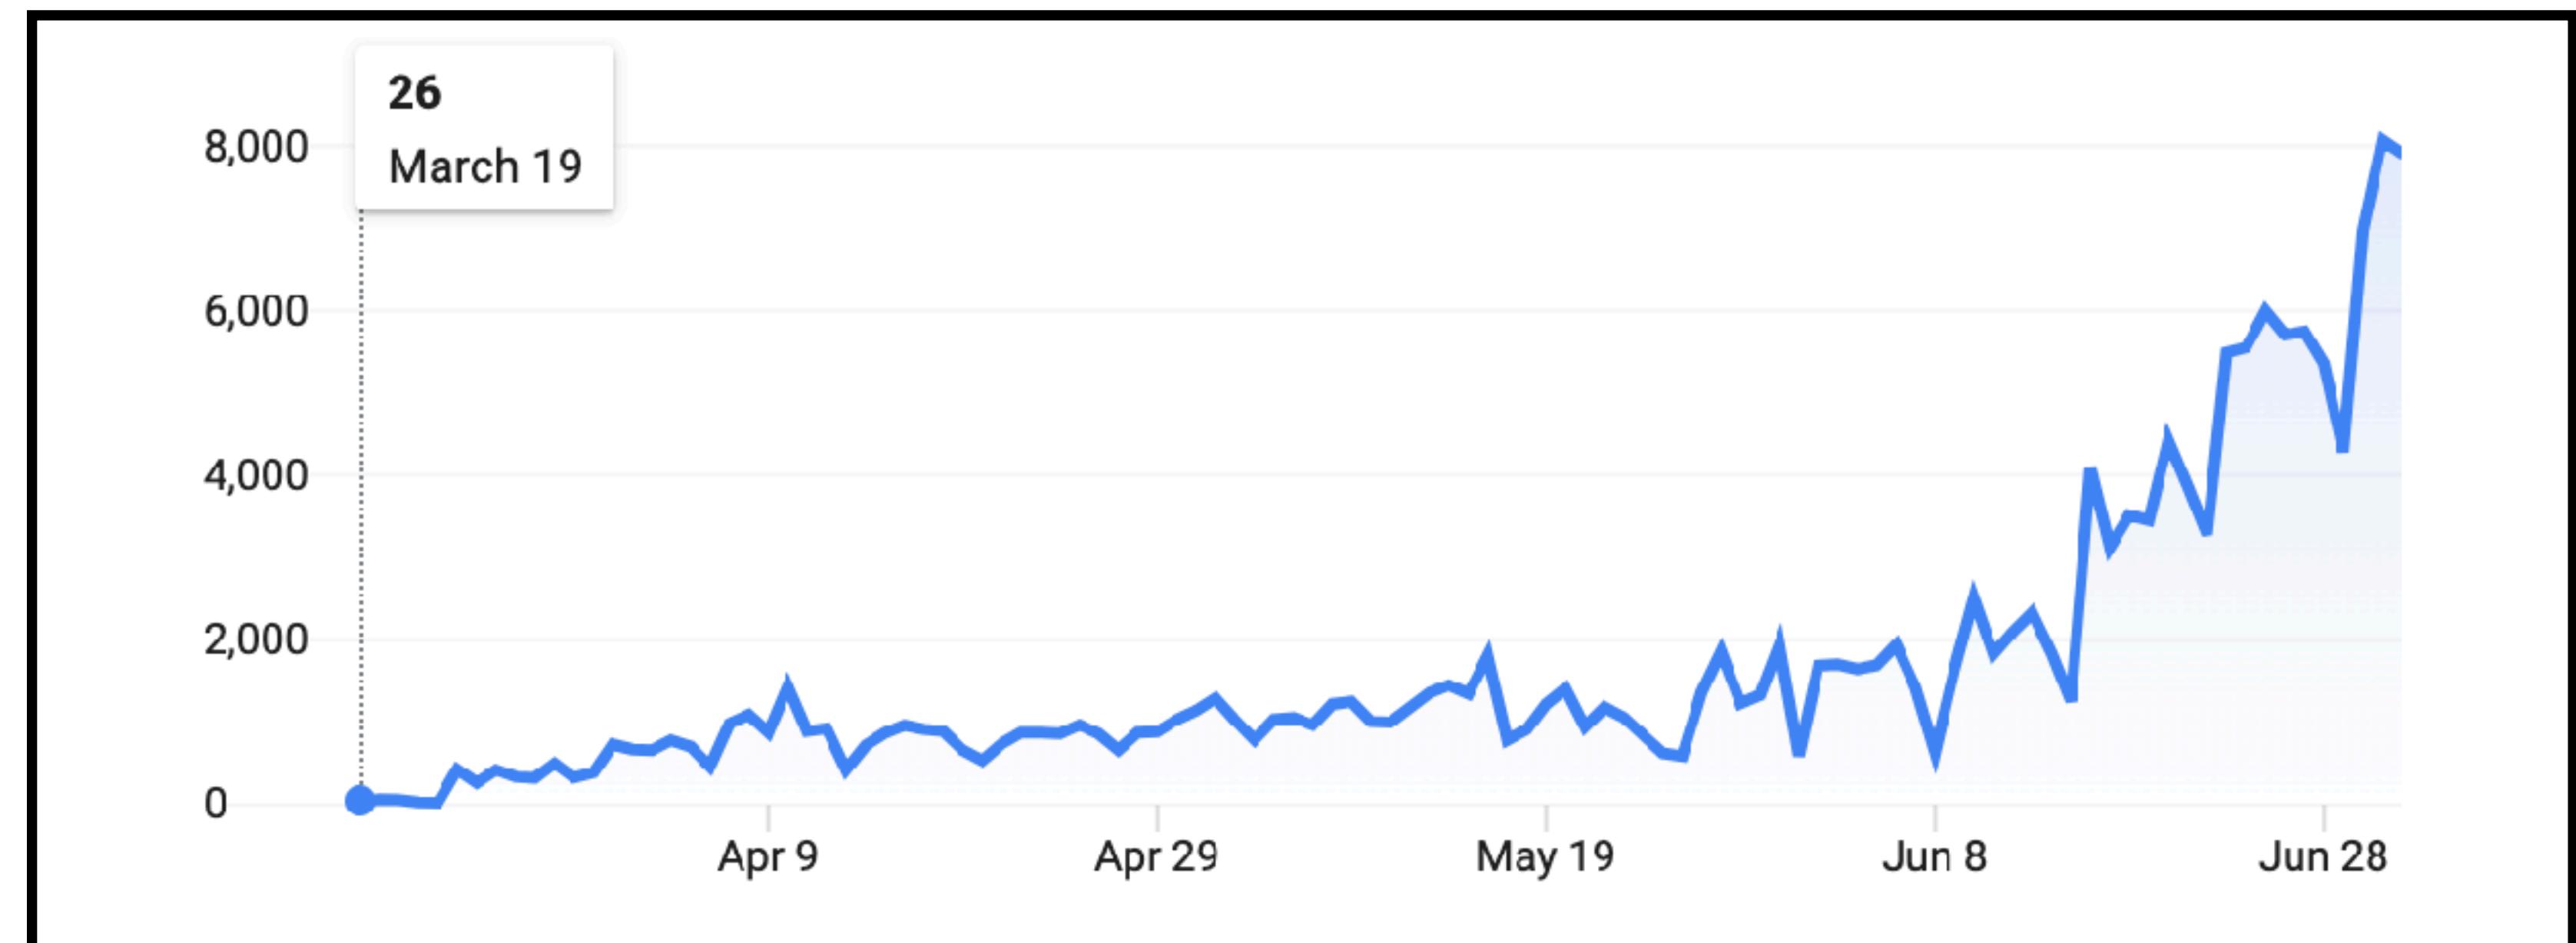
The Neural Network will determine which is more sensitive with small adjustments



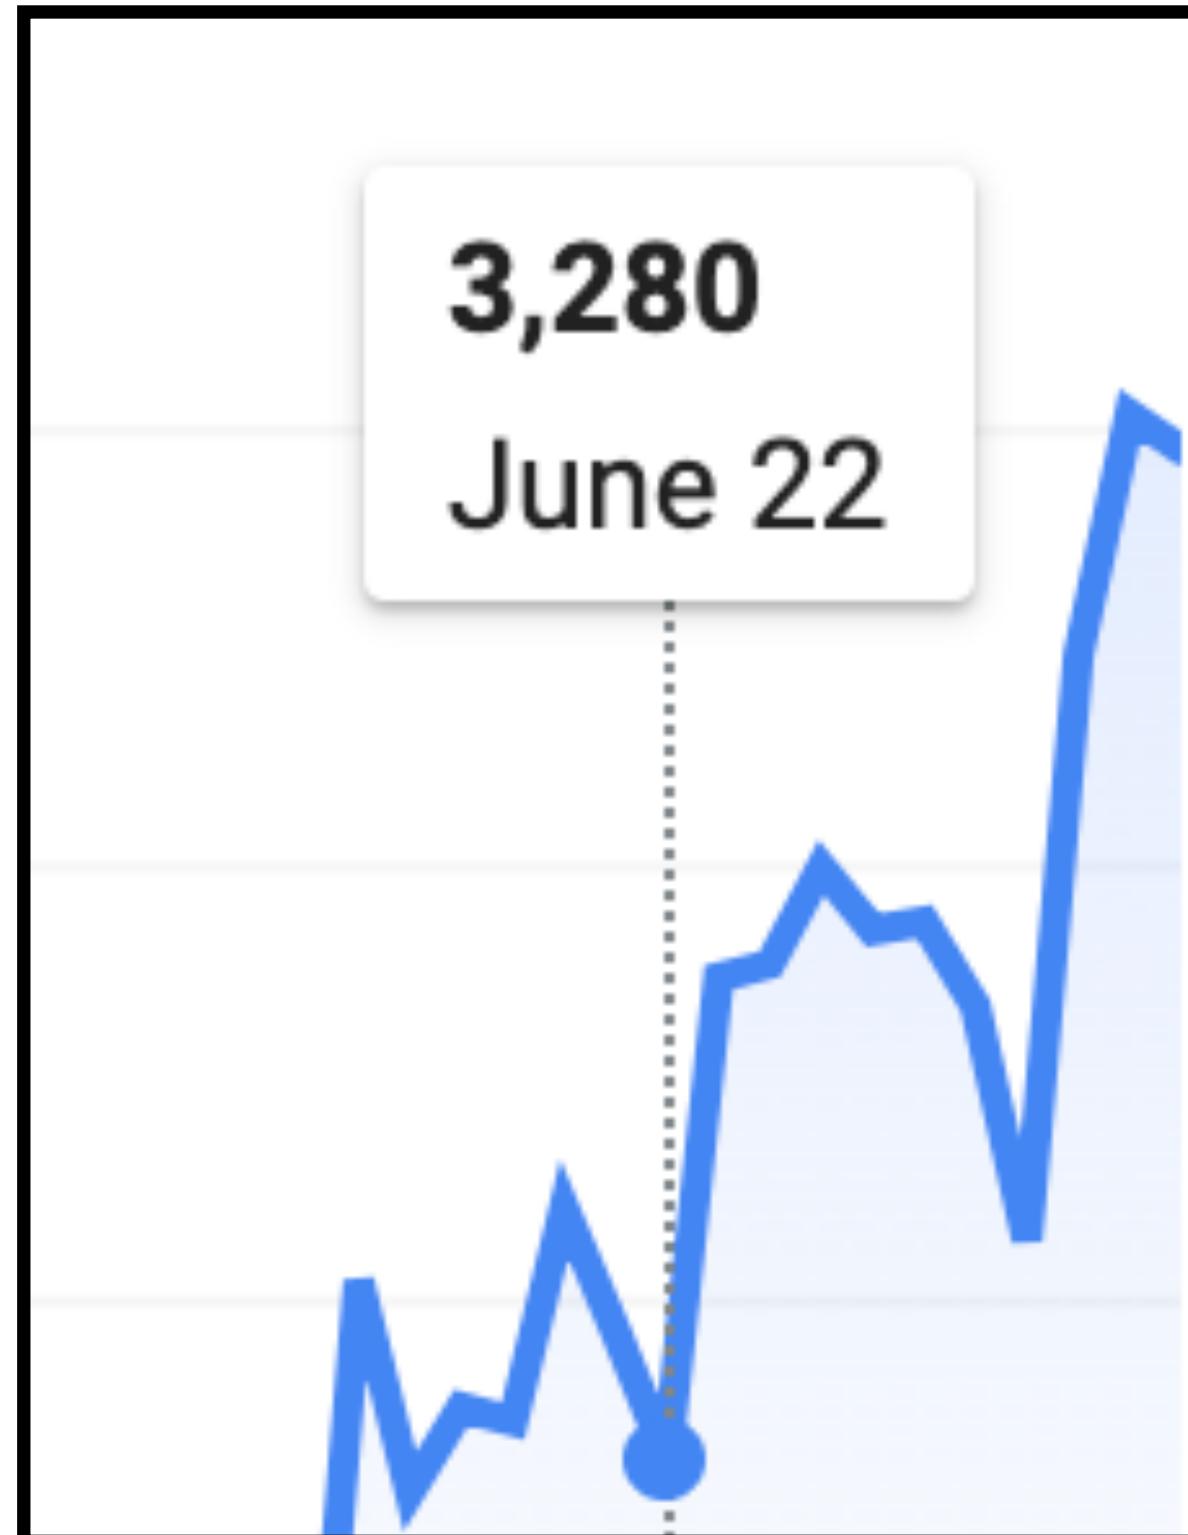
# Gradient Descent

# Local Minima/Global Minima

# Local Minima vs Global Minima

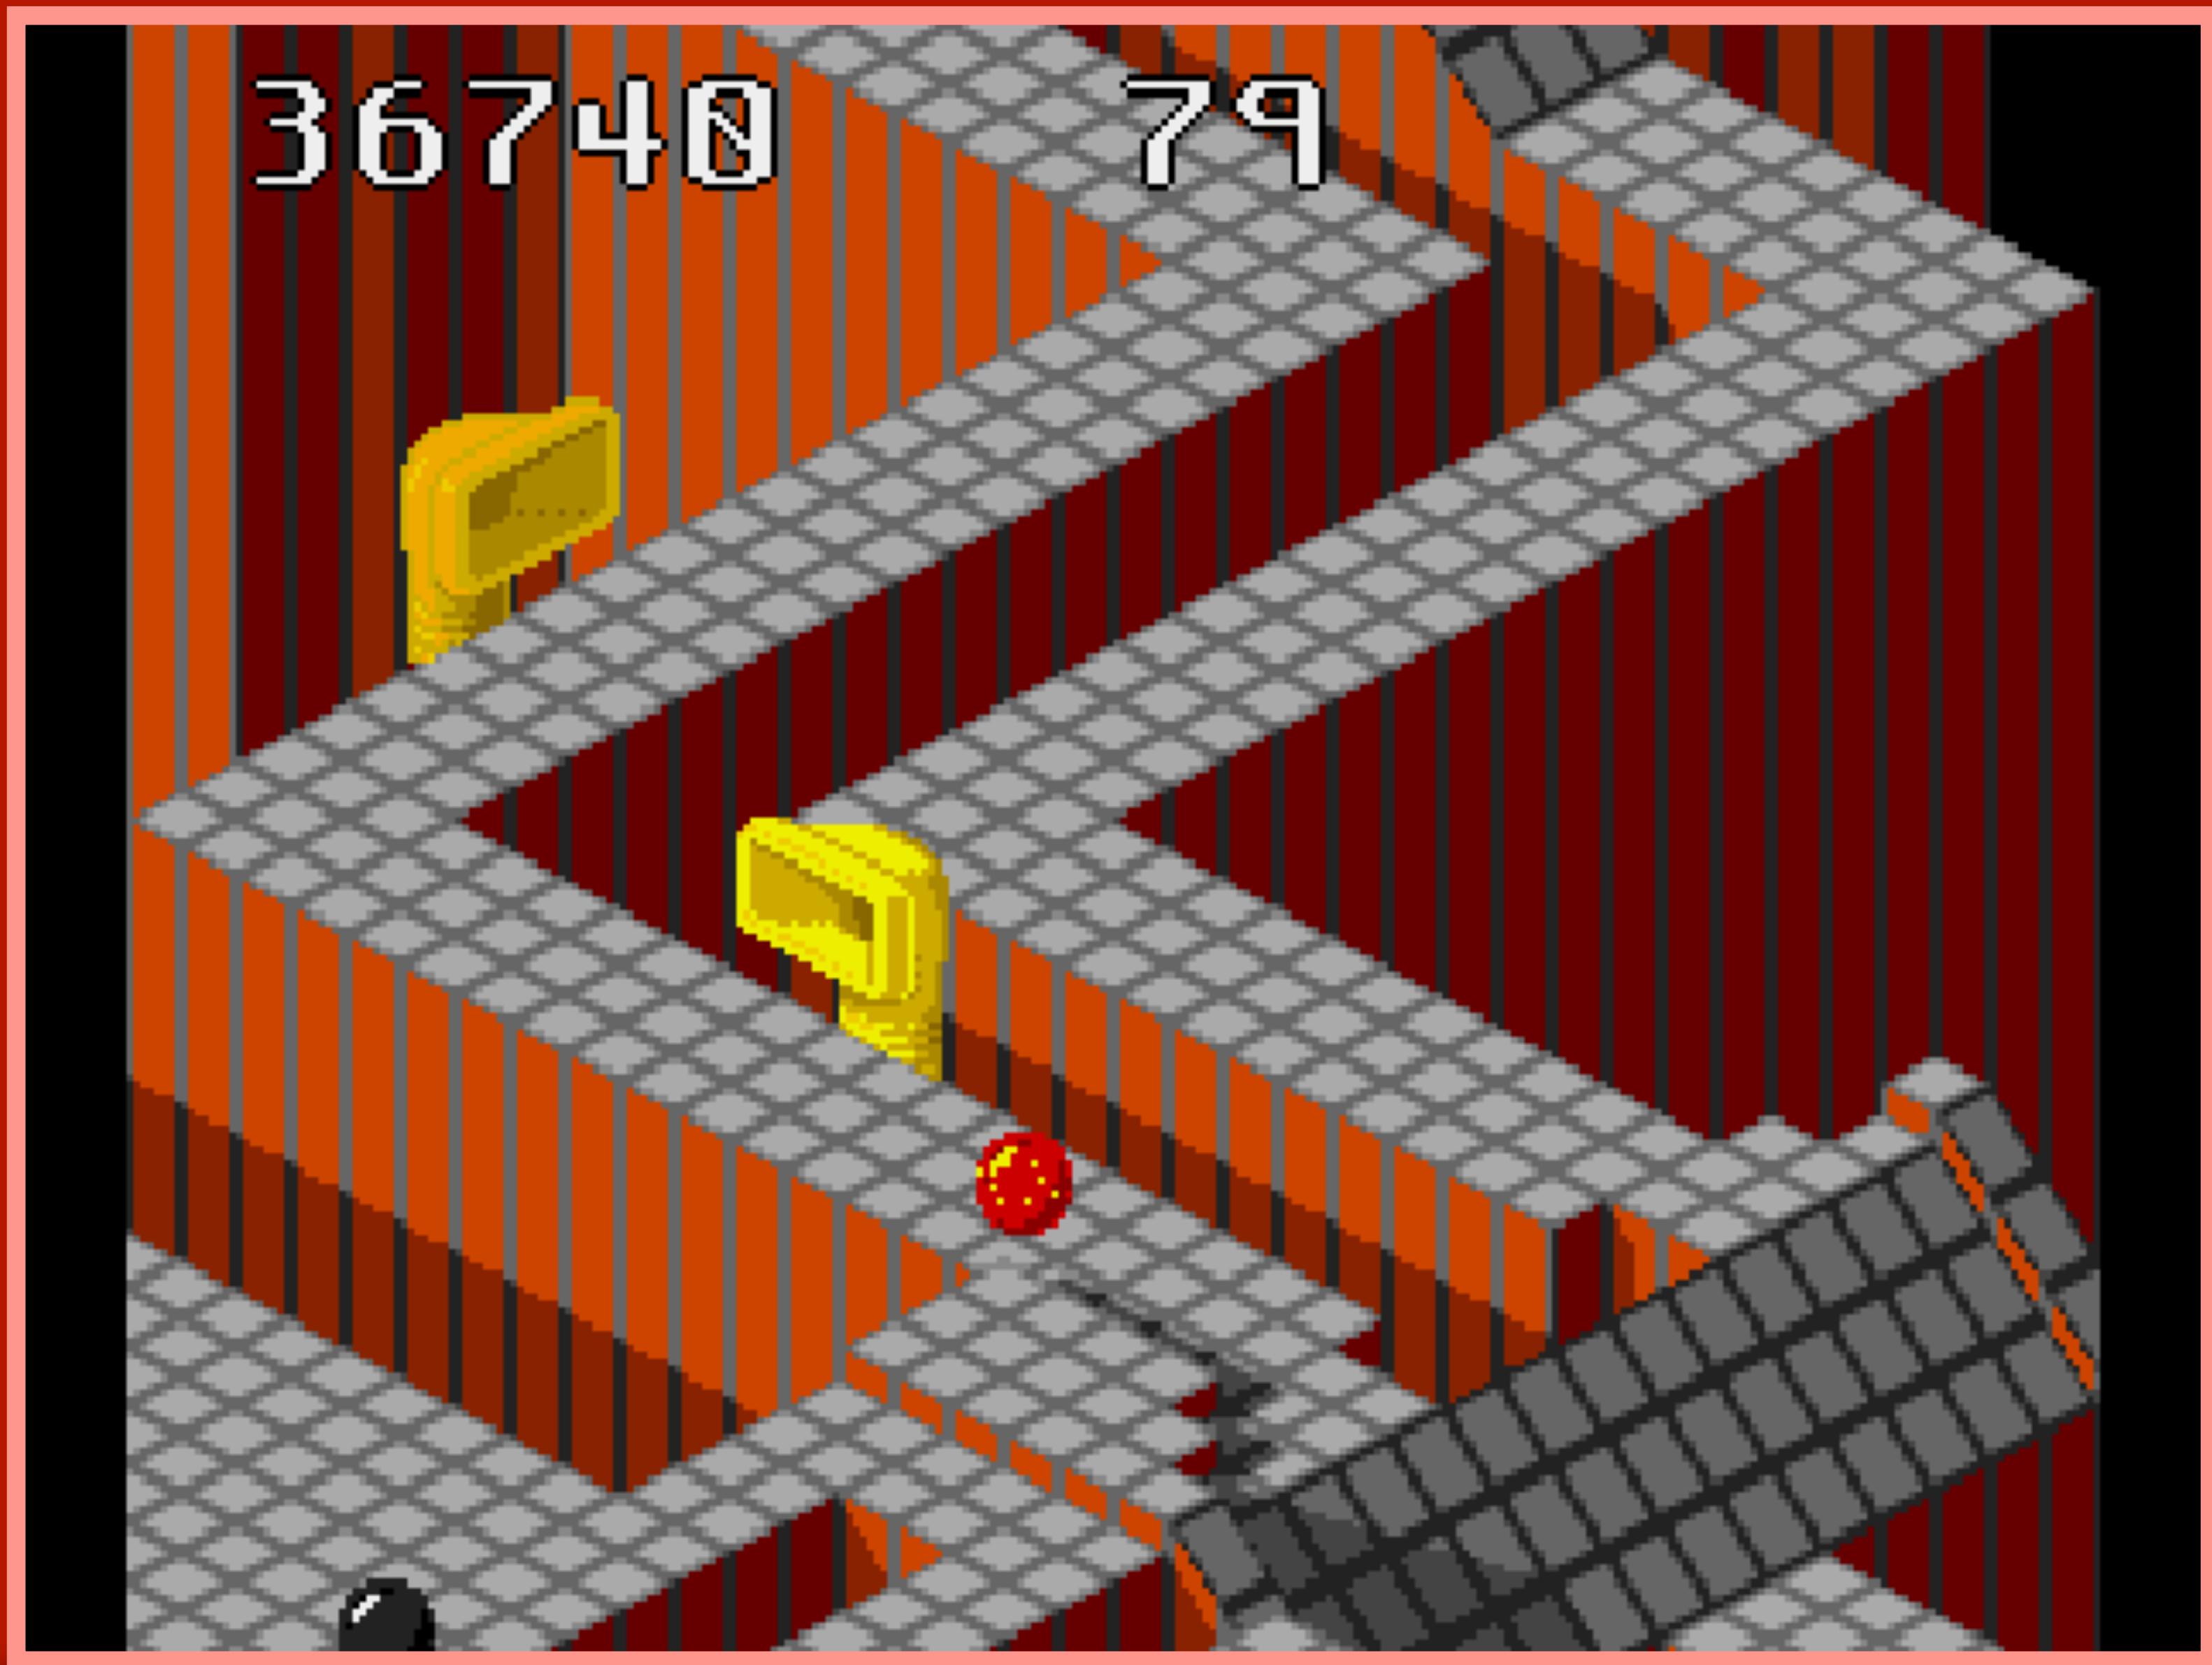


# Local Minima vs Global Minima

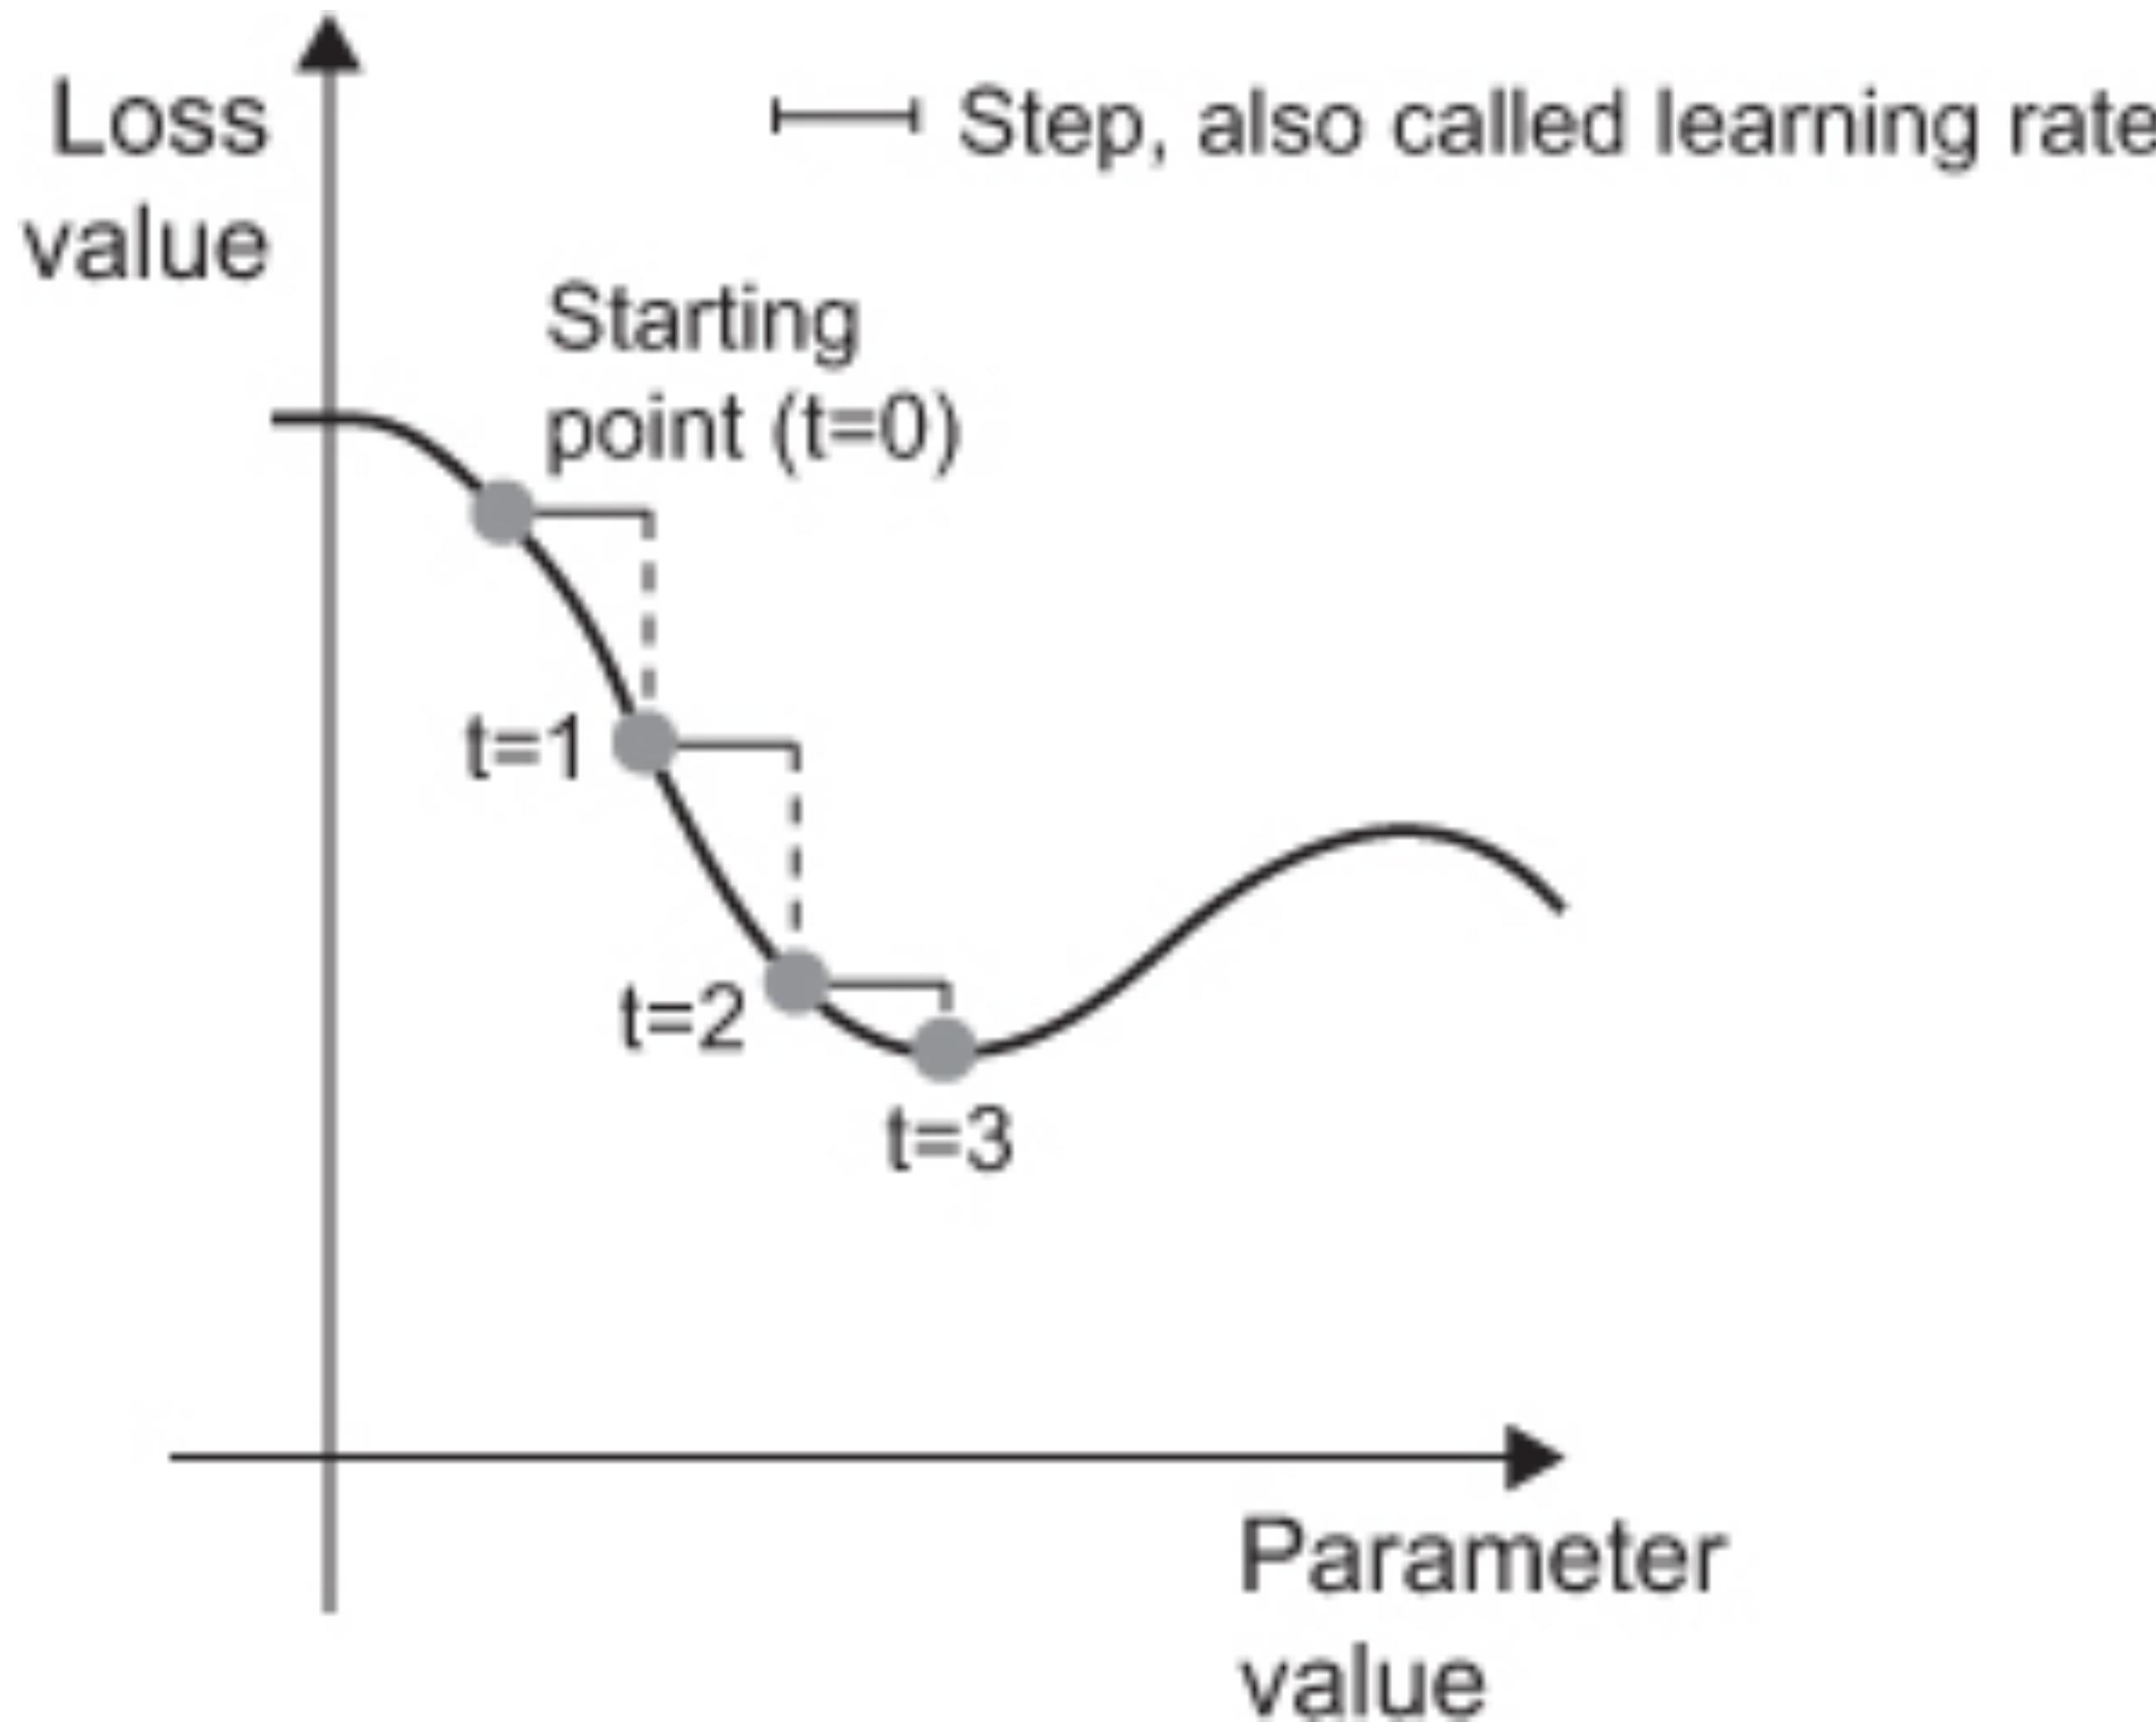


**Definition** - We pass the same dataset multiple times to the network in order to find optimal weights; one training session.

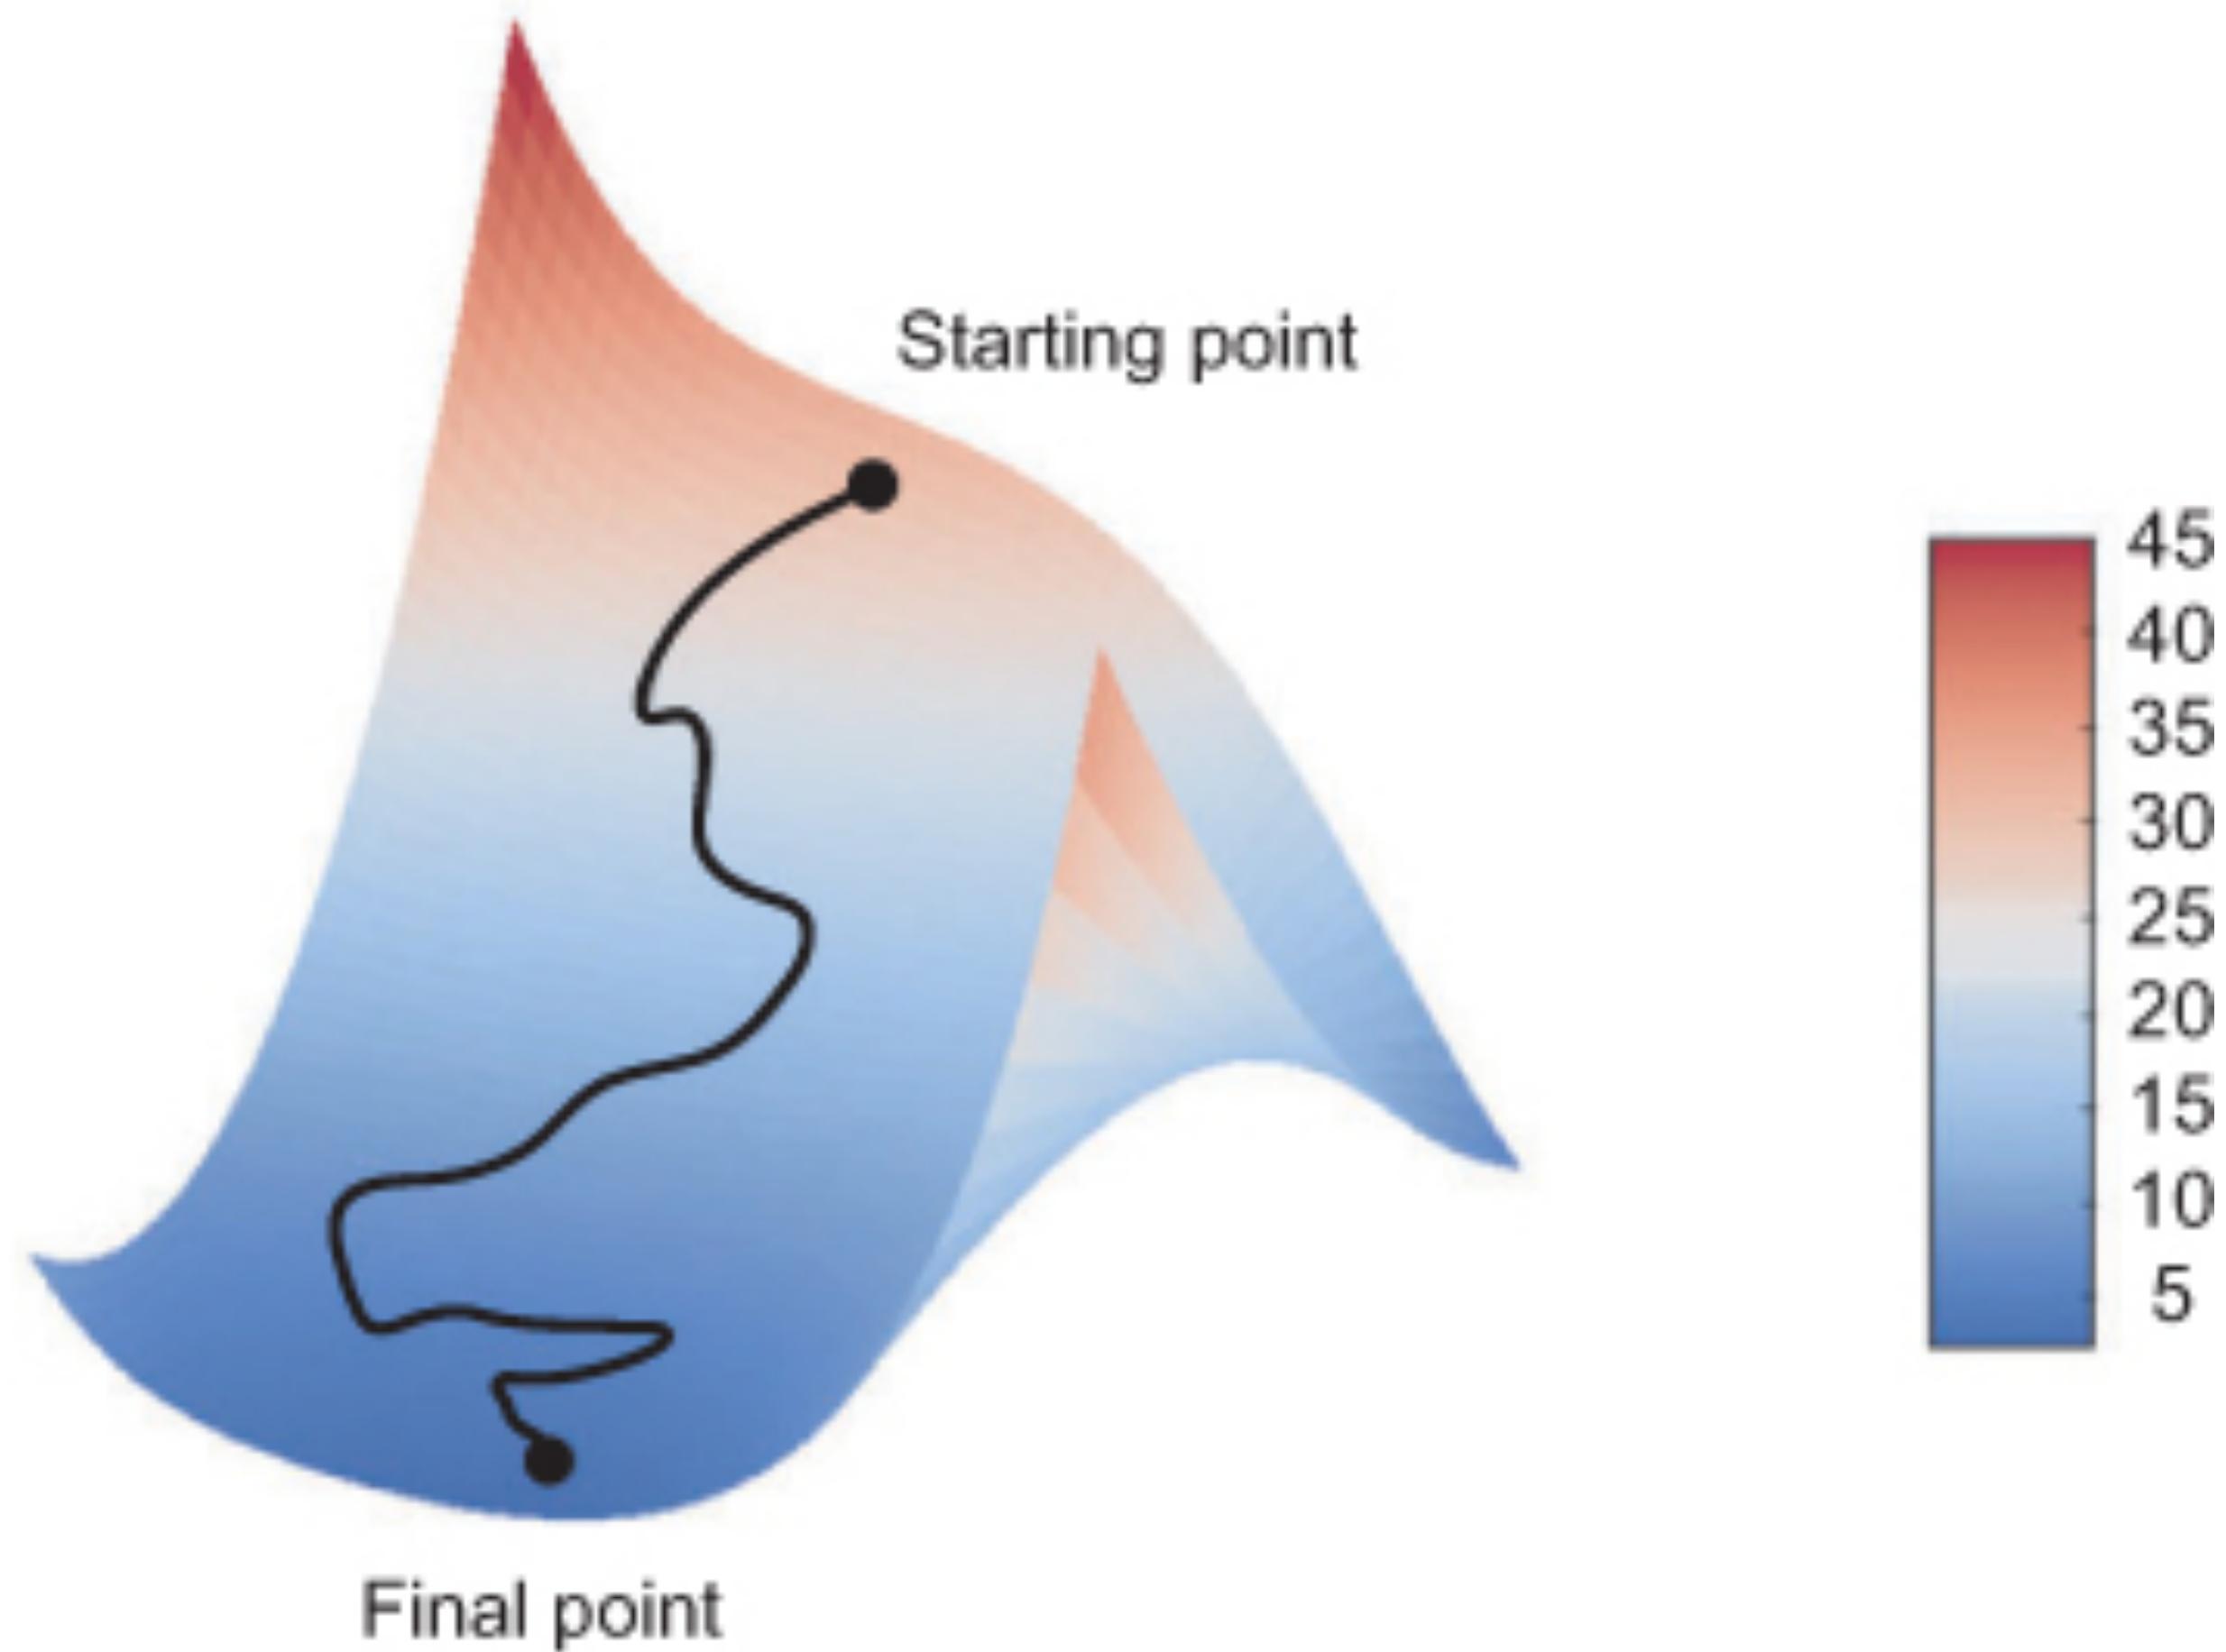
**The More the Merrier** - with just a few epochs you can reach only a local minimum, but with more epochs, you can reach a global minimum or at least a better local minimum



**Back to Gradient  
Descent**



Each Epoch we are learning and trying to lower our cost function



3d gradient descent, finding the lowest cost function



# Demo: Jupyter Lab & TensorFlow

# Other Neural Networks Patterns

# How Many Neurons?

- Determining the height, width, of a neural networks depends on the *complexity* of what you are trying to solve, the *amount of data* that you have, and the *resources* that you have.
- Adding more hidden layers increases the model's ability to **learn complex hierarchical features**
- **Increasing the number** of neurons in each layer helps the network capture more features at each level.

# Other Uses for Neural Networks

# Neural Network Types

**Neural Networks** are the backbone of many AI applications across different fields like **computer vision, language processing, healthcare, and more**. Let's explore the key types and their applications.

# Convolutional Neural Networks

- CNNs excel at image recognition, classification, and object detection
- Think of rod and cones in your eyes detecting a pattern, and think of the complexity of being able to see light vs. dark
- Applications:
  - Facial recognition (e.g., unlocking your phone)
  - Medical imaging (e.g., detecting tumors in X-rays)
  - Autonomous vehicles (e.g., recognizing pedestrians and road signs)

# Recurrent Neural Networks

- RNNs process data where order matters, like time series or language.
- State function of the input feeds back to 2nd layer so we can capture sequences of words: *We have seen this before*
- Applications:
  - Speech recognition (e.g., Siri, Google Assistant)
  - Language translation (e.g., Google Translate)
  - Stock price prediction

# Natural Language Processing

- Understanding and Generating Human Language
- This of how complex language is, “The book is on the table”. Later “It is brown”, what does “it” mean?
- Applications
  - Powered by: RNNs, Transformers, and Attention Mechanisms.
  - Chatbots (e.g., virtual assistants)
  - Sentiment analysis (e.g., detecting emotions in reviews)
  - Text generation (e.g., GPT models writing text)

# Generative Adversarial Networks

- GANs have two networks (generator + discriminator) that **create realistic data from scratch**.
- Deepfake generations
- Art and music creation (e.g., AI-generated paintings)
- Enhancing low-res images (e.g., image upscaling)

# Autoencoders

- Neural Network designed to compress data (encoding) and then reconstruct data to original
- Applications:
  - Data Compression and Anomaly Detection
  - Learn to compress and reconstruct data efficiently
  - Dimensionality reduction (e.g., simplifying complex datasets)
  - Denoising images

# Reinforcement Learning

- Decision making and control but with complex scenario.
- Train agents to make decisions by rewarding desirable outcomes. It does so by assigning credit:  
*What is the value by taking a step forward?*
- ***When teaching a teenager to drive, the teenager has a semblance of what not to do, like not drive off a cliff. ML doesn't know that.***
- Applications:
  - Game AI (e.g., AlphaGo beating human players)
  - Robotics (e.g., robots learning to walk)
  - Self-driving cars (e.g., navigating dynamic environments)

**What do Neural  
Networks have to do with  
LLMs?**

# Neural Networks and LLMS

- Neural Networks are the Foundation of LLMs: LLMs (Large Language Models) are built on deep neural networks—specifically, architectures like Transformers.
- Neural networks enable LLMs to learn patterns in language, grammar, and context from massive datasets.
- Scale and Complexity: LLMs have billions of parameters and require massive computing power to process and learn from large datasets.
- Models like ChatGPT analyze vast amounts of data to capture intricate patterns in language, enabling advanced capabilities like context understanding and text generation.

# How Neural Networks Power LLMs

- LLMs use deep architectures with millions or billions of parameters.
- Each layer of the neural network refines its understanding of the data.
- Transformers rely on self-attention mechanisms to understand relationships between words in a sentence, regardless of their position.
- This structure allows LLMs to handle complex language tasks like translation, summarization, and text generation.

# Transformers?

- Instead of processing words one by one (like older models), **self-attention** lets the model weigh the importance of all words in a sentence simultaneously.
- Example: In the sentence “The cat sat on the mat because it was tired,” the model uses self-attention to figure out that “it” refers to “the cat.”



<https://medium.com/@sharanharsoor/mastering-transformers-a-comprehensive-guide-to-transformer-architecture-questions-7694e52ba949>

# Transformers

- Large Scale Modeling, allows use to parallel to larger scale to Process and generate complex human language using attention mechanisms.
- Attention Mechanisms are relationships of words and tokens
- Applications:
  - Algorithms like GPT-4, BERT (*which is not an LLM, but a Transformer Architecture*)
  - Machine translation (e.g., Google Translate)
  - Summarization and Q&A Systems

**So, wait, what is  
difference between ML  
and AI?**

**Ai is about replicating  
human intelligence**

**ML is about learning from  
your data**

# Conclusion

# Thank You



- Email: [dhinojosa@evolutionnext.com](mailto:dhinojosa@evolutionnext.com)
- Github: <https://www.github.com/dhinojosa>
- Linked In: <http://www.linkedin.com/in/dhevolutionnext>